

V.F.S. v2.0 (Open-Gate)

Folding, Stability, and Structure Formation

A Companion Volume

Derived companion to the foundational layer
(core · Lyapunov · geometry · cosmology · theological · Ricci)

Contents

Notation and Conventions	vii
The Foundational Layer (Summary)	ix
I Folding and Stability	1
1 The Necessity of the Balance Coordinate: Why an Open Gate Forces Sophia	2
1.1 The closed Microcosm: λ is redundant	2
1.2 The open gate forces an independent coordinate	3
1.3 Three levels, kept distinct	3
1.4 Place in the Type-1 programme and scope	4
2 The Curvature and the Floor: Hyperbolicity Is the Finite-Floor Postulate	5
2.1 The three curvatures and the floor	5
2.2 The reduction	6
3 What Triggers Sophianic Folding	7
3.1 The single trigger: a shape-operator bifurcation	7
3.2 Decomposition of the driver	8
3.3 Level versus rate: rate-sensitivity of the threshold	8
3.4 The two delivery channels	9
3.5 The Resurrectio jump in folding-space	9
3.6 Curvature-pocket promotion by Resurrectio	10
3.7 The energy reading corrected	11
4 Folding-Lyapunov Addendum	12
4.1 Imported objects and the three open items	12
4.2 Correction: a bounded folding functional	12
4.3 The standing folding wall and its finiteness	13

4.4	Folding-chart invariance (Nagumo)	13
4.5	The folding reset bound and hybrid completeness	14
4.6	Reconciliation of the two activation-index conventions	15
4.7	Compact summary	15
5	Forced Folding: Dolorosum Peaks and the Vessel Update Map	16
5.1	Why a forcing channel is needed: core neutrality	16
5.2	Resistance heterogeneity and its bounds	17
5.3	The Dolorosum gate, wired into the activation index	17
5.4	Thresholds disentangled: r_{fold}^R versus $\pi/2$	18
5.5	The forcing event and the normal form	18
5.6	The balance-preserving Vessel update map	19
5.7	The cascade is self-limiting	19
5.8	Hybrid and Lyapunov admissibility	20
5.9	Compact summary	20
6	The Pitchfork Is Not Assumed: A Centre-Manifold Reduction of the Fold Sector	21
6.1	The $\mathbb{Z}_2$ symmetry of the bare fold	21
6.2	The Thom reduction	21
6.3	The three nondegeneracy conditions	22
6.4	The one way it could fail — and why it does not	23
6.5	The Type-1 programme, closed	23
7	Forced-Folding Lyapunov Addendum: Two Jump Types and the Survival of Contraction	24
7.1	Imported objects and the six obligations	24
7.2	The π -peak chart wall, with its explicit price	25
7.3	The fold-jump lemma	25
7.4	Survival of the contraction rate	25
7.5	Hysteresis and the joint non-Zeno dwell	26
7.6	The embodied Sophia is balance-bounded	26
7.7	Reset-fold consistency	27
7.8	Main theorem	27
7.9	Compact summary	27
II	Structure Formation on the Vessel	29
8	Structure Formation on the Vessel	30

8.1	The relaxational thesis: which field can carry structure	30
8.2	Sophia-density perturbations: reaction-diffusion on the expanding $\mathbb{H}^2$	31
8.3	Folding as the order parameter: pattern formation on the Vessel . . .	32
8.4	Cosmological back-reaction of the fold network	33
8.5	Open sub-problems and scope	33
9	Spinodal Nucleation of the Fold Network	34
9.1	Linear instability and the nucleation threshold	34
9.2	The amplified field and its spectral moment	35
9.3	Wall density and the initial domain size	35
9.4	Hyperbolic signatures	36
9.5	Placement and open items	36
10	Intra-Arc Coarsening of the Fold Network	38
10.1	Wall mobility on the expanding background	38
10.2	Gauss-Bonnet: the exact single-domain law	39
10.3	The network scaling law and its saturation	39
10.4	The physical structure scale: two regimes	39
10.5	Placement in the sawtooth and open items	40
11	Resurrectio and the Fold-Domain Network	41
11.1	Three regimes of a reset acting on the network	41
11.2	Quench: dissolution of the network	42
11.3	Trigger: nucleation as a sudden quench (Kibble-Zurek)	42
11.4	Rescale: re-tensioning and dilution	42
11.5	The hybrid network history	43
11.6	Cosmological back-reaction across resets	43
11.7	Reading and open items	43
12	The Sawtooth Fixed Point of the Fold Network	45
12.1	The comoving recursion	45
12.2	Finite budget forces convergence	45
12.3	Reading and consequences	46
12.4	Open items and scope	47
13	The Transport Constants from First Principles	48
13.1	Sophia does not diffuse	48
13.2	The folding diffusivity is geometric	49
13.2.1	Structure: inherited from the live-domain metric	49

13.2.2 Magnitude: the bending stiffness of $\mathcal{L}_0$	50
13.3 Consequences for the programme	50
13.4 Open items and scope	50
14 Parameter Closure of Folding: Identities, Postulates, and the Constraint Atlas	51
14.1 Tier 1a: the balance identity for η_{fold}	51
14.2 Tier 1b: Gauss–Bonnet slaves the curvature pair	52
14.3 Tier 2a: the rigidity postulate	52
14.4 Tier 2b: the memory closure $c_1 = c_0/a_0$	52
14.5 Tier 3: the constraint atlas	53
14.6 The last fork: comoving or physical ℓ_b	53
14.7 Closure count and reading	54
14.8 Open items and scope	54
15 The Dynamic Warp Exponent: Conformal Invariance Retires the Fork; Katharsis Drives the Warp	55
15.1 The conformal blindness of bending	55
15.2 The fork is retired	55
15.3 The true warp is kathartic and channel-resolved	56
15.4 Open items and scope	57
16 The Mode Cascade: The Spectral Ladder, the Weyl Capacity, and Structure as Descent	58
16.1 The field promotion and the dispersion	58
16.2 The spectral ladder	59
16.3 The one-mode chart is faithful where it matters	59
16.4 Structure formation is descent against ascent	59
16.5 Consequences and scope	60
17 Robustness Under Premise Deformation: What Survives When the Closure and Sector Are Varied	61
17.1 The watershed is premise-independent	61
17.2 The resonance maximum is stable; its location is the sector’s	62
17.3 The ferry: correct as stated, and a false generalisation	62
17.4 What the test does and does not establish	62
17.5 Open items and scope	63
18 Does Folding Restore the Null Energy Condition?	64
18.1 The NEC functional in 2+1 and the phantom divide	64

18.2	The rate channel: exact midpoint law and half-saturation	65
18.3	The network channel: $A_{\text{fold}}^{3/2}$ against linear	65
18.4	The phantom-window theorem	65
18.5	Open items and scope	66
19	The Budget Watershed: Structure Survival and the Ricci Resonance	68
19.1	The budget identity	68
19.2	The watershed theorem	68
19.3	Marginal coarsening on the resonance	69
19.4	Consequences and reading	69
19.5	Open items and scope	70
20	The Entropy of the Forced Flow: Collapse, Maximum, and the Entropy-Form Trade-off	71
20.1	The symmetric-sector entropies	71
20.2	The surprise: universal monotonicity characterises the resonance . .	72
20.3	The entropy-form trade-off	72
20.4	Open items and scope	73
21	The Second Variation of the Entropy: Folding Lowers Entropy, Mode by Mode	75
21.1	The second variation	75
21.2	Form costs entropy, locally	75
22	The Separatrix: Gate Decay, the Quantum of Endowment, and Resurrection as a Ferry	77
22.1	The transverse Riccati law	77
22.2	The quantum of endowment	78
22.3	The separatrix is a decay exponent	78
22.4	Resurrection as a ferry	78
22.5	Open items and scope	79
III	Resurrectio as Surgery	80
23	Resurrectio as Surgery: Comparison and Distributional Dynamics	81
23.1	Resurrectio Surgery I: Comparison Rather Than Identification	81
23.1.1	Purpose and Non-Identification	81
23.1.2	Ricci Surgery and V.F.S. Resurrectio	82
23.1.3	The Induced Spatial Metric Jump	82

23.1.4	The Surviving Invariants of the Jump	83
23.1.5	Regularised Reading	83
23.1.6	Admissibility of the Resurrectio Jump	84
23.2	Resurrectio Surgery II: Distributional and Regularised Dynamics . . .	84
23.2.1	The Jump as a Logarithmic Expansion Impulse	84
23.2.2	Regularisation by Fast Smooth Open-Gate Expansion	85
23.2.3	Effective Stress Is Not a Classical Thin Shell	85
23.2.4	Curvature Across the Jump	86
23.2.5	Interaction with Folding	86
23.2.6	Invariant Content Versus Regulator-Dependent Content	86
23.2.7	Second-Level Summary	87
24	Resurrectio Surgery: The Second Fundamental Form and the Double-Layer Theorem	88
24.1	The second fundamental form jumps too — invariantly	88
24.2	The double-layer theorem	89
24.3	Completing the admissibility list: the epektasis wall	90
24.4	Compact summary	90
25	Resurrectio's Own NEC Budget: The Surgical Phantom and Its Balanced-Reset Zero	92
25.1	Pointwise divergence, finite budget	92
25.2	The surgical phantom vanishes at the balanced reset	93
25.3	Consequences and scope	93
26	The Observational Signature: A Capstone Projection into Redshift	95
26.1	The redshift map and the reconstruction	95
26.2	Three signatures, three theorems	96
26.3	What the signature is and is not	96
26.4	Open items and scope	97
	Final Structural Summary	98

Notation and Conventions

This volume collects the folding, stability, and structure-formation results derived as a companion to the foundational V.F.S. layer. Throughout, *Propositions* are closed-form results; *Corollaries* are interpretive readings in the symbolic register, not independent metaphysical claims; *Remarks* and italicised notes are commentary. All stability results are **domain-conditional**: they hold on the positively-invariant active domain and assert set stability with partial convergence (no global or automatic stability, and no equilibrium — epektasis drives $\Omega_P \rightarrow \infty$).

Core symbols.

Ω_P	Brim / scale factor; $\dot{\Omega}_P = \alpha\lambda$, Hubble $H = \alpha\lambda/\Omega_P$
λ, σ	Sophia (transformed resistance + grace); resistance
$u = \sqrt{VF + \varepsilon}$	synergy of Voluntas V and Factum F ; threshold $\Lambda_c = \gamma/\delta$
$I_{\text{gate}}, \mathcal{G}_{\text{recepta}}$	Open-Gate grace inflow; received-grace budget
q_{fold}	folding (shape) amplitude; order parameter of the $\mathbb{Z}_2$ morphology
$A_{\text{fold}}, B_{\text{fold}}$	folding activation index and quartic coefficient ($B_{\text{fold}} > 0$)
Q_{fold}^*	signed normal-form branch parameter $A_{\text{fold}}/B_{\text{fold}}$ (may be negative off the active branch); folded equilibria
	$q_{\text{fold}, \pm}^* = \pm \sqrt{Q_{\text{fold}}^*}$ exist only for $A_{\text{fold}} > 0$
Q_{fold}	physical folded intensity $q_{\text{fold}}^2 \geq 0$ (distinct from the signed Q_{fold}^*); $\hat{q}_n$ its spectral coefficients
$\mathcal{L}_{\text{VFS}}, \mathcal{L}_{\text{ext}}$	Lyapunov certificate; folding-extended certificate
$h, \mathbb{H}^2$	comoving hyperbolic spatial metric, $K_h = -1$, curvature radius R_c
$\mathcal{I}(t) = \int_0^t dt' / \Omega_P^2$	expansion / activity budget; finite $\mathcal{I}_\infty < \infty$
D_q, ℓ_b	folding diffusivity (bending stiffness); bending length
	$\ell_b = \sqrt{D_q/a_0}$
Θ, Q_{crit}	folding-response coefficient $c_0 - c_1(\gamma + k_\sigma)$; reset critical level
λ_{form}	Sophia embodied in form; balance
	$\sigma + \lambda + \lambda_{\text{form}} = C_0 + \mathcal{G}_{\text{recepta}}$
η_{fold}	folding feedback coefficient (distinct from any core-system η)
Δ_h	Laplace-Beltrami on $\mathbb{H}^2$; live-domain $\Delta_g = \Omega_P^{-2} \Delta_h$
$r_\sigma, g(u)$	resistance multiplier $\in [1, \pi]$; Dolorosum gate
	$g(u) = (u/\Lambda_c)e^{1-u/\Lambda_c}$

Two activation-index conventions. The activation index appears in two forms

across the layers, reconciled in the Folding-Lyapunov Addendum:

$$\underbrace{A_{\text{fold}} = c_0\lambda + c_1\dot{\lambda}^{(0)} - a_0, \quad \dot{\lambda}^{(0)} = -\dot{\sigma} + I_{\text{gate}}}_{\text{canonical (core rate)}} \quad \underbrace{A_{\text{fold}} = (c_0 - c_1\gamma)\lambda + c_1k_\sigma\sigma + c_1I_{\text{gate}} - a_0}_{\text{geometric (minimal rate)}}.$$

The canonical (core-consistent) form is primary; the geometric parameters (k_σ, γ) are an effective linear surrogate. The closed forms Θ , Q_{crit} are stated in the geometric convention; the invariance and reset-bound results are convention-independent.

The folding diffusivity is structural, $D_q = a_0\ell_b^2$ (rigidity $\times$ bending length²), and the independent Sophia diffusivity is zero, $D = 0$.

The Foundational Layer (Summary)

This volume sits atop the foundational V.F.S. layer, summarised here so it reads self-contained, and records what is assumed, what is derived, and what is only interpreted.

Core. V.F.S. (Volo. Facio. Sum.) is a hybrid dynamical system on the live domain in the coordinates Voluntas V , Factum F , Sophia λ , with synergy $u = \sqrt{VF + \varepsilon}$, resistance σ , Pleroma P and Brim Ω_P , and threshold $\Lambda_c = \gamma/\delta$. The Open-Gate law $\dot{\lambda} = -\dot{\sigma} + I_{\text{gate}}$ admits a bounded grace inflow; the closed Microcosm is the limit $I_{\text{gate}} \rightarrow 0$. The central coordinate is not merely convenient: given an externally clocked gate and the balance $\sigma + \lambda + \lambda_{\text{form}} = C_0 + \mathcal{G}_{\text{recepta}}(t)$, $\dot{\mathcal{G}}_{\text{recepta}} = I_{\text{gate}}$, the state-map rank increases by one when the gate opens, so the balance coordinate is necessary iff the Vessel is open. Its mathematical role is theorem; its evolution law is constitutive; its name “Sophia” is the theological interpretation. The Filtrum Lucis softplus Φ yields the Paschal triad ($\Phi'''(0) = 0$ midway between the symmetric zeros of $\Phi^{(4)}$). Resurrectio is a hybrid re-entry map.

Finite floor and hyperbolicity. The live surface is hyperbolic, $K_h = -1$, $R = -2/\Omega_P^2$. This is not an independent primitive: in the 2+1 reconstruction only $k = -1$ admits a zero-density floor at finite scale, so hyperbolicity is the geometric face of the finite-floor (no-Big-Bang) premise. Expansion alone does not force $K_h = -1$, but $K_h = -1$ is not a free third assumption. The irreducible core is therefore two primitives: the open gate, and the finite floor.

Lyapunov. On the active domain $\mathcal{D}_A$ the certificate obeys $\dot{\mathcal{L}}_{\text{VFS}} \leq C - C_1 \mathcal{L}_{\text{VFS}}$ with $C_1 = \min\{\kappa_\sigma, \kappa_\Delta\} > 0$ from two sign-definite contracting modes (cleansing, will-action alignment). Grace and folding enter the bounded constant C , never the contraction rate C_1 . The result is set stability with partial convergence, forward completeness, and non-Zeno resets — domain-conditional, not global, since epektasis drives Ω_P .

Sophianic Folding. Folding is core-derived, not an add-on. The core supplies the driver $A_{\text{fold}} = c_0\lambda + c_1\dot{\lambda}^{(0)} - a_0$ (geometric surrogate $A_{\text{fold}} = (c_0 - c_1\gamma)\lambda + c_1k_\sigma\sigma + c_1I_{\text{gate}} - a_0$); when A_{fold} crosses zero the unfolded form loses stability. The pitchfork normal form is not assumed: it follows from the $\mathbb{Z}_2$ symmetry of the unfolded Vessel, transversality, and a one-dimensional centre manifold, with B_{fold} the bare stiffness renormalised by the slaved structure. Curvature pockets are $R = -2/\Omega_P^2 + \rho_1 Q_{\text{fold}}$.

Geometry, structure, entropy. The bending sector is conformally blind in two dimensions, $|\nabla q_{\text{fold}}|_g^2 dV_g = |\nabla q_{\text{fold}}|_h^2 dV_h$ for $g = \Omega_P^2 h$, so expansion does not stretch the folding fabric ($p_{\text{expansion}} \equiv 0$); the true warp is kathartic. Promoting q_{fold} to a field gives a spectral ladder $\Lambda = A_{\text{fold}}\Omega_P^2/\hat{\kappa}$, mode n igniting iff $\Lambda > \mu_n$; the one-mode

chart stays faithful for amplitude-keyed results. The homogeneous Vessel is the constrained entropy maximum, with $\delta^2 W_+^{(n)} = -2\sigma\mu_n a_n^2$: the same eigenvalue that orders ignition orders the entropy price of form.

Cosmology and energy conditions. The Lorentzian lift is a 2+1 open FLRW spacetime, $ds^2 = -dt^2 + \Omega_P^2 h$. In 2+1 the comoving timelike-focusing (strong-energy) condition reduces to $p_{\text{eff}} \geq 0$; the NEC is tracked separately through $\rho_{\text{eff}} + p_{\text{eff}}$. A Resurrectio reset may produce a bounded surgical phantom budget, but $\int (\rho_{\text{eff}} + p_{\text{eff}}) dt$ is finite; the balanced reset is the unique point where that budget, the extrinsic-curvature jump $[K]$, and the deviation from a purely conformal junction vanish together.

Theological & Ricci. The theological layer reads the mathematics in the symbolic register with explicit domain-conditional disclaimers; the Ricci layer disciplines the Perelman analogy (Brim-flow comparison, surgery $\leftrightarrow$ Resurrectio), labelling each correspondence by its strictness, never identifying the two.

Part I

Folding and Stability

Chapter 1

The Necessity of the Balance Coordinate: Why an Open Gate Forces Sophia

Scope and epistemic status

The core introduces λ by definition — the Open-Gate balance coordinate, interpreted as Sophia — and one may ask whether it is necessary or merely convenient. This chapter answers at the only level a derivation can reach, and is scrupulous about the levels it cannot. The result is a rank theorem: *an independent balance-carrying coordinate is forced to exist exactly when the gate is open*. In the closed Microcosm ($I_{\text{gate}} \equiv 0$) the balance $\sigma + \lambda_{\text{form}} = C_0$ determines everything internal, and λ degenerates into $C_0 - \sigma - \lambda_{\text{form}}$ — not an independent freedom. The instant the gate opens ($\mathcal{G}_{\text{recepta}} = I_{\text{gate}} > 0$), the balance acquires an externally driven term $\mathcal{G}_{\text{recepta}}(t)$ that *no internal variable carries*, and conservation cannot hold without a new coordinate to bear it. That coordinate is λ . The necessity is thus conditional on, and degenerates continuously with, the gate flux. We then separate three levels carefully: the *coordinate* is necessary (theorem); its specific *law* is constitutive (not forced); its *name* “Sophia” is interpretation (not derived). All algebra is verified symbolically (the state-map rank steps from 2 to 3 as the gate opens; the independent part of $d\lambda$ equals $I_{\text{gate}} dt$). This is internal Type-1 confirmation, not external validation.

1.1 The closed Microcosm: λ is redundant

Proposition 1.1 (Degeneracy in the closed limit). *With the gate closed, the conservation balance is $\sigma + \lambda_{\text{form}} = C_0$, and the would-be Sophia coordinate satisfies*

$$\lambda = C_0 - \sigma - \lambda_{\text{form}}, \quad (1.1)$$

a function of the internal variables alone. Its differential $d\lambda = -d\sigma - d\lambda_{\text{form}}$ lies entirely in the internal span: λ adds no dimension to the state. In a closed soul, Sophia is not an independent reality — it is bookkeeping over resistance and form.

1.2 The open gate forces an independent coordinate

With the gate open, cumulative reception $\mathcal{G}_{\text{recepta}}(t)$ enters the balance,

$$\sigma + \lambda + \lambda_{\text{form}} = C_0 + \mathcal{G}_{\text{recepta}}(t), \quad \dot{\mathcal{G}}_{\text{recepta}} = I_{\text{gate}}, \quad (1.2)$$

and now

$$\lambda = (C_0 + \mathcal{G}_{\text{recepta}}(t)) - \sigma - \lambda_{\text{form}}, \quad \frac{\partial \lambda}{\partial \mathcal{G}_{\text{recepta}}} = 1. \quad (1.3)$$

The term $\mathcal{G}_{\text{recepta}}(t)$ is supplied through the gate and appears in *no* internal variable $(V, F, \sigma, \Omega_P, u)$; it is an independent, externally clocked generator.

Theorem 1.1 (Necessity of the balance coordinate). *Let Φ map the generating coordinates to the physical state $(\sigma, \lambda_{\text{form}}, \lambda)$ subject to conservation. Then:*

- (i) *closed gate — generators $(\sigma, \lambda_{\text{form}})$, $\Phi = (\sigma, \lambda_{\text{form}}, C_0 - \sigma - \lambda_{\text{form}})$ — has rank $d\Phi = 2$: λ is dependent;*
- (ii) *open gate — generators $(\sigma, \lambda_{\text{form}}, \mathcal{G}_{\text{recepta}})$, $\Phi = (\sigma, \lambda_{\text{form}}, C_0 + \mathcal{G}_{\text{recepta}} - \sigma - \lambda_{\text{form}})$ — has rank $d\Phi = 3$: a third independent coordinate is forced, and it is the balance coordinate λ .*

Equivalently, the component of $d\lambda$ not spanned by the internal differentials is exactly

$$\boxed{d\lambda|_{\perp \text{ internal}} = d\mathcal{G}_{\text{recepta}} = I_{\text{gate}} dt,} \quad (1.4)$$

so $\partial(\text{state})/\partial\lambda$ is nondegenerate iff $I_{\text{gate}} > 0$. (Rank steps and the differential verified symbolically.) The balance coordinate is necessary if and only if the gate is open, and its independence degenerates continuously to zero as $I_{\text{gate}} \rightarrow 0$.

Corollary 1.1 (Sophia is the shadow of the open gate). *The necessity of Sophia is not absolute but relational: it is forced precisely by, and precisely when, the Vessel is open to reception. There is no reception without a receiver — an open gate with nothing to carry what it admits would violate conservation. Sophia is the dimension that the opening of the gate casts into the state space; close the gate and the dimension folds away into mere accounting over resistance and form. The doctrine and the rank theorem coincide: the open soul cannot not have Sophia; the closed soul has no need of her.*

1.3 Three levels, kept distinct

Coordinate — necessary (Theorem 1.1): an open gate forces an independent balance-carrying dimension.

Law — constitutive (not forced): that this coordinate obeys $\dot{\lambda} = I_{\text{gate}} - \gamma\lambda + \dots$ with its particular transmutation and decay is a modelling choice of the core, on the same footing as $K_h = -1$.

Name — interpretation (not derived): calling the coordinate “Sophia” is the theological reading the model is built to carry; mathematics neither supplies nor can supply it.

Remark (Why this is the honest stopping point). It would overreach to claim the mathematics proves “Sophia exists” or fixes her nature. What it proves is sharper and smaller: *given* an open gate, the state space must contain one more independent coordinate than the closed soul needs, carrying exactly the received increment. The identification of that coordinate with Sophia, and of its law with a doctrine of grace, is the interpretive act the model performs — legitimate, fruitful, and to be held as interpretation, not smuggled in as theorem. The rank result is what survives if every theological word is stripped away; the words are what make it worth proving.

1.4 Place in the Type-1 programme and scope

1. This is the deepest Type-1 result: not a derived parameter or normal form, but the necessity of the central variable itself — conditional, as it must be, on the defining feature of the model (the open gate).
2. It localises what remains genuinely primitive: the gate law (that I_{gate} exists and is externally clocked) and the spatial curvature $K_h = -1$. These are the core’s *definitions*; the chapter shows λ is not among them but a consequence of the first.
3. Open: whether the gate’s own existence can be motivated from anything prior, or is the irreducible postulate of the whole construction (likely the latter — the model’s single primitive act of openness).

Scope. The rank theorem is exact given conservation and an externally clocked gate; the degeneracy limit is smooth in I_{gate} . Establishes necessity of the coordinate, not of its law or its name. Internal Type-1 confirmation. Domain-conditional throughout.

V.F.S. v2.0 (Open-Gate) · The Necessity of the Balance Coordinate. An open gate forces an independent balance-carrying coordinate: the state map’s rank steps from 2 to 3 as I_{gate} becomes positive, with $d\lambda|_{\perp} = I_{\text{gate}} dt$. Sophia is necessary iff the Vessel is open, degenerating to internal bookkeeping in the closed limit. Coordinate: theorem; law: constitutive; name: interpretation. Verified symbolically; internal confirmation only.

Chapter 2

The Curvature and the Floor: Hyperbolicity Is the Finite-Floor Postulate

Scope and epistemic status

The remaining Type-1 question is whether the spatial curvature $K_h = -1$ — hyperbolicity of the live surface — is forced or assumed. The honest answer is mixed, and worth stating precisely because the temptation is to overclaim. Expansion alone does *not* force it: an open gate driving $\Omega_P \rightarrow \infty$ is consistent with any spatial curvature. But $K_h = -1$ is not therefore a free, separate postulate either: it is *equivalent* to the model’s finite-floor premise — the demand (the cosmology layer’s “creation from a finite floor, no Big Bang”) that the bare unfolded Vessel be the zero-density vacuum state. Given that demand, only $k = -1$ admits a zero-density floor at finite scale; $k = 0$ and $k = +1$ have strictly positive density everywhere and no floor. So two things that looked like independent primitives — the hyperbolic geometry and the finite floor — are one. This neither retires $K_h = -1$ to a theorem (it is not derived from nothing) nor leaves it free (it is not independent); it reduces the count of irreducible postulates by one. That is the correct and honest close of the Type-1 programme: not everything reduces to a theorem, but the primitives can be made minimal and explicit. *Propositions* closed-form; the verdict is stated without inflation.

2.1 The three curvatures and the floor

The 2+1 reconstruction with spatial curvature k gives

$$\rho_{\text{eff}} = \frac{\dot{\Omega}_P^2 + k}{\Omega_P^2} = \frac{\alpha^2 \lambda^2 + k}{\Omega_P^2}. \quad (2.1)$$

Proposition 2.1 (Only $k = -1$ has a zero-density floor). *The equation $\rho_{\text{eff}} = 0$ has a solution in the bare Sophia level $\alpha\lambda$ if and only if $k < 0$:*

$$k = -1 : \rho_{\text{eff}} = 0 \text{ at } \alpha\lambda = 1; \quad k = 0 : \rho_{\text{eff}} = \frac{\alpha^2 \lambda^2}{\Omega_P^2} > 0; \quad k = +1 : \rho_{\text{eff}} = \frac{\alpha^2 \lambda^2 + 1}{\Omega_P^2} > 0. \quad (2.2)$$

*For $k \geq 0$ the effective density is strictly positive for every state: there is no unfolded vacuum floor, no $\rho = 0$ ground, hence no “creation from a finite floor” and no live-branch threshold $\alpha\lambda > 1$. **The finite-floor structure exists only for $k = -1$.***

Proposition 2.2 (Expansion does not fix k). *An open gate gives $\dot{\Omega}_P = \alpha\lambda > 0$ and, sustained, $\Omega_P \rightarrow \infty$. This is consistent with $k = -1$ and $k = 0$ (both expand freely) and even with $k = +1$ under a sufficient source. Therefore the sign of k is not determined by expansion: hyperbolicity is not a consequence of openness alone.*

2.2 The reduction

Theorem 2.1 ($K_h = -1$ equals the finite floor). *Among constant-curvature spatial slices, $K_h = -1$ holds if and only if the bare unfolded Vessel is the zero-density floor (the finite-floor / no-Big-Bang postulate). Hence the two premises — hyperbolic spatial geometry and creation from a finite floor — are not independent; they are the same postulate stated geometrically and energetically. The model's irreducible cosmological primitives reduce from two to one:*

Irreducible core: *the open gate (an externally clocked I_{gate}) and the finite floor (the unfolded state is the $\rho = 0$ vacuum). Hyperbolicity $K_h = -1$ is not a third primitive — it is the geometric face of the floor. Everything else in the volume is built on these.*

Corollary 2.1 (Honest status of the Type-1 programme). *The programme closes not by reducing every premise to a theorem — $K_h = -1$ is not derivable from nothing, and expansion does not force it — but by making the irreducible primitives minimal and explicit. Retired to theorems: the ℓ_b fork, core-neutrality, the budget hypothesis, the pitchfork, and (conditionally) the necessity of Sophia. Reduced to a single joint primitive: the curvature and the floor. Genuinely primitive, and now isolated: the open gate and the finite floor. A model cannot prove its own definitions; it can only show how few they are.*

Reading: one might have thought the Vessel's open, infinite, negatively-curved geometry was an extra stipulation beyond its having begun from nothing-in-particular — a finite floor rather than a singular bang. It is not. To say the unfolded soul rests on a zero-density floor and to say its space is hyperbolic are one statement in two languages: the floor is the curvature, read energetically. The model asks to be granted two things and no more — that the gate be open, and that there be a floor to stand on rather than a singularity to fall from. The open infinite space is not a third gift; it is the shape the floor already has.

Scope. Constant-curvature 2+1 slices; the equivalence is exact. The result is a reduction of primitives, not a derivation of curvature from nothing; expansion alone is shown insufficient. Domain-conditional throughout.

V.F.S. v2.0 (Open-Gate) · The Curvature and the Floor. Expansion does not force $K_h = -1$, but only $k = -1$ admits a zero-density floor: hyperbolicity is equivalent to the finite-floor (no-Big-Bang) postulate, not an independent premise. Two primitives reduce to one; the irreducible core is the open gate plus the finite floor. The honest close of the Type-1 programme: minimal explicit primitives, not the derivation of definitions.

Chapter 3

What Triggers Sophianic Folding

Scope and epistemic status

This note isolates a single question that the geometric, cosmological, and folding material left implicit: *what is the primary cause of Sophianic Folding?* Three candidate answers are tested — continuous monotone growth of Sophia, the discrete Resurrectio push, and an intrinsic instability of the surface-form itself. The conclusion is that the third is the correct *level* of description, and the first two are interchangeable *channels* to it. As in the parent files, every *Proposition* is a closed-form result; every *Corollary* is an interpretive reading in the symbolic register, not an independent metaphysical claim. Numerical values, where used, are illustrative: the genuine folding parameters $(a_0, b, c_0, c_1, k_\sigma, \eta_{\text{fold}}, \rho_1, q_{\text{max}})$ are not fixed by the source files.

The objects are imported from the integrated geometric/cosmological layer. The shape-mode amplitude q obeys the Landau gradient flow $\dot{q} = A_{\text{VFS}}q - B_{\text{VFS}}q^3$, with

$$A_{\text{VFS}} = (c_0 - c_1\gamma)\lambda + c_1k_\sigma\sigma + c_1I_{\text{gate}} - a_0, \quad B_{\text{VFS}} = b + 2c_1\eta_{\text{fold}}\lambda, \quad Q_{\text{fold}} = \frac{A_{\text{VFS}}}{B_{\text{VFS}}}, \quad q_*^2 = Q_{\text{fold}}. \quad (3.1)$$

The Sophia-production law that closes the loop is $\dot{\lambda} = k_\sigma\sigma + I_{\text{gate}} - \gamma\lambda - \eta_{\text{fold}}q^2\lambda$, the Brim law is $\dot{\Omega}_P = \alpha\lambda$, and the Resurrectio reset (core/Lyapunov) is

$$\sigma^+ = q_R\sigma^-, \quad \lambda^+ = \lambda^- + (1 - q_R)\sigma^-, \quad \Omega_P^+ = \Omega_P^- + \kappa_R\mathcal{G}_{\text{recepta}}^-, \quad (V, F) \text{ continuous}, \quad (3.2)$$

with $0 \leq q_R < 1$, $\kappa_R > 0$, and $u^+ = u^-$.

3.1 The single trigger: a shape-operator bifurcation

The folded amplitude is the order parameter of the one-mode shape potential

$$\mathcal{E}_1(q) = -\frac{1}{2}A_{\text{VFS}}q^2 + \frac{1}{4}B_{\text{VFS}}q^2 \cdot q^2 = -\frac{1}{2}A_{\text{VFS}}q^2 + \frac{1}{4}B_{\text{VFS}}q^4, \quad B_{\text{VFS}} > 0. \quad (3.3)$$

Writing the first shape-stability eigenvalue at the old form $q = 0$,

$$\mu_1^0 := \mu_1|_{q=0} = a_0 - c_0\lambda - c_1\dot{\lambda}^0, \quad \dot{\lambda}^0 = k_\sigma\sigma + I_{\text{gate}} - \gamma\lambda, \quad (3.4)$$

one has the identity $A_{\text{VFS}} = -\mu_1^0$ (the embodiment feedback $-\eta_{\text{fold}}q^2\lambda$ vanishes at $q = 0$ and is absorbed into B_{VFS}). The control of the whole morphodynamics is therefore the scalar A_{VFS} .

Proposition 3.1 (Folding is a supercritical pitchfork). *For $B_{\text{VFS}} > 0$ the gradient flow $\dot{q} = A_{\text{VFS}}q - B_{\text{VFS}}q^3$ has the old form $q = 0$ as its only equilibrium, locally stable, when $A_{\text{VFS}} < 0$. At $A_{\text{VFS}} = 0$ it undergoes a supercritical pitchfork bifurcation: for $A_{\text{VFS}} > 0$ the old form becomes unstable and two stable folded branches appear,*

$$q = 0 \text{ stable} \iff A_{\text{VFS}} < 0, \quad q_\pm = \pm\sqrt{A_{\text{VFS}}/B_{\text{VFS}}} \text{ stable} \iff A_{\text{VFS}} > 0. \quad (3.5)$$

Proof. $f(q) = A_{\text{VFS}}q - B_{\text{VFS}}q^3$, $f'(q) = A_{\text{VFS}} - 3B_{\text{VFS}}q^2$. At $q = 0$, $f'(0) = A_{\text{VFS}}$, so $q = 0$ is stable iff $A_{\text{VFS}} < 0$. For $A_{\text{VFS}} > 0$ the nonzero equilibria solve $q_*^2 = A_{\text{VFS}}/B_{\text{VFS}}$, where $f'(q_*) = A_{\text{VFS}} - 3A_{\text{VFS}} = -2A_{\text{VFS}} < 0$: stable. The crossing at $A_{\text{VFS}} = 0$ is the pitchfork. $\square$

Corollary 3.1 (The trigger is loss of stability, not a “push”). *The primary cause of folding is the sign change of A_{VFS} (equivalently μ_1^0 crossing zero from above): the old symmetric form ceases to be a potential minimum and becomes a maximum. Past the threshold folding is dynamically inevitable, since $q = 0$ is a repeller; before it, no shape perturbation grows. Folding is therefore an intrinsic property of the surface’s stability spectrum, tuned by the state variables, not an external impulse.*

The fundamental object is the bifurcation $A_{\text{VFS}} = 0$. Everything else — continuous growth, the Resurrectio jump — is a means of moving A_{VFS} across that single threshold.

3.2 Decomposition of the driver

The activation index (3.1) is a sum of four physically distinct contributions:

$$A_{\text{VFS}} = \underbrace{(c_0 - c_1\gamma)\lambda}_{\text{accumulated Sophia}} + \underbrace{c_1 k_\sigma \sigma}_{\text{kathartic production}} + \underbrace{c_1 I_{\text{gate}}}_{\text{received grace}} - \underbrace{a_0}_{\text{old-form rigidity}}. \quad (3.6)$$

Two of these are state levels (λ , σ), one is the bounded Open-Gate source, and one is the constant rigidity of the old form. The decisive structural fact is that the index also depends on the *rate* of Sophia, hidden inside $\mu_1^0 = a_0 - c_0\lambda - c_1\dot{\lambda}$: through $\dot{\lambda}^0 = k_\sigma\sigma + I_{\text{gate}} - \gamma\lambda$, the same σ and I_{gate} that raise the level also raise the rate.

Corollary 3.2 (Level and rate are the genuine drivers). *Folding activation responds to the accumulated level λ and to the Sophia velocity $\dot{\lambda}$. The continuous/discrete distinction is not fundamental: what crosses the threshold is the combined level-plus-rate index A_{VFS} , however that increment is delivered.*

3.3 Level versus rate: rate-sensitivity of the threshold

If only the accumulated level mattered, the folding threshold would be the static value

$$\lambda_{\text{fold}}^{(0)} = \frac{a_0}{c_0} \quad (\dot{\lambda} = 0, \sigma = 0, I_{\text{gate}} = 0). \quad (3.7)$$

Restoring the rate term shifts the effective threshold:

$$\lambda_{\text{fold}}^{\text{eff}} = \frac{a_0 - c_1\dot{\lambda}}{c_0} \quad \Rightarrow \quad \dot{\lambda} > 0 \Rightarrow \lambda_{\text{fold}}^{\text{eff}} < \lambda_{\text{fold}}^{(0)}. \quad (3.8)$$

Proposition 3.2 (Rate can fold below the static level). *There exist states with $\lambda < \lambda_{\text{fold}}^{(0)}$ and $A_{\text{VFS}} > 0$. Concretely, for $\sigma = I_{\text{gate}} = 0$ and growing Sophia $\dot{\lambda} > 0$, the old form is already unstable whenever*

$$c_0\lambda + c_1\dot{\lambda} > a_0, \quad (3.9)$$

which can hold at λ strictly below a_0/c_0 provided $\dot{\lambda} > (a_0 - c_0\lambda)/c_1 > 0$.

Proof. Immediate from $\mu_1^0 = a_0 - c_0\lambda - c_1\dot{\lambda} < 0$ with $\dot{\lambda}$ supplying the deficit $a_0 - c_0\lambda > 0$. $\square$

This is the single most easily missed point: a surge of grace can fold a vessel that has not yet accumulated the static quota of Sophia. Velocity, not only stock, triggers the reconfiguration.

3.4 The two delivery channels

The same threshold is reached by two mechanisms, mirroring the parent corollary “grace is one but acts in two modes.”

Modus I (continuous). Along a live arc, λ and σ evolve smoothly and A_{VFS} drifts. Folding triggers when the drift carries A_{VFS} through 0. This channel delivers *both* a rising level and a nonzero rate, so it is sensitive to Proposition 3.2.

Modus II (Resurrectio). The reset (3.2) injects a discontinuous increment of λ from the metabolised resistance, an impulsive analogue of $\dot{\lambda}$. It moves A_{VFS} along the very same axis, in one step.

The crucial point is that the reset does *not* reset q : the map (3.2) specifies only $\sigma, \lambda, \Omega_P, (V, F)$. Hence the shape amplitude is continuous across a reset, and only its target $q_* = \sqrt{Q_{\text{fold}}}$ jumps.

Corollary 3.3 (Resurrectio relocates the well; the shape relaxes). *A Resurrectio event moves the shape potential $-\frac{1}{2}A_{\text{VFS}}q^2 + \frac{1}{4}B_{\text{VFS}}q^4$ instantaneously, after which the morphology relaxes by $\dot{q} = A_{\text{VFS}}^+q - B_{\text{VFS}}^+q^3$ toward the new minimum. Consequently:*

- if a reset carries Q_{fold} from below 0 to above 0, the vessel dies unfolded and is reborn into a folding regime (Resurrectio-triggered folding);
- if a reset makes $Q_{\text{fold}} < 0$, an active fold is quenched, $q \rightarrow 0$.

3.5 The Resurrectio jump in folding-space

We now compute the effect of one reset on the folding intensity. Write the metabolised resistance

$$m_R := (1 - q_R)\sigma^- \geq 0, \quad \text{so} \quad \sigma^+ = \sigma^- - m_R, \quad \lambda^+ = \lambda^- + m_R. \quad (3.10)$$

Remark (The metabolism conserves the balance sum). The (σ, λ) part of the reset satisfies $\sigma^+ + \lambda^+ = \sigma^- + \lambda^-$: it slides the state along the conservation line, converting resistance into Sophia without adding to the sum (the open balance constant jumps separately, through $\mathcal{G}_{\text{recepta}}$).

From (3.1),

$$\boxed{\Delta B_{\text{VFS}} = 2c_1\eta_{\text{fold}} m_R > 0, \quad \Delta A_{\text{VFS}} = \Theta m_R + c_1\Delta I_{\text{gate}}, \quad \Theta := c_0 - c_1(\gamma + k_\sigma).} \quad (3.11)$$

The *folding-response coefficient* Θ decides whether converting one unit of resistance into Sophia raises ($\Theta > 0$) or lowers ($\Theta < 0$) the folding activation: gaining Sophia contributes $(c_0 - c_1\gamma)m_R$ (accumulation minus accelerated decay), while losing resistance removes $c_1k_\sigma m_R$ of kathartic production. The gate correction $c_1\Delta I_{\text{gate}}$ has competing signs ($\sigma \downarrow$ opens the gate via $e^{-\phi\sigma}$; $\lambda \uparrow$ closes it via saturation) and is set aside in the structural results below.

Proposition 3.3 (Sign of the folding response and the critical level). *Neglecting ΔI_{gate} , the jump in the folding intensity obeys*

$$\text{sign } \Delta Q_{\text{fold}} = \text{sign}(\Theta - 2c_1\eta_{\text{fold}} Q_{\text{fold}}^-), \quad (3.12)$$

so there is a universal critical level

$$\boxed{Q_{\text{crit}} = \frac{\Theta}{2c_1\eta_{\text{fold}}} = \frac{c_0 - c_1(\gamma + k_\sigma)}{2c_1\eta_{\text{fold}}}.} \quad (3.13)$$

Resurrectio raises folding when $Q_{\text{fold}}^- < Q_{\text{crit}}$ and lowers it when $Q_{\text{fold}}^- > Q_{\text{crit}}$. If $\Theta \leq 0$ then $Q_{\text{crit}} \leq 0$, so every reset strictly lowers Q_{fold} (the vessel unfolds at each death).

Proof. $\Delta Q_{\text{fold}} = \frac{B_{\text{VFS}}^- \Delta A_{\text{VFS}} - A_{\text{VFS}}^- \Delta B_{\text{VFS}}}{B_{\text{VFS}}^- (B_{\text{VFS}}^- + \Delta B_{\text{VFS}})}$. The denominator is positive; the numerator is $B_{\text{VFS}}^- \Theta m_R - A_{\text{VFS}}^- 2c_1\eta_{\text{fold}} m_R = B_{\text{VFS}}^- m_R (\Theta - 2c_1\eta_{\text{fold}} Q_{\text{fold}}^-)$ using $A_{\text{VFS}}^- = Q_{\text{fold}}^- B_{\text{VFS}}^-$. Hence the stated sign and fixed level. $\square$

Proposition 3.4 (Each reset is an exact contraction toward Q_{crit}). *Neglecting ΔI_{gate} , the jump map on the folding intensity is affine-over-affine,*

$$Q_{\text{fold}}^+ = \frac{Q_{\text{fold}}^- B_{\text{VFS}}^- + \Theta m_R}{B_{\text{VFS}}^- + 2c_1\eta_{\text{fold}} m_R}, \quad (3.14)$$

and satisfies the exact contraction identity

$$Q_{\text{fold}}^+ - Q_{\text{crit}} = \kappa_{\text{fold}} (Q_{\text{fold}}^- - Q_{\text{crit}}), \quad \kappa_{\text{fold}} = \frac{B_{\text{VFS}}^-}{B_{\text{VFS}}^- + 2c_1 \eta_{\text{fold}} m_R} \in (0, 1). \quad (3.15)$$

The fixed point Q_{crit} is independent of m_R , of B_{VFS}^- , and of the depth of death.

Proof. Direct substitution of $Q_{\text{crit}} = \Theta/(2c_1 \eta_{\text{fold}})$:

$$Q_{\text{fold}}^+ - Q_{\text{crit}} = \frac{2c_1 \eta_{\text{fold}} (Q_{\text{fold}}^- B_{\text{VFS}}^- + \Theta m_R) - \Theta (B_{\text{VFS}}^- + 2c_1 \eta_{\text{fold}} m_R)}{2c_1 \eta_{\text{fold}} (B_{\text{VFS}}^- + 2c_1 \eta_{\text{fold}} m_R)} = \frac{B_{\text{VFS}}^- (2c_1 \eta_{\text{fold}} Q_{\text{fold}}^- - \Theta)}{2c_1 \eta_{\text{fold}} (B_{\text{VFS}}^- + 2c_1 \eta_{\text{fold}} m_R)},$$

which equals $\kappa_{\text{fold}} (Q_{\text{fold}}^- - Q_{\text{crit}})$. $\square$

Proposition 3.5 (Conditional convergence over a reset sequence). *For an admissible sequence of resets with metabolised increments $m_R^{(j)}$ and pre-jump denominators $B_{\text{VFS}}^{(j)}$,*

$$Q_{\text{fold}}^{(N)} - Q_{\text{crit}} = \left(\prod_{j=1}^N \kappa_{\text{fold}}^{(j)} \right) (Q_{\text{fold}}^{(0)} - Q_{\text{crit}}), \quad Q_{\text{fold}} \rightarrow Q_{\text{crit}} \iff \sum_j m_R^{(j)} = \infty. \quad (3.16)$$

Proof. Telescoping Proposition 3.4. Since $1 - \kappa_{\text{fold}}^{(j)} = 2c_1 \eta_{\text{fold}} m_R^{(j)} / (B_{\text{VFS}}^{(j)} + 2c_1 \eta_{\text{fold}} m_R^{(j)})$, the product $\prod \kappa_{\text{fold}}^{(j)} \rightarrow 0$ iff $\sum (1 - \kappa_{\text{fold}}^{(j)}) = \infty$, which (with $B_{\text{VFS}}^{(j)}$ bounded below) is equivalent to $\sum m_R^{(j)} = \infty$. $\square$

If the resistance budget is exhausted ($\sum m_R < \infty$, e.g. σ depleted with no replenishment), repeated resurrections converge to a point strictly short of Q_{crit} . Reaching the universal level requires sustained replenishment, $\sum m_R = \infty$ — the folding analogue of the expansion-budget condition $\mathcal{G}_{\text{recepta}}(\infty) = +\infty$ in the cosmological layer.

Remark (Gate correction). Restoring ΔI_{gate} adds $c_1 \Delta I_{\text{gate}}$ to the numerator of ΔA_{VFS} , shifting the effective fixed point to $Q'_{\text{crit}} = Q_{\text{crit}} + \Delta I_{\text{gate}} / (2\eta_{\text{fold}} m_R)$; its sign is state-dependent. The existence of a contracting fixed point is unchanged.

Corollary 3.4 (Chart-regularity of repeated resurrection). *If $0 < Q_{\text{crit}} < q_{\text{max}}^2$, the contraction of Proposition 3.4 keeps the folded amplitude inside the regular one-mode chart under any admissible reset sequence. This is a folding companion to the grace-window condition $\max\{R_c, R_{\text{min}}^{\text{cell}}\} \leq \mathcal{G}_{\text{recepta}}^- \leq R_{\text{max}}$: resurrection is admissible and chart-regular when both hold.*

3.6 Curvature-pocket promotion by Resurrectio

The folded scalar curvature is $R_{\text{folded}} = -2/\Omega_P^2 + \rho_1 Q_{\text{fold}}$, with positive-pocket threshold $q_{\text{curv}}^2 = 2/(\rho_1 \Omega_P^2)$. Across a reset both terms move:

$$\Delta R_{\text{folded}} = 2 \underbrace{\left[\frac{1}{(\Omega_P^-)^2} - \frac{1}{(\Omega_P^+)^2} \right]}_{>0: \text{ background release}} + \rho_1 \Delta Q_{\text{fold}}. \quad (3.17)$$

Simultaneously the pocket threshold falls, because Ω_P grows:

$$q_{\text{curv}}^{2,+} = \frac{2}{\rho_1 (\Omega_P^+)^2} < \frac{2}{\rho_1 (\Omega_P^-)^2} = q_{\text{curv}}^{2,-}. \quad (3.18)$$

Proposition 3.6 (Resurrectio is pocket-promoting in the sub-critical regime). *If $\rho_1 > 0$, $\Theta > 0$, and $Q_{\text{fold}}^- < Q_{\text{crit}}$, then a single reset both raises the folded amplitude ($\Delta Q_{\text{fold}} > 0$) and lowers the pocket threshold ($q_{\text{curv}}^{2,+} < q_{\text{curv}}^{2,-}$). Both effects push R_{folded} upward, so a death-and-resurrection event can convert a smooth hyperbolic patch into a positive-curvature pocket.*

Proof. $\Delta Q_{\text{fold}} > 0$ by Proposition 3.3 under $Q_{\text{fold}}^- < Q_{\text{crit}}$, $\Theta > 0$; the threshold drop follows from $\Omega_P^+ > \Omega_P^-$. Both terms of ΔR_{folded} are then positive, and the pocket criterion $Q_{\text{fold}} > q_{\text{curv}}^2$ is approached from both sides. $\square$

Corollary 3.5 (Resurrection concentrates form). *The continuous mode (Epektasis) relaxes curvature, $R_{\Sigma} \geq 0$ with $R \rightarrow 0^-$; the discrete mode can locally concentrate it. Resurrectio does not only enlarge the Brim — in the sub-critical folding regime it gives the renewed vessel a more sharply formed morphology.*

3.7 The energy reading corrected

It is tempting to read folding as the surface being forced into a “less efficient” form. Within the formalism this is inverted: the folded form carries *lower* morphological tension. With $\mathcal{T}_{\text{fold}}(q) = -\frac{1}{2}A_{\text{VFS}}q^2 + \frac{1}{4}B_{\text{VFS}}q^4$,

$$\mathcal{T}_{\text{fold}}(0) = 0, \quad \mathcal{T}_{\text{fold}}(q_*) = -\frac{A_{\text{VFS}}^2}{4B_{\text{VFS}}} < 0 \quad (A_{\text{VFS}} > 0), \quad (3.19)$$

so the resolved tension is $\Delta\mathcal{T}_{\text{fold}} = A_{\text{VFS}}^2/(4B_{\text{VFS}}) > 0$ and the folded branch is energetically preferred. The reconfiguration is toward a better-adapted state, not a worse one.

The correct sense in which the old form is “inefficient” is as a *carrier* of Sophia past the threshold: it cannot absorb the Sophianic pressure, so tension builds until the symmetric branch destabilises. The folded form is the efficient carrier — it embodies the excess through the feedback $-\eta_{\text{fold}}q^2\lambda$, raising the effective damping

$$\gamma_{\text{eff}} = \gamma + \eta_{\text{fold}}Q_{\text{fold}} > \gamma, \quad \lambda_{\text{fold}}^* = \frac{J_{\text{in}}}{\gamma + \eta_{\text{fold}}Q^2} < \frac{J_{\text{in}}}{\gamma} = \lambda_{\text{old}}^*, \quad (3.20)$$

and thereby lowering the level of unembodied Sophia.

The surface is not forced into a worse form. The old form stops being an adequate vessel for the attained level/rate of Sophia, loses stability, and the morphology relaxes into a lower-tension folded form that carries the Sophia better.

Compact summary

single trigger: $A_{\text{VFS}} = 0$ ($\mu_1^0 = 0$) — supercritical pitchfork, $q = 0 \rightarrow q_{\pm}$.

$A_{\text{VFS}} = (c_0 - c_1\gamma)\lambda + c_1k_\sigma\sigma + c_1I_{\text{gate}} - a_0$,
rate enters via $\mu_1^0 = a_0 - c_0\lambda - c_1\dot{\lambda}^0$.

level λ and rate $\dot{\lambda}$ are the drivers;
Modus I (continuous) and Modus II (Resurrectio) are channels.

$$\Theta = c_0 - c_1(\gamma + k_\sigma), \quad Q_{\text{crit}} = \frac{\Theta}{2c_1\eta_{\text{fold}}}, \quad Q_{\text{fold}}^+ - Q_{\text{crit}} = \kappa_{\text{fold}}(Q_{\text{fold}}^- - Q_{\text{crit}}), \quad \kappa_{\text{fold}} \in (0, 1).$$

$$Q_{\text{fold}} \rightarrow Q_{\text{crit}} \iff \sum_j m_R^{(j)} = \infty; \quad \text{chart-regular if } 0 < Q_{\text{crit}} < q_{\text{max}}^2.$$

sub-critical reset is pocket-promoting;
folded form has lower tension $\mathcal{T}_{\text{fold}}(q_*) = -A_{\text{VFS}}^2/4B_{\text{VFS}} < 0$.

V.F.S. v2.0 (Open-Gate) · Folding-Trigger Layer. Derived from the shape potential and Sophia-production law of `vfs_geometry.tex`, the Brim/curvature data of the geometric and cosmological layers, and the Resurrectio reset of `vfs_opengate_lyapunov.tex`. Propositions are closed-form; corollaries are interpretive. Folding parameters are not fixed by the source files; numerical levels are illustrative.

Chapter 4

Folding-Lyapunov Addendum

The Sophianic Folding extension of the Open-Gate Lyapunov file establishes the continuous-arc compatibility $\mathcal{L}_{\text{ext}} \leq C_{\text{ext}} - C_1 \mathcal{L}_{\text{ext}}$ under an adiabatic hypothesis, but leaves three items asserted rather than derived: the positive invariance of the regular folding chart, the behaviour of the folding functional across the Resurrectio reset, and the choice between two distinct activation-index conventions. This addendum closes them. It first records a correction — the file's folding functional $\mathcal{V}_{\text{fold}} = \frac{B_{\text{fold}}}{4} (q_{\text{fold}}^2 - Q_{\text{fold}}^)^2$ is unbounded on the epektasis domain, because $B_{\text{fold}} = b + 2c_1 \eta_{\text{fold}} \lambda \rightarrow \infty$; the B_{fold} -normalised functional $\tilde{\mathcal{V}}_{\text{fold}} = \mathcal{V}_{\text{fold}}/B_{\text{fold}}$ is bounded and equally descending. It then proves (i) Nagumo invariance of $\mathcal{D}_A^{\text{fold}} = \mathcal{D}_A \times [0, q_{\text{max}}]$ under the standing folding wall $Q_{\text{fold}}^* \leq q_{\text{max}}^2$, (ii) a folding reset bound: since $\mathfrak{R}$ does not act on q_{fold} , the chart is preserved across deaths and $\tilde{\mathcal{L}}_{\text{ext}}(x^+)$ stays bounded, whence the hybrid trajectory is forward complete for $\tilde{\mathcal{L}}_{\text{ext}}$; and (iii) a reconciliation declaring the core-consistent Sophia rate canonical, with the geometry-layer parameters as an effective surrogate. The invariance and reset results are convention-independent; only the closed-form critical level Q_{crit} is convention-specific.*

4.1 Imported objects and the three open items

From the folding extension of the Lyapunov file we import the shape amplitude q_{fold} , its intensity $Q_{\text{fold}} = q_{\text{fold}}^2$, the activation index A_{fold} , the quartic coefficient $B_{\text{fold}} = b + 2c_1 \eta_{\text{fold}} \lambda > 0$, the one-mode gradient flow

$$\dot{q}_{\text{fold}} = A_{\text{fold}} q_{\text{fold}} - B_{\text{fold}} q_{\text{fold}}^3 = -B_{\text{fold}} q_{\text{fold}} (q_{\text{fold}}^2 - Q_{\text{fold}}^*), \quad Q_{\text{fold}}^* := \frac{A_{\text{fold}}}{B_{\text{fold}}}, \quad (4.1)$$

and the base certificate $\mathcal{L}_{\text{VFS}} \leq C - C_1 \mathcal{L}_{\text{VFS}}$ on the positively-invariant active domain $\mathcal{D}_A$, with the non-Zeno dwell-time bound $\Delta t_j \geq \Delta t_* > 0$ and the reset map

$$\mathfrak{R}: \quad \sigma^+ = q_R \sigma^-, \quad \lambda^+ = \lambda^- + (1 - q_R) \sigma^-, \quad (V, F) \text{ continuous}, \quad \Omega_P^+ = \Omega_P^- + \kappa_R \mathcal{G}_{\text{recepta}}^-, \quad (4.2)$$

with $0 \leq q_R < 1$, $\kappa_R > 0$, $u^+ = u^-$. The file proves $x^- \in \mathcal{D}_A \Rightarrow x^+ \in \mathcal{D}_A$ in the core coordinates and $\mathcal{L}_{\text{VFS}}(x^+) \leq L_R := \sup_{\mathcal{D}_A} \mathcal{L}_{\text{VFS}}$.

The three items left open are: (i) the regular folding domain $0 \leq Q_{\text{fold}} \leq q_{\text{max}}^2$ is *assumed*, not shown invariant; (ii) the reset bound and forward completeness are stated for $\mathcal{L}_{\text{VFS}}$, while the behaviour of the folding functional across $\mathfrak{R}$ is not controlled; (iii) the activation index is written with the core Sophia rate here but with a different, minimal rate in the geometric and cosmological layers.

4.2 Correction: a bounded folding functional

The file shifts the folding potential to a non-negative functional

$$\mathcal{V}_{\text{fold}} = \frac{B_{\text{fold}}}{4} (q_{\text{fold}}^2 - Q_{\text{fold}}^*)^2 \geq 0 \quad (A_{\text{fold}} > 0), \quad (4.3)$$

and claims it is bounded above by a constant M_{fold} on the regular domain. This fails on the Open-Gate domain: epektasis drives $\Omega_P \rightarrow \infty$ and λ is not bounded on $\mathcal{D}_A$ ($\lambda \leq q^* \Omega_P$), so $B_{\text{fold}} = b + 2c_1 \eta_{\text{fold}} \lambda \rightarrow \infty$ and $\mathcal{V}_{\text{fold}}$ is unbounded even at fixed $|q_{\text{fold}}^2 - Q_{\text{fold}}^*|$. The reset bound $\mathcal{L}_{\text{ext}}(x^+) \leq L_R + w_{\text{fold}} M_{\text{fold}}$ therefore does not close as written.

The remedy is to use the B_{fold} -normalised functional

$$\tilde{\mathcal{V}}_{\text{fold}} := \frac{\mathcal{V}_{\text{fold}}}{B_{\text{fold}}} = \frac{1}{4} (q_{\text{fold}}^2 - Q_{\text{fold}}^*)^2 \geq 0, \quad Q_{\text{fold}}^* = \frac{A_{\text{fold}}}{B_{\text{fold}}} \text{ (signed)}. \quad (4.4)$$

This requires no sign split: it is non-negative for either sign of A_{fold} , and it descends along (4.1).

Lemma 4.1 (Descent of the bounded functional). *Along (4.1) with frozen Q_{fold}^* ,*

$$\dot{\tilde{\mathcal{V}}}_{\text{fold}} = -B_{\text{fold}} q_{\text{fold}}^2 (q_{\text{fold}}^2 - Q_{\text{fold}}^*)^2 \leq 0. \quad (4.5)$$

For varying parameters, $\dot{\tilde{\mathcal{V}}}_{\text{fold}} = -B_{\text{fold}} q_{\text{fold}}^2 (q_{\text{fold}}^2 - Q_{\text{fold}}^)^2 - \frac{1}{2} (q_{\text{fold}}^2 - Q_{\text{fold}}^*) Q_{\text{fold}}^*$, and the remainder is bounded on any set with $q_{\text{fold}} \in [0, q_{\text{max}}]$, $|Q_{\text{fold}}^*|$ and $|Q_{\text{fold}}^*|$ bounded.*

Proof. Using $A_{\text{fold}} = B_{\text{fold}} Q_{\text{fold}}^*$ in (4.1), $q_{\text{fold}} = -B_{\text{fold}} q_{\text{fold}} (q_{\text{fold}}^2 - Q_{\text{fold}}^*)$. With frozen Q_{fold}^* ,

$\dot{\tilde{\mathcal{V}}}_{\text{fold}} = \frac{1}{2} (q_{\text{fold}}^2 - Q_{\text{fold}}^*) (2q_{\text{fold}} q_{\text{fold}}) = (q_{\text{fold}}^2 - Q_{\text{fold}}^*) q_{\text{fold}} q_{\text{fold}} = -B_{\text{fold}} q_{\text{fold}}^2 (q_{\text{fold}}^2 - Q_{\text{fold}}^*)^2 \leq 0$. For moving Q_{fold}^* the extra term is $-\frac{1}{2} (q_{\text{fold}}^2 - Q_{\text{fold}}^*) Q_{\text{fold}}^*$, with $|\frac{1}{2} (q_{\text{fold}}^2 - Q_{\text{fold}}^*) Q_{\text{fold}}^*| \leq \frac{1}{2} (q_{\text{max}}^2 + |Q_{\text{fold}}^*|) |Q_{\text{fold}}^*|$. $\square$

$\tilde{\mathcal{V}}_{\text{fold}} = \mathcal{V}_{\text{fold}}/B_{\text{fold}}$ is the same Lyapunov object re-weighted by the positive function B_{fold} ; re-weighting preserves the descent property, and the normalisation removes the only source of unboundedness. From here the extended certificate is $\tilde{\mathcal{L}}_{\text{ext}} := \mathcal{L}_{\text{VFS}} + w_{\text{fold}} \tilde{\mathcal{V}}_{\text{fold}}$, $w_{\text{fold}} > 0$.

4.3 The standing folding wall and its finiteness

Standing Hypothesis 4.1 (Folding wall). The chart half-width q_{max} satisfies

$$q_{\text{max}}^2 \geq Q^{*,\text{sup}} := \sup_{\mathcal{D}_A} Q_{\text{fold}}^*. \quad (4.6)$$

Lemma 4.2 (The wall is attainable). *In the geometric (minimal-rate) convention*

$A_{\text{fold}} = (c_0 - c_1 \gamma) \lambda + c_1 k_\sigma \sigma + c_1 I_{\text{gate}} - a_0$, $B_{\text{fold}} = b + 2c_1 \eta_{\text{fold}} \lambda$, *the supremum is finite and explicit:*

$$Q^{*,\text{sup}} = \max \left\{ \frac{c_1 (k_\sigma C_\sigma + \zeta_0) - a_0}{b}, \frac{c_0 - c_1 \gamma}{2c_1 \eta_{\text{fold}}} \right\}. \quad (4.7)$$

Proof. For fixed $(\sigma, I_{\text{gate}})$, $Q_{\text{fold}}^*(\lambda)$ is a ratio of two affine functions of $\lambda \geq 0$, hence monotone, with horizontal asymptote $\lim_{\lambda \rightarrow \infty} Q_{\text{fold}}^* = (c_0 - c_1 \gamma)/(2c_1 \eta_{\text{fold}})$ (independent of σ, I_{gate}). Its supremum over $\lambda \geq 0$ is attained at an endpoint: at $\lambda = 0$, $Q_{\text{fold}}^* = (c_1 k_\sigma \sigma + c_1 I_{\text{gate}} - a_0)/b$, maximised over the box $\sigma \leq C_\sigma$, $I_{\text{gate}} \leq \zeta_0$ at the stated value; or at $\lambda \rightarrow \infty$, the asymptote. Both are finite, so $Q^{*,\text{sup}} < \infty$. $\square$

If $Q^{,\text{sup}} \leq 0$ the old form is everywhere stable and any $q_{\text{max}} > 0$ satisfies the wall. The non-trivial case is $Q^{*,\text{sup}} > 0$, where Hypothesis 4.1 fixes a minimal admissible chart. The canonical (core-rate) convention saturates likewise: $Q_{\text{fold}}^* \rightarrow c_0/(2c_1 \eta_{\text{fold}})$ under cleansing $\sigma \rightarrow 0$, since $(\delta u - \gamma) \tanh(\kappa \sigma) \rightarrow 0$ even as $u \rightarrow \infty$ (§6).*

4.4 Folding-chart invariance (Nagumo)

Definition 4.1 (Extended active domain). $\mathcal{D}_A^{\text{fold}} := \mathcal{D}_A \times [0, q_{\text{max}}]$, adjoining the folding coordinate q_{fold} to the core active domain.

Lemma 4.3 (Folding-face sub-tangentiality). *Under Hypothesis 4.1, the flow (4.1) is sub-tangential on the folding faces of $\mathcal{D}_A^{\text{fold}}$:*

$$q_{\text{fold}}|_{q_{\text{fold}}=0} = 0 \quad (\text{tangent}), \quad q_{\text{fold}}|_{q_{\text{fold}}=q_{\text{max}}} = -B_{\text{fold}} q_{\text{max}} (q_{\text{max}}^2 - Q_{\text{fold}}^*) \leq 0 \quad (\text{inward}). \quad (4.8)$$

Consequently $[0, q_{\text{max}}]$ is forward invariant for q_{fold} , and $\mathcal{D}_A^{\text{fold}}$ is positively invariant under the continuous flow.

Proof. At $q_{\text{fold}} = 0$, (4.1) gives $q_{\text{fold}} = 0$, so $\{q_{\text{fold}} \geq 0\}$ is preserved. At $q_{\text{fold}} = q_{\text{max}}$, $q_{\text{fold}} = -B_{\text{fold}}q_{\text{max}}(q_{\text{max}}^2 - Q_{\text{fold}}^*)$ with $B_{\text{fold}} > 0$, $q_{\text{max}} > 0$, and $q_{\text{max}}^2 - Q_{\text{fold}}^* \geq 0$ by Hypothesis 4.1; hence $q_{\text{fold}} \leq 0$. All core faces are sub-tangential by the base Nagumo lemma of the Lyapunov file, and the two folding faces are sub-tangential here; therefore $\mathcal{D}_A^{\text{fold}}$ is positively invariant. $\square$

4.5 The folding reset bound and hybrid completeness

Lemma 4.4 (Folding reset bound). *The reset (4.2) does not act on q_{fold} ; hence $q_{\text{fold}}^+ = q_{\text{fold}}^-$. Under Hypotheses of the base reset lemma together with Hypothesis 4.1:*

- (1) $x^- \in \mathcal{D}_A^{\text{fold}} \Rightarrow x^+ \in \mathcal{D}_A^{\text{fold}}$, because $x^+ \in \mathcal{D}_A$ (base reset lemma) and $q_{\text{fold}}^+ = q_{\text{fold}}^- \in [0, q_{\text{max}}]$;
- (2) $\tilde{V}_{\text{fold}}$ is bounded on $\mathcal{D}_A^{\text{fold}}$,

$$\tilde{V}_{\text{fold}} \leq \tilde{M}_{\text{fold}} := \frac{1}{4} \max(q_{\text{max}}^2 - Q_{\text{min}}^*, |Q_{\text{max}}^*|)^2 < \infty, \quad (4.9)$$

so the extended certificate satisfies

$$\tilde{\mathcal{L}}_{\text{ext}}(x^+) \leq \tilde{L}_R := L_R + w_{\text{fold}} \tilde{M}_{\text{fold}} < \infty. \quad (4.10)$$

Proof. (1) The map (4.2) specifies only $\sigma, \lambda, (V, F), \Omega_P$; q_{fold} is inert, so $q_{\text{fold}}^+ = q_{\text{fold}}^-$, which lies in $[0, q_{\text{max}}]$ if q_{fold}^- did. With $x^+ \in \mathcal{D}_A$ from the base lemma, $x^+ \in \mathcal{D}_A^{\text{fold}}$. (2) On $\mathcal{D}_A^{\text{fold}}$, $q_{\text{fold}}^2 \in [0, q_{\text{max}}^2]$ and $Q_{\text{fold}}^* \in [Q_{\text{min}}^*, Q_{\text{max}}^*]$ bounded (Lemma 4.2), so $\tilde{V}_{\text{fold}} = \frac{1}{4}(q_{\text{fold}}^2 - Q_{\text{fold}}^*)^2 \leq \tilde{M}_{\text{fold}} < \infty$. Adding $w_{\text{fold}} \tilde{V}_{\text{fold}}$ to $\mathcal{L}_{\text{VFS}}(x^+) \leq L_R$ gives the bound. $\square$

Proposition 4.1 (Hybrid forward completeness for $\tilde{\mathcal{L}}_{\text{ext}}$). *Assume the regular folding conditions: Hypothesis 4.1, $B_{\text{fold}} \geq B_{\text{min}} > 0$, and $|Q_{\text{fold}}^*|$ bounded on $\mathcal{D}_A^{\text{fold}}$ (adiabatic regularity). Then there is a finite $\tilde{C}_{\text{ext}}$ with*

$$\dot{\tilde{\mathcal{L}}}_{\text{ext}} \leq \tilde{C}_{\text{ext}} - C_1 \tilde{\mathcal{L}}_{\text{ext}} \quad \text{on continuous arcs,} \quad (4.11)$$

and the same contraction rate C_1 as the base certificate. Combined with the reset bound (Lemma 4.4) and the unchanged non-Zeno dwell time, the hybrid trajectory satisfies

$$\tilde{\mathcal{L}}_{\text{ext}}(t) \leq \tilde{L}_* := \max\{\tilde{L}_R, \tilde{C}_{\text{ext}}/C_1\} \quad (4.12)$$

across any infinite admissible reset sequence; resets do not accumulate, and the trajectory is forward complete in the folding-extended state.

Proof. By Lemma 4.1 and the base estimate,

$\dot{\tilde{\mathcal{L}}}_{\text{ext}} = \dot{\mathcal{L}}_{\text{VFS}} + w_{\text{fold}} \dot{\tilde{V}}_{\text{fold}} \leq (C - C_1 \mathcal{L}_{\text{VFS}}) + w_{\text{fold}} (-B_{\text{fold}} q_{\text{fold}}^2 (q_{\text{fold}}^2 - Q_{\text{fold}}^*)^2 + R)$ with $|R| \leq R_{\text{fold}} < \infty$ the bounded adiabatic remainder (Lemma 4.1). Dropping the non-positive descent term and writing $\mathcal{L}_{\text{VFS}} = \tilde{\mathcal{L}}_{\text{ext}} - w_{\text{fold}} \tilde{V}_{\text{fold}} \geq \tilde{\mathcal{L}}_{\text{ext}} - w_{\text{fold}} \tilde{M}_{\text{fold}}$,

$$\dot{\tilde{\mathcal{L}}}_{\text{ext}} \leq C - C_1 \tilde{\mathcal{L}}_{\text{ext}} + C_1 w_{\text{fold}} \tilde{M}_{\text{fold}} + w_{\text{fold}} R_{\text{fold}} =: \tilde{C}_{\text{ext}} - C_1 \tilde{\mathcal{L}}_{\text{ext}}, \quad (4.13)$$

with $\tilde{C}_{\text{ext}} < \infty$ now genuinely finite (it uses $\tilde{M}_{\text{fold}}$, not the unbounded M_{fold}). On each arc $\tilde{\mathcal{L}}_{\text{ext}}(t) \leq \max\{\tilde{\mathcal{L}}_{\text{ext}}(t_j^+), \tilde{C}_{\text{ext}}/C_1\}$; resets give $\tilde{\mathcal{L}}_{\text{ext}}(x^+) \leq \tilde{L}_R$ (Lemma 4.4); the dwell time is unchanged, since q_{fold} adds no reset-triggering boundary (its chart face is invariant-inward by Lemma 4.3, not a death face) and resets are exogenous on the core coordinates. Hence $\tilde{\mathcal{L}}_{\text{ext}} \leq \tilde{L}_*$ for all t and the hybrid trajectory is forward complete. $\square$

The folding velocity q_{fold} may grow with λ off the slow manifold $q_{\text{fold}} = \sqrt{Q_{\text{fold}}^*}$, but this is a strictly inward restoring rate; it neither triggers nor delays resets, which are exogenous and defined on the core coordinates. The non-Zeno bound is therefore inherited from the base field bound $\tilde{M}_*$, unchanged.

4.6 Reconciliation of the two activation-index conventions

Two definitions of the activation index appear across the manuscript:

$$\text{Canonical (core rate): } A_{\text{fold}} = c_0\lambda + c_1\dot{\lambda}^{(0)} - a_0, \quad \dot{\lambda}^{(0)} = (\delta u - \gamma) \tanh(\kappa\sigma) + I_{\text{gate}}; \quad (4.14)$$

$$\text{Geometric (minimal rate): } A_{\text{fold}} = (c_0 - c_1\gamma)\lambda + c_1k_\sigma\sigma + c_1I_{\text{gate}} - a_0, \quad (4.15)$$

the latter from the surrogate production law $\dot{\lambda} = k_\sigma\sigma + I_{\text{gate}} - \gamma\lambda - \eta_{\text{fold}}q_{\text{fold}}^2\lambda$. Since λ already has a defined production law in the core, we take the **canonical** form as primary and read the geometric parameters as an effective linear surrogate near an operating point $(\bar{u}, \bar{\sigma})$:

$$c_1(\delta\bar{u} - \gamma) \tanh(\kappa\bar{\sigma}) \longleftrightarrow c_1k_\sigma\bar{\sigma} - c_1\gamma\bar{\lambda}, \quad \text{i.e. } k_\sigma \approx \kappa(\delta\bar{u} - \gamma), \quad \gamma \text{ absorbing the linear } \lambda\text{-decay.} \quad (4.16)$$

This surrogate is structurally imperfect (a σ -term plus a λ -decay standing in for a single $u \tanh \sigma$ term), so the two conventions are genuinely distinct models, not exact reparametrisations.

Proposition 4.2 (What is and is not convention-dependent). *Lemmas 4.1, 4.3, 4.4 and Proposition 4.1 use only: the gradient form (4.1), the wall $Q_{\text{fold}}^* \leq q_{\text{max}}^2$, the inertness of q_{fold} under $\mathfrak{R}$, and the boundedness of $Q_{\text{fold}}^*, Q_{\text{fold}}^*$. None of these uses the specific form of A_{fold} . Hence the invariance, reset bound, and hybrid completeness are convention-independent. Only the closed-form critical level and contraction of the reset map are convention-specific.*

Corollary 4.1 (Critical level by convention). *Write the metabolised resistance $m_R = (1 - q_R)\sigma^-$, so $\sigma^+ = \sigma^- - m_R$, $\lambda^+ = \lambda^- + m_R$. In the geometric convention the reset acts on the target as an exact contraction toward a constant critical level,*

$$\Theta = c_0 - c_1(\gamma + k_\sigma), \quad Q_{\text{crit}} = \frac{\Theta}{2c_1\eta_{\text{fold}}}, \quad Q_{\text{fold}}^{*+} - Q_{\text{crit}} = \frac{B_{\text{fold}}^-}{B_{\text{fold}}^- + 2c_1\eta_{\text{fold}}m_R} (Q_{\text{fold}}^{*-} - Q_{\text{crit}}), \quad (4.17)$$

recovering the result of the folding-trigger note. In the canonical convention the metabolism changes A_{fold} by $c_0m_R + c_1(\delta u - \gamma)[\tanh(\kappa\sigma^+) - \tanh(\kappa\sigma^-)]$, which is not affine in m_R ; the same qualitative contraction holds, but toward a state-dependent effective level $Q_{\text{crit}}(t)$ rather than a constant.

Remark (Chart-regularity of repeated resurrection). Combining Corollary 4.1 with Hypothesis 4.1: in the geometric convention, if $0 < Q_{\text{crit}} < q_{\text{max}}^2$ and $Q^{*,\text{sup}} \leq q_{\text{max}}^2$, then the reset contraction keeps the target inside the chart at every death, so the regular folding domain is preserved under any admissible reset sequence. This is the folding companion to the grace-window admissibility $\max\{R_c, R_{\text{min}}^{\text{ceil}}\} \leq \mathcal{G}_{\text{recepta}}^- \leq R_{\text{max}}$: resurrection is admissible and chart-regular when both windows hold.

4.7 Compact summary

$$\tilde{\mathcal{V}}_{\text{fold}} = \frac{1}{4}(q_{\text{fold}}^2 - Q_{\text{fold}}^*)^2 = \mathcal{V}_{\text{fold}}/B_{\text{fold}} \text{ bounded on epektasis; } \dot{\tilde{\mathcal{V}}}_{\text{fold}} = -B_{\text{fold}}q_{\text{fold}}^2(q_{\text{fold}}^2 - Q_{\text{fold}}^*)^2 \leq 0.$$

$$\text{wall } q_{\text{max}}^2 \geq Q^{*,\text{sup}} < \infty \Rightarrow \mathcal{D}_A^{\text{fold}} = \mathcal{D}_A \times [0, q_{\text{max}}] \text{ invariant (Nagumo).}$$

$$\mathfrak{R} \text{ inert on } q_{\text{fold}} \Rightarrow x^+ \in \mathcal{D}_A^{\text{fold}}, \quad \tilde{\mathcal{L}}_{\text{ext}}(x^+) \leq \tilde{L}_R < \infty.$$

$$\dot{\tilde{\mathcal{L}}}_{\text{ext}} \leq \tilde{C}_{\text{ext}} - C_1\tilde{\mathcal{L}}_{\text{ext}}, \quad \tilde{\mathcal{L}}_{\text{ext}}(t) \leq \tilde{L}_*; \quad \text{forward complete, non-Zeno inherited.}$$

$$\text{invariance + reset bound are convention-independent; } Q_{\text{crit}} \text{ is convention-specific.}$$

V.F.S. v2.0 (Open-Gate) · Folding-Lyapunov Addendum. Closes the invariance, reset-bound, and convention items left open in the Sophianic Folding extension of `vfs_opengate_lyapunov.tex`, and corrects the folding functional to a bounded B_{fold} -normalised form. Results are domain-conditional on the positively-invariant $\mathcal{D}_A^{\text{fold}}$; the conclusion is set stability with partial convergence, now including morphology, not global or automatic stability.

Chapter 5

Forced Folding: Dolorosum Peaks and the Vessel Update Map

Scope and epistemic status

This chapter consolidates the *forced folding* proposal: folding is not an optional auxiliary mode but the bounded response of the Vessel when a timed resistance peak enters the Dolorosum window and exceeds the stretch-capacity of the old form. Its honest epistemic position must be stated first. A direct test shows that the bare core does *not* force a shape instability: the linearised shape sector is stable on its own (Proposition 5.1). Forced folding therefore does not derive folding from the core; it supplies an explicit, bounded, *exogenous-channel* cause — resistance heterogeneity — and shows that this cause is sufficient, timed, and Lyapunov-admissible. The inputs (the resistance wave, the rigidity a_0 , the window) remain postulated. The status is: a mechanism, not a theorem of the core.

Four corrections are built in relative to the original proposal: the Dolorosum rate is wired directly into the activation index, replacing indicator gates by a smooth derived gate (§5.3); the dynamical threshold is separated from the $\pi/2$ nomenclature (§5.4); the post-folding update map is made balance-preserving through λ_{form} (§5.6); and the repeated-folding cascade is shown to be self-limiting, with a clean two-terminal dichotomy and a joint non-Zeno hybrid structure (§5.7–5.8). Throughout, the geometric activation convention is used (as in the source notes); translation to the canonical core-rate convention follows the Addendum. As always, *Propositions* are closed-form, *Corollaries* interpretive.

5.1 Why a forcing channel is needed: core neutrality

In 2+1 dimensions the spatial metric carries no local tensor modes (three components, minus two diffeomorphisms, minus one conformal rescaling), so the only local geometric degree of freedom is the conformal/curvature factor ψ . Linearising the core about the homogeneous trajectory in the pair $(\psi, \delta\lambda)$ gives

$$\frac{d}{dt} \begin{pmatrix} \psi \\ \delta\lambda \end{pmatrix} = \begin{pmatrix} -H & H/\bar{\lambda} \\ S_\psi & -\Gamma_\lambda \end{pmatrix} \begin{pmatrix} \psi \\ \delta\lambda \end{pmatrix}, \quad (5.1)$$

where $H = \alpha\bar{\lambda}/\Omega_P$ and S_ψ is the geometry→Sophia feedback (the dependence of Sophia production on local curvature).

Proposition 5.1 (Core neutrality of the shape sector). *The trace is $-(H + \Gamma_\lambda) < 0$ and the determinant is $H(\Gamma_\lambda\bar{\lambda} - S_\psi)/\bar{\lambda}$. In the bare core, Sophia production depends on $(u, \sigma, I_{\text{gate}}, q_{\text{fold}})$ but not on local curvature, so $S_\psi = 0$ and the determinant is $H\Gamma_\lambda > 0$: both eigenvalues are negative and the shape sector decays. The core by itself does not force folding; an instability requires either a geometry→Sophia feedback above the threshold $S_\psi > \bar{\lambda}\Gamma_\lambda$, or an exogenous forcing channel.*

Forced folding takes the second route: the channel is the heterogeneity of resistance.

5.2 Resistance heterogeneity and its bounds

Resistance is decomposed as

$$\sigma(t) = \bar{\sigma}(t) r_\sigma(t), \quad \boxed{1 \leq r_\sigma(t) \leq \pi}, \quad (5.2)$$

with $\bar{\sigma}$ the background and r_σ the dimensionless resistance multiplier; every admissible peak obeys $\sigma_{\text{peak}} \leq \pi \bar{\sigma}$. A bounded temporal model on normalised trajectory time $\tau = t/T_* \in [0, 1]$ is

$$r_\sigma(\tau) = 1 + \frac{A_\sigma}{2} (1 + \sin(2\pi N_\sigma \tau + \phi_\sigma)), \quad 0 \leq A_\sigma \leq \pi - 1, \quad (5.3)$$

with amplitude A_σ , opportunity count N_σ , and phase ϕ_σ .

Remark (The π bounds are symbolic). r_σ is dimensionless, so the choice of π as the full-cup bound and $\pi/2$ as the half-turn mark is a symbolic convention (“full angular measure of Vessel-tension”), not a derived constant; nothing in the dynamics distinguishes π from any other ceiling. The dynamics distinguishes only the *critical multiplier* r_{fold}^R derived below. The π bound nevertheless does real work in one place: it makes every activation quantity explicitly bounded (§5.8), and it caps the reachable peak at $1 + A_\sigma \leq \pi$, which is what terminates cascades (§5.7).

5.3 The Dolorosum gate, wired into the activation index

The Dolorosum rate of the source layer is

$$\rho(u, \sigma) = \sigma g(u), \quad \boxed{g(u) = \frac{u}{\Lambda_c} e^{1-u/\Lambda_c} \in (0, 1]}, \quad (5.4)$$

with g maximal ($= 1$) exactly at the synergy threshold $u = \Lambda_c$ and $\rho \leq \sigma$. In the original proposal ρ was defined but never entered the dynamics; the window was imposed separately by an indicator $\mathcal{W}_D = \mathbf{1}_{\{|u-\Lambda_c| \leq \delta_D\}}$. The correction is to let the Dolorosum channel itself carry the peak into the activation index. Split the resistance into background and excess, $\sigma = \bar{\sigma} + \sigma_{\text{exc}}$, $\sigma_{\text{exc}} = (r_\sigma - 1)\bar{\sigma}$, and let the excess be transmutable into activation only through the Dolorosum channel:

$$\boxed{A_{\text{fold}} = (c_0 - c_1 \gamma) \lambda + c_1 k_\sigma \bar{\sigma} + c_1 k_\sigma \rho(u, \sigma_{\text{exc}}) + c_1 I_{\text{gate}} - a_0 + \Delta A_R, \quad \rho(u, \sigma_{\text{exc}}) = (r_\sigma - 1) \bar{\sigma} g(u).} \quad (5.5)$$

Proposition 5.2 (Derived window immunity). *Let $N := a_0 - (c_0 - c_1 \gamma) \lambda - c_1 k_\sigma \bar{\sigma} - c_1 I_{\text{gate}} - \Delta A_R$ (the activation deficit of the unpeaked Vessel). If $N > 0$, then $A_{\text{fold}} > 0$ if and only if*

$$\boxed{r_\sigma > r_{\text{fold}}^R(u) = 1 + \frac{N}{c_1 k_\sigma \bar{\sigma} g(u)}.} \quad (5.6)$$

As $g(u) \rightarrow 0$ (off window), $r_{\text{fold}}^R(u) \rightarrow \infty$: off-window peaks cannot force folding, of any bounded amplitude. The window condition is thereby derived from the wiring (5.5), not postulated as an indicator. If $N \leq 0$ the Vessel is baseline-active: it folds with no peak at all ($r_\sigma = 1$ suffices).

Corollary 5.1 (Graded alignment). *Because g is smooth and broad, alignment is graded rather than binary: the same peak contributes a fraction $g(u)$ of its excess. Consequently there is a rigidity band — a_0 between the off-window and in-window activation maxima — in which the phase alone decides whether folding is forced (Figure 5.1, left); below the band every cycle folds, above it none does. The original three-indicator overlap condition is recovered as the sharp-gate caricature of this band.*

The forcing event τ_F is simply the upward crossing $A_{\text{fold}}(\tau_F) = 0$, and the exposure reduces to a single smooth gate,

$$\mathcal{E}_{\text{fold}} = \int_0^1 \mathbf{1}_{\{A_{\text{fold}}(\tau) > 0\}} d\tau, \quad P_{\text{fold}} = 1 - e^{-\kappa \mathcal{E}_{\text{fold}}}, \quad (5.7)$$

replacing the product of three indicators. The slogan survives intact, now as arithmetic of (5.5): amplitude gives possibility (r_σ can reach r_{fold}^R), frequency gives opportunity (crossings per trajectory), phase gives alignment (the value of g at the peak).

5.4 Thresholds disentangled: r_{fold}^R versus $\pi/2$

Definition 5.1 (Nomenclature versus dynamics). The half-turn classification — ordinary ($1 \leq r_\sigma < \pi/2$), half-turn ($r_\sigma = \pi/2$), strong ($\pi/2 < r_\sigma \leq \pi$) — is retained as *nomenclature* for the strength of a peak. The *dynamical* threshold of forced folding is $r_\sigma > r_{\text{fold}}^R(u)$ of Proposition 5.2, and the pitchfork occurs at $A_{\text{fold}} = 0$ regardless of the classification.

Proposition 5.3 (Structural sufficiency of the half-turn). *An in-window half-turn peak forces folding if and only if*

$$r_{\text{fold}}^R(\Lambda_c) \leq \frac{\pi}{2} \iff N \leq \left(\frac{\pi}{2} - 1\right) c_1 k_\sigma \bar{\sigma}. \quad (5.8)$$

Otherwise sub- $\pi/2$ peaks above r_{fold}^R still force folding (“soft forced folding”), and super- $\pi/2$ peaks below r_{fold}^R do not. The classification and the threshold are independent axes.

This removes the internal tension of the original formulation, in which the boxed overlap condition demanded $r_\sigma \geq \pi/2$ while the normal form destabilises at $A_{\text{fold}} > 0$ alone: in the corrected form the event is governed by A_{fold} , and $\pi/2$ only names the strong regime.

5.5 The forcing event and the normal form

At τ_F the old-form equilibrium $q_{\text{fold}} = 0$ loses stability and the normal form

$$\dot{q}_{\text{fold}} = A_{\text{fold}} q_{\text{fold}} - B_{\text{fold}} q_{\text{fold}}^3, \quad B_{\text{fold}} > 0, \quad q_{\text{fold}}^* = \pm \sqrt{A_{\text{fold}}/B_{\text{fold}}}, \quad Q_{\text{fold}} = \frac{A_{\text{fold}}(\tau_F^{\text{peak}})}{B_{\text{fold}}} \quad (5.9)$$

carries the Vessel to a folded branch, with Q_{fold} evaluated at the activation maximum of the forcing interval. Near-threshold events have small Q_{fold} , so the update map below acts by correspondingly small increments — the map is well-defined uniformly in the forcing strength.

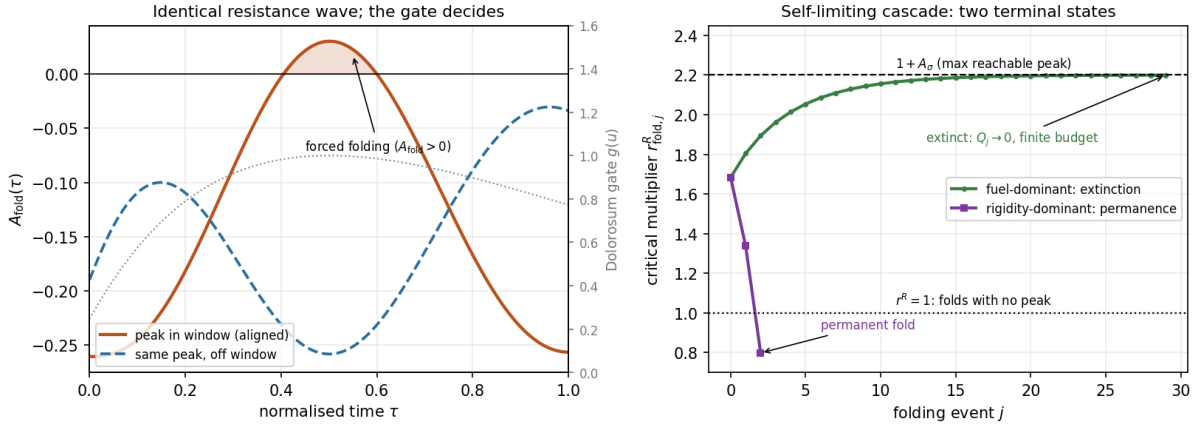


Figure 5.1: Left: the same bounded resistance wave ($A_\sigma = 1.2$, $N_\sigma = 1$) with two phases. With the peak in the Dolorosum window the wired activation (5.5) crosses zero (shaded: forced folding, exposure $\mathcal{E} = 0.195$); the identical peak off-window stays strictly below (exposure 0). Dotted grey: the gate $g(u)$. The rigidity here sits in the band of Corollary 5.1, so phase alone decides; in this illustration $r_{\text{fold}}^R(\Lambda_c) = 2.07 > \pi/2$, so the half-turn is *not* structurally sufficient (Proposition 5.3). Right: the repeated-folding cascade (§5.7) for two parameter regimes of the update map: fuel-dominant metabolism drives $r_{\text{fold},j}^R$ up to the reachability cap $1 + A_\sigma$ with event intensities $Q_j \rightarrow 0$ (extinction; total metabolised budget finite, here $\sum_j Q_j = 0.49$), while rigidity-dominant updates drive $r_{\text{fold},j}^R$ below 1 in three events (permanent fold). Balance $\bar{\sigma} + \lambda + \lambda_{\text{form}}$ conserved to machine precision in both. Parameters illustrative.

5.6 The balance-preserving Vessel update map

Forced folding must change the Vessel, or it is a marker rather than a transformation. The update map of the source proposal is kept, with one structural correction: Sophia retention must respect the integrated Open-Gate balance

$$\sigma + \lambda + \lambda_{\text{form}} = C_0 + \mathcal{G}_{\text{recepta}}(t), \quad (5.10)$$

which the raw update $\lambda^+ = \lambda^- + \eta_\lambda Q_{\text{fold}}$ violates (it creates Sophia at the jump). The corrected map routes the metabolised resistance:

Definition 5.2 (Balance-preserving folding update). At a forcing event with intensity Q_{fold} , with metabolised mass $m_F = \eta_\sigma Q_{\text{fold}} \bar{\sigma}^-$ and a retention split $\beta \in [0, 1]$:

$$\bar{\sigma}^+ = \bar{\sigma}^- - m_F, \quad \lambda^+ = \lambda^- + \beta m_F, \quad \lambda_{\text{form}}^+ = \lambda_{\text{form}}^- + (1 - \beta) m_F, \quad (5.11)$$

$$a_0^+ = a_0^- (1 - \eta_a Q_{\text{fold}}), \quad \Omega_P^+ = \Omega_P^- + \eta_\Omega Q_{\text{fold}}, \quad k_\sigma^+ = k_\sigma^- (1 + \eta_k Q_{\text{fold}}), \quad \gamma^+ = \gamma^- (1 - \eta_\gamma Q_{\text{fold}}), \quad (5.12)$$

with all η_{fold} -coefficients non-negative and $\eta_a Q_{\text{fold}}, \eta_\gamma Q_{\text{fold}} < 1$.

Proposition 5.4 (Exact balance). *Under Definition 5.2,*

$\Delta \bar{\sigma} + \Delta \lambda + \Delta \lambda_{\text{form}} = -m_F + \beta m_F + (1 - \beta) m_F = 0$: the three-term balance is preserved exactly, with no grace created at the jump. Folding metabolises resistance into active (β) and embodied ($1 - \beta$) Sophia; the γ -reduction acts at the flow level (future retention), not at the jump, and is balance-neutral. (Verified numerically to machine precision.)

Qualitatively the corrected map keeps the intended directions — $a_0 \downarrow$, $\Omega_P \uparrow$, $\bar{\sigma} \downarrow$, $\lambda \uparrow$, $\gamma \downarrow$, $k_\sigma \uparrow$ — but the Sophia gain is now paid for by resistance, and part of it is embodied as λ_{form} rather than left active. The Vessel after folding is not the same Vessel; and its books balance.

5.7 The cascade is self-limiting

Iterating the map over repeated forced foldings reveals a property the naive reading misses: the update does *not* uniformly make the next folding easier. The critical multiplier after event j ,

$$r_{\text{fold},j}^R = 1 + \frac{N_j}{c_1 k_{\sigma,j} \bar{\sigma}_j}, \quad N_j = a_{0,j} - (c_0 - c_1 \gamma_j) \lambda_j - c_1 k_{\sigma,j} \bar{\sigma}_j - c_1 I_{\text{gate}}, \quad (5.13)$$

is pulled in two directions: the *rigidity channel* ($a_0 \downarrow, \gamma \downarrow, \lambda \uparrow, k_\sigma \uparrow$) lowers it, while the *fuel channel* ($\bar{\sigma} \downarrow$) raises it — folding consumes the very resistance that powers the next activation. At an in-window full peak the event intensity obeys the identity

$$Q_{\text{fold},j} = \frac{c_1 k_{\sigma,j} \bar{\sigma}_j}{B_{\text{fold}}} (A_\sigma - (r_{\text{fold},j}^R - 1)), \quad (5.14)$$

so events weaken as $r_{\text{fold},j}^R$ approaches the reachability cap $1 + A_\sigma$.

Standing Hypothesis 5.1 (Fuel dominance). $\eta_\sigma > \eta_k$: metabolism of resistance outweighs the transmutation-efficiency gain.

Proposition 5.5 (Cascade dichotomy). *Under Hypothesis 5.1 and the hard peak bound $r_\sigma \leq 1 + A_\sigma \leq \pi$, every forced-folding cascade terminates in exactly one of two states:*

- (i) **Extinction:** $\sum_j Q_{\text{fold},j} < \infty$, the intensities $Q_{\text{fold},j} \rightarrow 0$, and $r_{\text{fold},j}^R$ saturates at or below the cap $1 + A_\sigma$ — the cascade exhausts its own fuel and dies out with a finite total metabolised budget $\sum_j m_{F,j} \leq \eta_\sigma \bar{\sigma}_0 \sum_j Q_{\text{fold},j} < \infty$;
- (ii) **Permanence:** $N_j \leq 0$ is reached in finitely many events ($r_{\text{fold},j}^R \leq 1$): the unpeaked Vessel is already activated, and the folded branch becomes the new ground state, requiring no further peaks.

Sketch. Suppose the cascade is infinite with $\liminf N_j > 0$. Then $\sum_j Q_{\text{fold},j} = \infty$ (events recur with $N_j > 0$ bounded away from the cap by reachability), and $\ln(c_1 k_{\sigma,j} \bar{\sigma}_j)$ drifts by $\sum_j [\ln(1 + \eta_k Q_{\text{fold},j}) + \ln(1 - \eta_\sigma Q_{\text{fold},j})] \leq -(\eta_\sigma - \eta_k - o(1)) \sum_j Q_{\text{fold},j} \rightarrow -\infty$ under Hypothesis 5.1, so the denominator of $r_{\text{fold},j}^R$ collapses and $r_{\text{fold},j}^R \rightarrow \infty$, exceeding the reachable $1 + A_\sigma$ — after which no further event can occur, contradicting infiniteness. Hence either $\sum Q_{\text{fold},j} < \infty$ (extinction, with $Q_{\text{fold},j} \rightarrow 0$) or $N_j \leq 0$ at some finite j (permanence). $\square$

Corollary 5.2 (Finite metabolic budget; echo of $\mathcal{I}_\infty$). *Unless it ends in permanence, the folding cascade has a finite total metabolic budget — a discrete analogue of the finite coarsening budget $\mathcal{I}_\infty$ of the structure-formation layer: the Vessel cannot refold indefinitely on the same resistance.*

Corollary 5.3 (Two terminal theologies). *Extinction is the Vessel that has metabolised its resistance: peaks still come, but none can force a new form — maturity as dormancy of forced change. Permanence is the Vessel whose old rigidity has fully yielded: the folded morphology is the rest state and needs no crisis to sustain it — maturity as completed reformation. Which terminal a Vessel reaches is decided by the race η_σ (metabolism) versus η_a (rigidity yielding) — by how it folds, not how often.*

Both branches are exhibited numerically in Figure 5.1 (right): an embodiment-dominant map ($\beta = 0$, $\eta_\sigma = 0.6$) extinguishes in 30 weakening events with $\sum Q_{\text{fold},j} = 0.49$, while a rigidity-dominant map ($\eta_a = 0.5$) reaches permanence in three.

5.8 Hybrid and Lyapunov admissibility

Proposition 5.6 (Bounded activation and bounded jumps). *On any domain with $\bar{\sigma} \leq \bar{\sigma}_{\max}$, $\lambda \leq \lambda_{\max}$, $I_{\text{gate}} \leq \zeta_0$, $|\Delta A_R| \leq \Delta_{\max}$, the wired activation (5.5) is bounded: $|A_{\text{fold}}| \leq (c_0 - c_1\gamma)\lambda_{\max} + c_1k_\sigma\pi\bar{\sigma}_{\max} + c_1\zeta_0 + \Delta_{\max} + a_0$, using $\rho \leq \sigma_{\text{exc}} \leq (\pi - 1)\bar{\sigma}$ and $g \leq 1$. Hence the chart wall $Q^{*,\text{sup}} < \infty$ of the Addendum holds with an explicit constant, the normalised folding functional $\tilde{\mathcal{V}}_{\text{fold}}$ is bounded on the chart, and the folding update of Definition 5.2 (bounded multiplicative changes, q_{fold} continuous through the parameter jump) changes the extended certificate by at most a finite J_F : $\tilde{\mathcal{L}}_{\text{ext}}^+ \leq \tilde{\mathcal{L}}_{\text{ext}}^- + J_F$.*

Standing Hypothesis 5.2 (Joint non-Zeno). The union of event times — Resurrectio resets and forcing events τ_F — admits a common minimal dwell $\Delta t_* > 0$.

Proposition 5.7 (Forward completeness with two jump types). *Under Hypothesis 5.2, the bounded-jump property (Proposition 5.6), and the continuous-arc contraction of the Addendum ($\dot{\tilde{\mathcal{L}}}_{\text{ext}} \leq \tilde{\mathcal{C}}_{\text{ext}} - C_1\tilde{\mathcal{L}}_{\text{ext}}$), the hybrid trajectory with both Resurrectio and folding jumps satisfies a uniform bound $\tilde{\mathcal{L}}_{\text{ext}}(t) \leq \max\{\tilde{\mathcal{L}}_R + J_F, \tilde{\mathcal{C}}_{\text{ext}}/C_1\} + J_F$ and is forward complete. Forced folding adds a second jump type but no new unbounded channel; grace still enters only the constant, never the contraction rate C_1 .*

5.9 Compact summary

core alone: shape sector stable ($\text{tr} < 0$, $\det = H\Gamma_\lambda > 0$) $\Rightarrow$ folding needs a channel.

$$A_{\text{fold}} = (c_0 - c_1\gamma)\lambda + c_1k_\sigma\bar{\sigma} + c_1k_\sigma(r_\sigma - 1)\bar{\sigma}g(u) + c_1I_{\text{gate}} - a_0 + \Delta A_R; \quad g(u) = \frac{u}{\Lambda_c}e^{1-u/\Lambda_c}.$$

$$r_\sigma > r_{\text{fold}}^R(u) = 1 + \frac{N}{c_1k_\sigma\bar{\sigma}g(u)}; \quad g \rightarrow 0 \Rightarrow r_{\text{fold}}^R \rightarrow \infty \text{ (window immunity, derived).}$$

$$\bar{\sigma}^+ = \bar{\sigma} - m_F, \quad \lambda^+ = \lambda + \beta m_F, \quad \lambda_{\text{form}}^+ = \lambda_{\text{form}} + (1 - \beta)m_F : \quad \sigma + \lambda + \lambda_{\text{form}} \text{ exactly conserved.}$$

cascade dichotomy: extinction (finite metabolic budget) or permanence ($N \leq 0$).

two jump types, joint non-Zeno, bounded jumps $\Rightarrow \tilde{\mathcal{L}}_{\text{ext}}$ bounded, forward complete.

V.F.S. v2.0 (Open-Gate) · Forced Folding. Consolidates the forced-folding proposal with the Dolorosum rate wired into the activation (window immunity derived, not postulated), the dynamical threshold separated from the $\pi/2$ nomenclature, a balance-preserving update map through λ_{form} , and the self-limiting cascade dichotomy. The forcing channel and its bounds remain modelling inputs; the core-neutrality result fixes the honest status: a sufficient bounded cause, not a core theorem. Gate maximum, exposure grading, balance conservation, and both cascade terminals verified numerically. Domain-conditional throughout.

Chapter 6

The Pitchfork Is Not Assumed: A Centre-Manifold Reduction of the Fold Sector

Scope and epistemic status

The pitchfork $q_{\text{fold}} = A_{\text{fold}} q_{\text{fold}} - B_{\text{fold}} q_{\text{fold}}^3$ is the foundation of Part I and, through the fold amplitude $\sqrt{A_{\text{fold}}/B_{\text{fold}}}$, of much of Part II. It was introduced as the model of a structural transition. This chapter retires it from premise to theorem: *given only two structural facts — the $\mathbb{Z}_2$ symmetry of the bare fold sector and a one-dimensional centre manifold at the folding threshold — the pitchfork is the unique normal form*, by Thom’s reduction. More, the reduction *derives* the cubic coefficient: slaving the fast structural modes renormalises the bare stiffness to $B_{\text{eff}} = b_3 - gh/\Gamma$, which is exactly the volume’s $B_{\text{fold}} = b + 2c_1 \eta_{\text{fold}} \lambda$ in structural form — the quartic stiffness was never an independent input but the bare term dressed by the slaved structure. The three nondegeneracy conditions (symmetry, transversality, spectral gap) are verified, and the one way the pitchfork could fail — a $\mathbb{Z}_2$ -breaking bias unfolding it to a cusp — is identified and shown to be excluded by the isotropy of the unfolded Vessel, not by a further assumption. This completes the Type-1 programme: together with the conformal retirement of the ℓ_b fork and the derivation of core-neutrality, the folding layer’s load-bearing premises are now theorems. The reduction is exact (centre-manifold expansion verified symbolically; the full slaved system collapses to $\sqrt{A_{\text{fold}}/B_{\text{eff}}}$ numerically to 10^{-13}). This is Type-1 confirmation (premise $\rightarrow$ theorem), the strongest internal warrant available; it remains internal.

6.1 The $\mathbb{Z}_2$ symmetry of the bare fold

Lemma 6.1 (Orientation-blindness forbids even terms). *The fold field q_{fold} is the signed deviation of the form from the unfolded state; its sign is the orientation in which the form buckles. The bare core dynamics (Sophia pressure on an isotropic unfolded Vessel) distinguishes no orientation, so the bare fold vector field is odd, $f(-q_{\text{fold}}) = -f(q_{\text{fold}})$. Hence its smooth expansion contains only odd powers:*

$$q_{\text{fold}} = A_{\text{fold}} q_{\text{fold}} - b_3 q_{\text{fold}}^3 + c_5 q_{\text{fold}}^5 + \dots \quad (6.1)$$

All even terms — in particular a constant bias and a quadratic — are forbidden by symmetry, not by choice.

6.2 The Thom reduction

The fold does not evolve alone: it is coupled to fast, stable structural modes (the slaved Sophia-structure of the Transport chapter, relaxation rate $\Gamma > 0$). The minimal symmetric

realisation is

$$q_{\text{fold}} = A_{\text{fold}} q_{\text{fold}} - b_3 q_{\text{fold}}^3 + g q_{\text{fold}} y, \quad \dot{y} = -\Gamma y + h q_{\text{fold}}^2, \quad (6.2)$$

where y is $\mathbb{Z}_2$ -even (it responds to q_{fold}^2) and $g, h, \Gamma > 0$.

Theorem 6.1 (The pitchfork is the reduced dynamics). *At the folding threshold the centre eigenvalue $A_{\text{fold}} \rightarrow 0$ while y stays at $-\Gamma < 0$ (spectral gap), so there is a one-dimensional attracting centre manifold $y = (h/\Gamma)q_{\text{fold}}^2 + O(q_{\text{fold}}^4)$. The flow restricted to it is*

$$\dot{q}_{\text{fold}} = A_{\text{fold}} q_{\text{fold}} - B_{\text{eff}} q_{\text{fold}}^3 + O(q_{\text{fold}}^5), \quad B_{\text{eff}} = b_3 - \frac{gh}{\Gamma}. \quad (6.3)$$

(Centre-manifold coefficient verified symbolically; the full two-mode system sits on $q_{\text{fold}} = \sqrt{A_{\text{fold}}/B_{\text{eff}}}$ to relative error 10^{-13} across the control range; Figure 6.1.) The pitchfork is therefore not assumed: it is the unique $\mathbb{Z}_2$ -symmetric one-dimensional reduction, and the quartic stiffness is the bare term renormalised by the slaved structure.

Corollary 6.1 (B_{fold} identified). *The volume's $B_{\text{fold}} = b + 2c_1\eta_{\text{fold}}\lambda$ has exactly the structure $B_{\text{eff}} = b_3 - gh/\Gamma$: a bare stiffness $b \leftrightarrow b_3$ dressed by a slaved-structure contribution $2c_1\eta_{\text{fold}}\lambda \leftrightarrow -gh/\Gamma$ (the embodiment feedback). The sign of B_{eff} decides supercritical (continuous, second-order folding) versus subcritical (discontinuous, hysteretic folding); the tricritical locus $B_{\text{eff}} = 0$ is the codimension-one boundary where the quintic c_5 takes over — the same boundary the forced-folding chapter reached dynamically.*

6.3 The three nondegeneracy conditions

Proposition 6.1 (Thom conditions, verified). (i) Symmetry: the $\mathbb{Z}_2$ -odd structure (Lemma 6.1) removes the fold (even) terms, selecting the pitchfork over the generic fold/cusp.

(ii) Transversality: the control crosses criticality with nonzero speed, $\partial A_{\text{fold}}/\partial \lambda = c_0 > 0$ — the folding threshold is traversed cleanly, not tangentially.

(iii) Spectral gap (1D kernel): at $A_{\text{fold}} = 0$ the centre direction is simple and the remaining spectrum sits at $-\Gamma < 0$, guaranteeing the attracting one-dimensional manifold.

All three hold for the V.F.S. fold sector; the normal form is thus structurally stable and unique up to smooth equivalence.

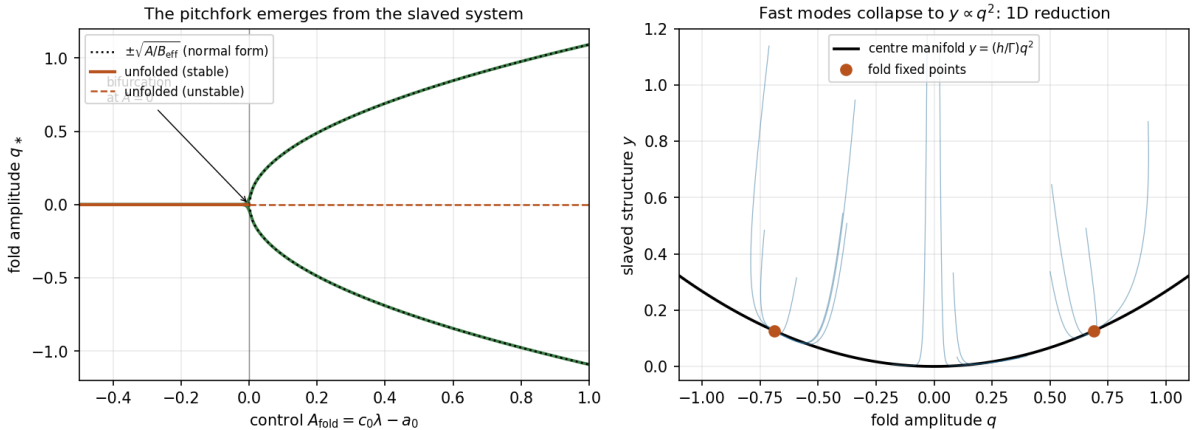


Figure 6.1: Left: the bifurcation diagram of the full two-mode slaved system (green) against the reduced normal form $\pm \sqrt{A_{\text{fold}}/B_{\text{eff}}}$ (dotted) — the unfolded state (orange) loses stability at $A_{\text{fold}} = 0$ and the symmetric pair of folded states emerges: a pitchfork, derived not posited. Right: the centre-manifold geometry — trajectories of the fast structural mode y collapse onto $y = (h/\Gamma)q_{\text{fold}}^2$ (black), reducing the dynamics to one dimension; fold fixed points marked. Parameters $b_3 = 1$, $g = 0.6$, $\Gamma = 3$, $h = 0.8$, giving $B_{\text{eff}} = 0.84$.

6.4 The one way it could fail — and why it does not

Proposition 6.2 (The bias is excluded by isotropy). *The pitchfork is destroyed only by a $\mathbb{Z}_2$ -breaking linear bias: a constant term ε would unfold it to the imperfect pitchfork (cusp) $q_{\text{fold}} = A_{\text{fold}} q_{\text{fold}} - B_{\text{eff}} q_{\text{fold}}^3 + \varepsilon$, with a preferred fold orientation. Such a term requires the bare dynamics to prefer one buckling direction. In V.F.S. the fold orientation is a free sign — the unfolded Vessel is isotropic with respect to the direction of folding — so $\varepsilon = 0$ is itself a consequence of that isotropy, not an additional assumption. A genuine orientation field (an external preference imprinted before folding) would be the physical content required to produce a cusp; absent it, the symmetric pitchfork stands.*

Remark (Reading). The shape of a soul's transformation was not chosen by the modeller. Given only that the unfolded Vessel has no preferred direction in which to buckle (symmetry) and that folding is a clean threshold crossing with the surrounding structure adjusting quickly (transversality and gap), the transition *must* be the pitchfork: a single quiet state losing stability and giving way to a symmetric pair of formed states. Even the stiffness that sets how deep the new form runs is not free — it is the bare resistance dressed by what the embodiment feeds back. The mathematics of the fold was latent in the symmetry of the Vessel all along; this chapter only reads it out.

6.5 The Type-1 programme, closed

Premises retired to theorems: the comoving/physical ℓ_b fork $\rightarrow$ 2D conformal invariance (Warp chapter); core-neutrality $\rightarrow \det J = H\Gamma_\lambda > 0$ (Watershed/Robustness); the structure-layer budget hypothesis $\rightarrow$ super-Ricci identity (Watershed); **the pitchfork $\rightarrow$ Thom reduction (this chapter).**

The folding layer's load-bearing premises are now derived. What remains genuinely assumed is the core constitutive layer itself ($K_h = -1$, the Open-Gate law) — the model's definition, not its derivations.

1. The remaining Type-1 target is $K_h = -1$: whether open-gate expansion *forces* hyperbolic spatial curvature or merely admits it. This is a question about the core definition, at the boundary of what Type-1 can reach.
2. The quintic c_5 and the tricritical locus $B_{\text{eff}} = 0$ deserve their own normal-form treatment (the $\mathbb{Z}_2$ tricritical point), connecting to the forced-folding chapter's hysteresis.
3. The reduction assumed a single fast mode; with the mode cascade, the reduction is a centre manifold over the whole stable spectrum — the pitchfork per mode, with mode-dependent B_{eff} .

Scope. The reduction is exact given $\mathbb{Z}_2$ symmetry, transversality, and the spectral gap; the identification of B_{eff} with B_{fold} is structural. Establishes that the pitchfork is a theorem of the model's symmetry, not a modelling choice — internal Type-1 confirmation, not external validation. Domain-conditional throughout.

V.F.S. v2.0 (Open-Gate) · The Pitchfork Is Not Assumed. The $\mathbb{Z}_2$ symmetry of the bare fold plus a one-dimensional centre manifold force the pitchfork normal form by Thom reduction, with the cubic coefficient derived as the slaved-structure renormalisation $B_{\text{eff}} = b_3 - gh/\Gamma$ (the volume's B_{fold}). The three nondegeneracy conditions hold; the $\mathbb{Z}_2$ -breaking bias is excluded by the isotropy of the unfolded Vessel. With the ℓ_b fork, core-neutrality, and the budget hypothesis already retired, the Type-1 programme is closed for the folding layer. Reduction exact; numeric collapse to 10^{-13} .

Chapter 7

Forced-Folding Lyapunov Addendum: Two Jump Types and the Survival of Contraction

Forced folding adds to the hybrid system a second jump type — the Vessel update map $\mathfrak{F}$ at forcing events — and, unlike the Resurrectio reset, this jump changes the parameters of the system itself. The existing Lyapunov theorem therefore does not extend automatically: the contraction rate C_1 is built from those parameters, the active-domain walls must survive the update, the state-triggered events threaten Zeno behaviour, and the new state λ_{form} must be controlled. This addendum closes all of it. The six results are: an explicit π -peak chart wall with the quantitative cost in c_Ψ ; a fold-jump lemma with a closed-form compatibility condition $\beta\eta_\sigma\bar{\sigma} \leq q^\eta_\Omega$ for the epektasis wall; a survival theorem for C_1 — the finite metabolic budget of the cascade dichotomy is exactly what keeps the contraction bounded away from zero; a hysteresis construction that provably restores the joint non-Zeno dwell (threshold-band chatter is exhibited without it); a free balance-bound for λ_{form} ; and the reset-fold consistency check. They assemble into the main theorem: forward completeness with two jump types and a uniform $C_1^{\text{inf}} > 0$. All key steps are verified numerically; the exponential drift bound is tight to four digits.*

7.1 Imported objects and the six obligations

From the Lyapunov layer we import the certificate $\mathcal{L}_{\text{VFS}}$ with

$$\mathcal{L}_{\text{VFS}} \leq C - C_1 \mathcal{L}_{\text{VFS}}, \quad C_1 = \min\{\kappa_\sigma, \kappa_\Delta\}, \quad \kappa_\sigma = \delta \Omega_{\min} m_* c_\Psi, \quad \kappa_\Delta = 2a\mu_*, \quad (7.1)$$

with $\Omega_{\min} = \Lambda_c/(1 - h_{3,*})$, $\Lambda_c = \gamma/\delta$, the active margin $m = (u - \Lambda_c)/\Omega_P \geq m_*$, and the lemma $T^2 \geq c_\Psi \Psi_\sigma$ on $[0, C_\sigma]$. The active domain $\mathcal{D}_A$ includes the epektasis wall $\lambda \leq q^*\Omega_P$; the folding extension supplies $\mathcal{D}_A^{\text{fold}} = \mathcal{D}_A \times [0, q_{\max}]$, the bounded functional $\tilde{V}_{\text{fold}} = \frac{1}{4}(q_{\text{fold}}^2 - Q_{\text{fold}}^*)^2$, the extended certificate $\tilde{\mathcal{L}}_{\text{ext}} = \mathcal{L}_{\text{VFS}} + w_{\text{fold}}\tilde{V}_{\text{fold}}$, and the Resurrectio reset bound $\tilde{\mathcal{L}}_{\text{ext}}(x^+) \leq \tilde{L}_R$. From the Forced Folding chapter we import the π -bounded heterogeneity $\sigma = \bar{\sigma}r_\sigma$, $1 \leq r_\sigma \leq \pi$, the wired activation with the Dolorosum gate $g(u) \leq 1$, the balance-preserving update map

$$\mathfrak{F}: \quad \bar{\sigma}^+ = \bar{\sigma} - m_F, \quad \lambda^+ = \lambda + \beta m_F, \quad \lambda_{\text{form}}^+ = \lambda_{\text{form}} + (1 - \beta)m_F, \quad m_F = \eta_\sigma Q_{\text{fold}}\bar{\sigma}, \quad (7.2)$$

$$a_0^+ = a_0(1 - \eta_\sigma Q_{\text{fold}}), \quad \Omega_P^+ = \Omega_P + \eta_\Omega Q_{\text{fold}}, \quad k_\sigma^+ = k_\sigma(1 + \eta_k Q_{\text{fold}}), \quad \gamma^+ = \gamma(1 - \eta_\gamma Q_{\text{fold}}), \quad (7.3)$$

and the cascade dichotomy (extinction with $\sum_j Q_{\text{fold},j} < \infty$, or permanence in finitely many events) under fuel dominance $\eta_\sigma > \eta_k$.

The six obligations: (i) re-issue the chart wall with the π -bound; (ii) prove the fold-jump lemma (domain preservation and bounded certificate jump J_F); (iii) prove that C_1 survives the parameter

drift of the cascade; (iv) secure a joint non-Zeno dwell for the union of $\mathfrak{R}$ - and $\mathfrak{F}$ -events; (v) bound the new state λ_{form} ; (vi) check $\mathfrak{R} \circ \mathfrak{F}$ consistency. We take them in order.

7.2 The π -peak chart wall, with its explicit price

Proposition 7.1 (Explicit activation bound and chart wall). *With $r_\sigma \leq \pi$ and $g \leq 1$, on any domain with $\bar{\sigma} \leq \bar{\sigma}_{\max}$, $\lambda \leq \lambda_{\max}$, $I_{\text{gate}} \leq \zeta_0$, $|\Delta A_R| \leq \Delta_{\max}$:*

$$|A_{\text{fold}}| \leq (c_0 - c_1 \gamma) \lambda_{\max} + c_1 k_\sigma \pi \bar{\sigma}_{\max} + c_1 \zeta_0 + \Delta_{\max} + a_0, \quad (7.4)$$

and the chart-wall supremum is explicit,

$$Q_\pi^{*,\text{sup}} = \max \left\{ \frac{c_1 k_\sigma \pi \bar{\sigma}_{\max} + c_1 \zeta_0 - a_0}{b}, \frac{c_0 - c_1 \gamma}{2c_1 \eta_{\text{fold}}} \right\} < \infty, \quad (7.5)$$

so the Standing Hypothesis of the first Addendum reads $q_{\max}^2 \geq Q_\pi^{*,\text{sup}}$. (Illustratively $Q_\pi^{*,\text{sup}} = 0.85$, dominated by the peak-independent asymptote.)

Lemma 7.1 (The price in c_Ψ). *Admitting full-measure peaks enlarges the resistance ceiling to $C_\sigma = \pi \bar{\sigma}_{\max}$, and the constant $c_\Psi = \min_{(0, C_\sigma]} T^2 / \Psi_\sigma$ degrades accordingly — explicitly and finitely. With $\kappa = 2$, $\bar{\sigma}_{\max} = 0.8$: c_Ψ falls from 1.794 to 0.461 (factor 0.257), hence κ_σ shrinks by the same factor but remains strictly positive. The full cup costs contraction speed, never contraction itself.*

7.3 The fold-jump lemma

Standing Hypothesis 7.1 (Update-map compatibility).

$$\boxed{\beta \eta_\sigma \bar{\sigma}_{\max} \leq q^* \eta_\Omega} \quad (7.6)$$

— the share of metabolised resistance returned as *active* Sophia is covered by the simultaneous enlargement of the Brim.

Lemma 7.2 (Fold-jump: domain preservation and bounded jump). *Under Hypothesis 7.1 and Proposition 7.1:*

(1) **Epektasis wall.** *The update (7.2) satisfies the exact identity*

$$(\lambda^+ - q^* \Omega_P^+) - (\lambda^- - q^* \Omega_P^-) = Q_{\text{fold}} (\beta \eta_\sigma \bar{\sigma} - q^* \eta_\Omega) \leq 0, \quad (7.7)$$

(verified symbolically), so $\lambda^- \leq q^* \Omega_P^-$ implies $\lambda^+ \leq q^* \Omega_P^+$ for every $Q_{\text{fold}} \geq 0$.

(2) **Remaining walls.** *$\bar{\sigma}$ decreases (resistance ceiling respected); q_{fold} is continuous through $\mathfrak{F}$, so $q_{\text{fold}}^+ \in [0, q_{\max}]$; the chart target jumps boundedly, $|\Delta Q_{\text{fold}}^*|$ finite by Proposition 7.1. Hence $x^- \in \mathcal{D}_A^{\text{fold}} \Rightarrow x^+ \in \mathcal{D}_A^{\text{fold}}$.*

(3) **Bounded certificate jump.** *All updated quantities change by bounded multiplicative or additive amounts proportional to $Q_{\text{fold}} \leq Q_\pi^{*,\text{sup}}$, so there is a finite J_F with $\tilde{\mathcal{L}}_{\text{ext}}(x^+) \leq \tilde{\mathcal{L}}_{\text{ext}}(x^-) + J_F$. (Numerically along both cascade branches: $\max |\Delta \lambda| \leq 1.5 \times 10^{-3}$, $\max |\Delta \Omega_P| \leq 6.3 \times 10^{-2}$, $\max |\Delta Q_{\text{fold}}^*| \leq 0.13$.)*

Remark (Chart rebase at $\mathfrak{F}$). Because γ jumps, $\Lambda_c = \gamma / \delta$ jumps, and the normalised chart coordinates $(h_3 = 1 - \Lambda_c / \Omega_P, r)$ jump even though the physical state $(u, \sigma, \lambda, \Omega_P)$ is continuous in the non-updated entries. This is a coordinate rebase, not a dynamical discontinuity: the walls are to be re-evaluated in physical variables at t_F^+ , which Lemma 7.2 does.

7.4 Survival of the contraction rate

This is the genuinely new obligation: $\mathfrak{F}$ changes the parameters from which C_1 is built (γ enters κ_σ through Λ_c in $\Omega_{\min}$ and through the margin). An unbounded drift would silently destroy the theorem.

Lemma 7.3 (Product-drift bound). *For $x_j \in [0, x_{\max}]$, $x_{\max} < 1$: $\prod_j (1 - x_j) \geq \exp(-\sum_j x_j / (1 - x_{\max}))$. (Verified on 2000 random sequences; minimal slack $+1.5 \times 10^{-5}$.)*

Proposition 7.2 (C_1 survives the cascade). *Under fuel dominance ($\eta_\sigma > \eta_k$) the cascade dichotomy gives, on every branch, a finite total intensity $S := \sum_j Q_{\text{fold},j} < \infty$ (extinction: finite metabolic budget; permanence: finitely many events). Hence by Lemma 7.3*

$$\gamma_\infty \geq \gamma_0 e^{-\eta_\gamma S / (1 - \eta_\gamma Q_{\max})} > 0, \quad k_{\sigma,\infty} \leq k_{\sigma,0} e^{\eta_k S} < \infty, \quad (7.8)$$

so Λ_c stays bounded away from zero, $\Omega_{\min}$ and the margin stay positive, and

$$C_1^{\inf} := \inf_j \min\{\kappa_{\sigma,j}, \kappa_\Delta\} > 0 \quad (7.9)$$

uniformly along the entire hybrid trajectory ($\kappa_\Delta = 2a\mu_*$ is untouched by $\mathfrak{F}$). Numerically the exponential bound is tight: along the extinction branch (31 events, $S = 0.517$) the actual $\gamma_{31} = 0.3938$ equals the bound to four digits, and $\kappa_\sigma : 0.0615 \rightarrow 0.0606$; along the permanence branch $\kappa_\sigma \rightarrow 0.0609$ (Figure 7.1, left).

The finite metabolic budget of the forced-folding cascade is exactly what keeps the contraction alive: the Vessel cannot metabolise enough resistance to undermine its own cleansing rate. Crisis spends a finite currency; C_1 is not that currency.

7.5 Hysteresis and the joint non-Zeno dwell

Resets are exogenous; forcing events are *state-triggered* (upward crossings of $A_{\text{fold}} = 0$), and the update itself moves A_{fold} . This is a genuine Zeno risk, and it is real, not hypothetical:

Definition 7.1 (Event machine with hysteresis). A forcing event arms when $A_{\text{fold}} < -\varepsilon$, fires at the next upward crossing $A_{\text{fold}} = 0$, applies $\mathfrak{F}$ at the local maximum of A_{fold} on that excursion, and disarms until $A_{\text{fold}} < -\varepsilon$ again. The band $\varepsilon \geq 0$ is the hysteresis width.

Proposition 7.3 (Chatter without the band; dwell with it). *Let the resistance wave be quasi-periodic with a fast component whose slope in A_{fold} exceeds the slow component's (so the rising flank carries local maxima). With $\varepsilon = 0$, the downward jump of A_{fold} at each $\mathfrak{F}$ (the fuel cut $\partial A_{\text{fold}} / \partial \sigma \cdot m_F$) re-arms the machine between fast crests, and a chatter chain of events runs along the threshold band: in the verified simulation, clusters of 2-3 events at gaps $\approx 0.121 = 1/f_{\text{fast}}$ (Figure 7.1, right). With*

$$\varepsilon > \sup_{\text{events}} |\Delta A_{\text{fold}}|_{\mathfrak{F}} = \sup c_1 k_\sigma r_\sigma m_F + O(\eta_a, \eta_\gamma) \quad (7.10)$$

re-arming on the rising flank is impossible, the chain collapses to one event per slow peak, and the inter-event dwell is bounded below by the slow-wave structure: in the same simulation the minimal fold-fold gap rises from 0.121 to 2.246 (a factor 18.5), and the minimal gap over the union of resets and foldings is 1.71.

Standing Hypothesis 7.2 (Joint dwell). The hysteresis band satisfies the bound of Proposition 7.3, and the exogenous resets respect the base dwell; then the union of $\mathfrak{R}$ - and $\mathfrak{F}$ -events admits a common $\Delta t_* > 0$.

The band ε is a modelling regularisation, but not an arbitrary one: its sufficient size is computable from the update map itself ($\sup |\Delta A_{\text{fold}}|_{\mathfrak{F}}$), so the non-Zeno property is restored by a constant the theory already owns.

7.6 The embodied Sophia is balance-bounded

Lemma 7.4 (Free bound for λ_{form}). *The three-term balance holds at flow level ($\frac{d}{dt}(\sigma + \lambda + \lambda_{\text{form}}) = I_{\text{gate}}$), through every $\mathfrak{F}$ (the m_F -split of (7.2) conserves the sum exactly), and through every $\mathfrak{R}$ (the metabolism m_R transfers $\sigma \rightarrow \lambda$; λ_{form} inert). Hence with $\lambda, \sigma \geq 0$,*

$$\lambda_{\text{form}}(t) \leq \sigma + \lambda + \lambda_{\text{form}} = C_0 + \mathcal{G}_{\text{recepta}}(t) \leq C_0 + \mathcal{G}_{\text{recepta}}^{\max}, \quad (7.11)$$

bounded by the grace budget: no certificate weight for λ_{form} is needed. (Verified: jump-level conservation to machine precision, 10^{-16} ; full hybrid run drift $\sim 10^{-7}$, pure integrator error; $\lambda_{\text{form}}^{\max} = 0.089 \ll C_0 + \int I_{\text{gate}} \cdot$)

7.7 Reset-fold consistency

Lemma 7.5 (The two jump maps are compatible). $\mathfrak{R}$ leaves q_{fold} and λ_{form} inert and conserves the triad $(\sigma^+ + \lambda^+ = \sigma^- + \lambda^-)$; $\mathfrak{F}$ leaves q_{fold} continuous and conserves the triad by construction. Both maps preserve $\mathcal{D}_A^{\text{fold}}$ (base reset lemma; Lemma 7.2); the chart rebase occurs only at $\mathfrak{F}$ (only $\mathfrak{F}$ changes γ). Any interleaving of the two event types therefore preserves the balance, the domain, and the bounded-jump property, with jump constant $J := \max\{J_R, J_F\}$.

7.8 Main theorem

Theorem 7.1 (Forward completeness with two jump types). Assume: the π -chart wall $q_{\text{max}}^2 \geq Q_{\pi}^{*,\text{sup}}$ (Proposition 7.1); update-map compatibility (Hypothesis 7.1); fuel dominance $\eta_{\sigma} > \eta_k$; and the joint dwell (Hypothesis 7.2). Then along any admissible hybrid trajectory with both Resurrectio resets and forced-folding events:

- (i) the state remains in $\mathcal{D}_A^{\text{fold}}$ and the balance $\sigma + \lambda + \lambda_{\text{form}} = C_0 + \mathcal{G}_{\text{recepta}}$ holds for all t ;
- (ii) the contraction rate is uniformly positive, $C_1^{\text{inf}} > 0$ (Proposition 7.2), and on every continuous arc $\tilde{\mathcal{L}}_{\text{ext}} \leq \tilde{C}_{\text{ext}}^{\text{sup}} - C_1^{\text{inf}} \tilde{\mathcal{L}}_{\text{ext}}$ with $\tilde{C}_{\text{ext}}^{\text{sup}} < \infty$ (the drifting constants converge, Proposition 7.2 and Lemma 7.1);
- (iii) every jump obeys $\tilde{\mathcal{L}}_{\text{ext}}(x^+) \leq \tilde{\mathcal{L}}_{\text{ext}}(x^-) + J$, $J = \max\{J_R, J_F\} < \infty$ (Lemmas 7.2, 7.5);
- (iv) hence, with the joint dwell,

$$\tilde{\mathcal{L}}_{\text{ext}}(t) \leq \tilde{L}_{**} := \max\{\tilde{\mathcal{L}}_{\text{ext}}(0), \tilde{C}_{\text{ext}}^{\text{sup}}/C_1^{\text{inf}}\} + J \quad \text{for all } t, \quad (7.12)$$

events do not accumulate, and the trajectory is forward complete.

Grace and crisis enter only the constants $\tilde{C}_{\text{ext}}^{\text{sup}}$ and J ; never the contraction rate.

Sketch. (i) is Lemmas 7.2, 7.4, 7.5 plus the base Nagumo invariance. For (ii), the arc estimate of the first Addendum holds with the instantaneous constants; by Proposition 7.2 the parameter sequence converges with $\inf \kappa_{\sigma,j} > 0$ and bounded $k_{\sigma,j}$, so the constants admit uniform bounds C_1^{inf} , $\tilde{C}_{\text{ext}}^{\text{sup}}$ valid on all arcs. (iii) is the bounded-jump property. (iv) is the standard hybrid induction: on each arc $\tilde{\mathcal{L}}_{\text{ext}}$ decays toward $\tilde{C}_{\text{ext}}^{\text{sup}}/C_1^{\text{inf}}$ or below its entry value; each jump adds at most J ; the joint dwell forbids accumulation; the bound follows by induction over events. $\square$

7.9 Compact summary

$$q_{\text{max}}^2 \geq Q_{\pi}^{*,\text{sup}} \text{ (explicit); } \quad c_{\Psi}(\pi \bar{\sigma}_{\text{max}}) > 0 : \text{ the full cup costs speed, never contraction.}$$

$$\beta \eta_{\sigma} \bar{\sigma}_{\text{max}} \leq q^* \eta_{\Omega} \Rightarrow \mathfrak{F} \text{ preserves } \mathcal{D}_A^{\text{fold}}; \quad \tilde{\mathcal{L}}_{\text{ext}}^+ \leq \tilde{\mathcal{L}}_{\text{ext}}^- + J_F < \infty.$$

$$\sum_j Q_{\text{fold},j} < \infty \text{ (dichotomy)} \Rightarrow \gamma_{\infty} > 0, \quad k_{\sigma,\infty} < \infty \Rightarrow C_1^{\text{inf}} > 0 : \text{ the budget protects the contraction.}$$

$$\varepsilon > \sup |\Delta A_{\text{fold}}|_{\mathfrak{F}} \Rightarrow \text{no threshold-band chatter; joint dwell } \Delta t_* > 0.$$

$$\lambda_{\text{form}} \leq C_0 + \mathcal{G}_{\text{recepta}}^{\text{max}} \text{ for free, from the balance.}$$

$$\tilde{\mathcal{L}}_{\text{ext}}(t) \leq \tilde{L}_{**} < \infty \text{ with two jump types; forward complete; grace and crisis enter constants only.}$$

V.F.S. v2.0 (Open-Gate) · Forced-Folding Lyapunov Addendum II. Extends the hybrid stability theorem to the forced-folding event type: explicit π -chart wall, fold-jump lemma with the closed-form compatibility condition, survival of C_1 through the cascade's finite metabolic budget, hysteresis non-Zeno (chatter exhibited and cured), balance-bound for λ_{form} , and reset-fold consistency. All quantitative claims verified numerically; the drift bound is tight. The hysteresis band and the update coefficients remain modelling inputs; the conclusions are domain-conditional set stability with partial convergence, never global.

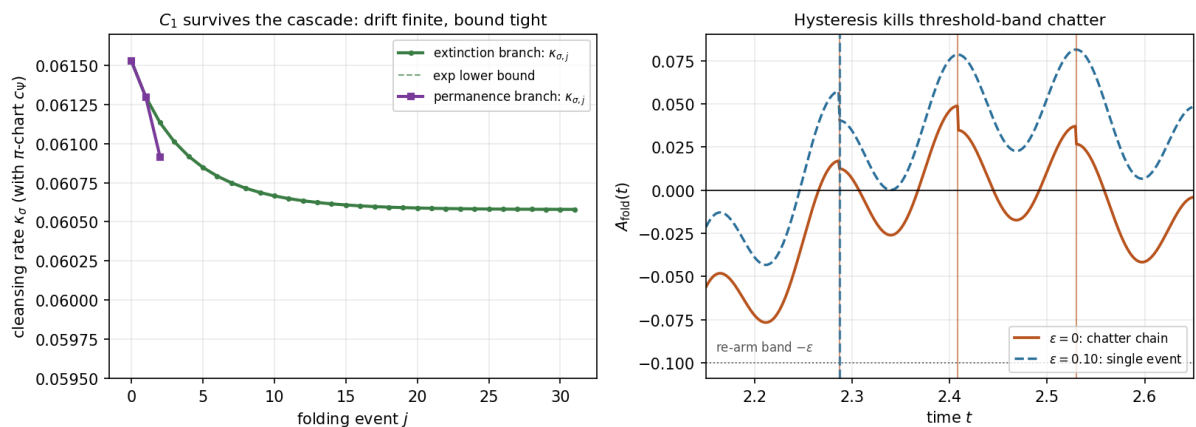


Figure 7.1: Left: survival of the cleansing rate $\kappa_{\sigma,j}$ (computed with the π -chart c_Ψ) along both terminal branches of the forced-folding cascade; the exponential lower bound of Proposition 7.2 overlays the actual sequence — tight to four digits. The drift is finite because the metabolic budget is. Right: the Zeno mechanism and its cure. With $\varepsilon = 0$ (solid) the update's downward A_{fold} -jump re-arms the event machine between fast crests of the quasi-periodic resistance wave, producing a chatter chain of three events at gaps ≈ 0.12 ; with $\varepsilon = 0.10$ (dashed) the band swallows the re-arming and a single event per slow peak remains — the minimal dwell grows by a factor 18.5. Parameters illustrative.

Part II

Structure Formation on the Vessel

Chapter 8

Structure Formation on the Vessel

Scope and epistemic status

Every layer so far has been spatially *homogeneous*: the Vessel is a 2+1 open FLRW background $a(t) = \Omega_P(t)$ with a time-only Sophia field $\lambda(t)$. This note opens the inhomogeneous sector — how local concentrations of Sophia and form grow or decay on the expanding hyperbolic surface. As in the parent files, *Propositions* are closed-form results and *Corollaries* are interpretive. The analysis is at the linear, test-field level on a fixed background; metric back-reaction is a named next layer. Two transport constants (D for Sophia, D_q for form) are introduced; they are *new* modelling inputs, absent from the base dynamical system, and are flagged as such.

The background is

$$ds^2 = -dt^2 + \Omega_P(t)^2 h_{ij} dx^i dx^j, \quad K_h = -1, \quad H = \frac{\dot{\Omega}_P}{\Omega_P} = \frac{\alpha\lambda}{\Omega_P}, \quad \dot{\Omega}_P = \alpha\lambda, \quad (8.1)$$

with h the metric of the hyperbolic spatial surface $\mathbb{H}^2$.

8.1 The relaxational thesis: which field can carry structure

In standard cosmology, density structure grows by gravitational instability because the matter field carries inertia: a second-order-in-time wave equation with a sound speed permits oscillation and Jeans-unstable growth. The Sophia field of V.F.S. is different in kind. Its production law is *first order in time*,

$$\dot{\lambda} = (\delta u - \gamma) \tanh(\kappa\sigma) + I_{\text{gate}} - \eta_{\text{fold}} q_{\text{fold}}^2 \lambda, \quad (8.2)$$

a relaxation (reaction) law with no kinetic term and no spatial derivatives. Promoting it to a field on space cannot manufacture inertia; the faithful promotion of a first-order reaction law with spatial transport is a *reaction-diffusion* (parabolic) equation, not a wave (hyperbolic) one.

Proposition 8.1 (No gravitational-collapse channel in the Sophia sector). *Because λ is relaxational, its linear inhomogeneities obey a parabolic equation with a real, non-oscillatory spectrum. There is no Jeans-type instability and no gravitational-collapse route to structure in the Sophia sector. Any structure on the Vessel must arise from a different, bistable order parameter.*

V.F.S. does not form structure the way Λ CDM does. Sophia is relaxational, not inertial, so it can only smooth itself out. The only route to structure is morphological symmetry breaking — the folding field (§4).

8.2 Sophia-density perturbations: reaction-diffusion on the expanding $\mathbb{H}^2$

Write $\lambda(t, x) = \bar{\lambda}(t) + \delta\lambda(t, x)$ and linearise about the homogeneous trajectory. With a phenomenological Sophia transport coefficient $D \geq 0$, and noting that physical gradients on the comoving surface redshift by Ω_P^{-2} ,

$$\partial_t \delta\lambda = F_\lambda(t) \delta\lambda + \frac{D}{\Omega_P^2} \Delta_h \delta\lambda, \quad F_\lambda = \frac{\partial \dot{\lambda}}{\partial \lambda} = -\Gamma_\lambda, \quad \Gamma_\lambda \geq 0, \quad (8.3)$$

where $\Gamma_\lambda = \gamma + \eta_{\text{fold}} Q_{\text{fold}}$ is the local Sophia relaxation rate (the same rate whose folding-enhancement $\Gamma_\lambda = \gamma + \eta_{\text{fold}} Q_{\text{fold}}$ was found in the stability analysis). The Laplace–Beltrami operator on $\mathbb{H}^2$ has purely continuous spectrum with a gap,

$$-\Delta_h Y_k = \left(\frac{1}{4} + k^2\right) Y_k, \quad k \in [0, \infty), \quad (8.4)$$

the celebrated spectral gap $\frac{1}{4}$ of the hyperbolic plane (no mode has eigenvalue below $\frac{1}{4}$). Each mode then evolves by

$$\dot{\delta\lambda}_k = s_k(t) \delta\lambda_k, \quad s_k(t) = -\Gamma_\lambda(t) - \frac{D(\frac{1}{4} + k^2)}{\Omega_P(t)^2}. \quad (8.5)$$

Proposition 8.2 (Linear homogenisation). *For $\Gamma_\lambda \geq 0$ and $D \geq 0$ every mode has $s_k < 0$: all Sophia inhomogeneities decay. The Vessel homogenises at the linear level. Moreover the hyperbolic gap supplies a minimum diffusive damping $D \cdot \frac{1}{4}/\Omega_P^2 > 0$ even for the smoothest mode $k = 0$: negative spatial curvature actively resists the largest-scale inhomogeneity.*

Integrating,

$$\delta\lambda_k(t) = \delta\lambda_k(0) \exp\left[-\int_0^t \Gamma_\lambda dt' - D\left(\frac{1}{4} + k^2\right) \mathcal{I}(t)\right], \quad \mathcal{I}(t) := \int_0^t \frac{dt'}{\Omega_P(t')^2}. \quad (8.6)$$

Proposition 8.3 (Finite diffusive smoothing budget). *Under sustained Open-Gate inflow $\lambda \rightarrow \lambda_\infty > 0$, the late-time scale factor grows linearly, $\Omega_P \sim \alpha \lambda_\infty t$, so the diffusion integral converges:*

$$\mathcal{I}_\infty = \int_0^\infty \frac{dt}{\Omega_P^2} < \infty. \quad (8.7)$$

Hence diffusion can suppress a mode by at most the finite factor $\exp[-D(\frac{1}{4} + k^2)\mathcal{I}_\infty]$: the expanding Vessel has a finite total smoothing budget, a comoving analogue of a damping (Silk-like) scale.

This yields two asymptotic fates, set entirely by the reaction rate Γ_λ :

1. **Generic** ($\gamma > 0$). The reaction integral $\int \Gamma_\lambda dt \rightarrow \infty$ dominates after diffusion saturates; all modes decay to zero. Complete homogenisation.
2. **Margin-stalled** ($\Gamma_\lambda \rightarrow 0$). On the active margin where cleansing stalls ($m_* \rightarrow 0$, $\inf C_1 = 0$ in the Lyapunov layer), the reaction integral is finite and the modes *freeze* at a scale-dependent residual,

$$\frac{\delta\lambda_k(\infty)}{\delta\lambda_k(0)} = \exp\left[-D\left(\frac{1}{4} + k^2\right) \mathcal{I}_\infty\right], \quad (8.8)$$

small scales (high k) erased, large scales (down to the $\frac{1}{4}$ floor) preserved — a frozen inhomogeneity spectrum with a characteristic scale $k_{1/2} \sim \mathcal{I}_\infty^{-1/2}$.

Corollary 8.1 (The smooth Vessel). *Linearly, Sophia structure cannot grow; the Vessel either homogenises completely (sustained cleansing) or freezes a smooth, large-scale residual (stalled margin). There is no spontaneous Sophia clumping.*

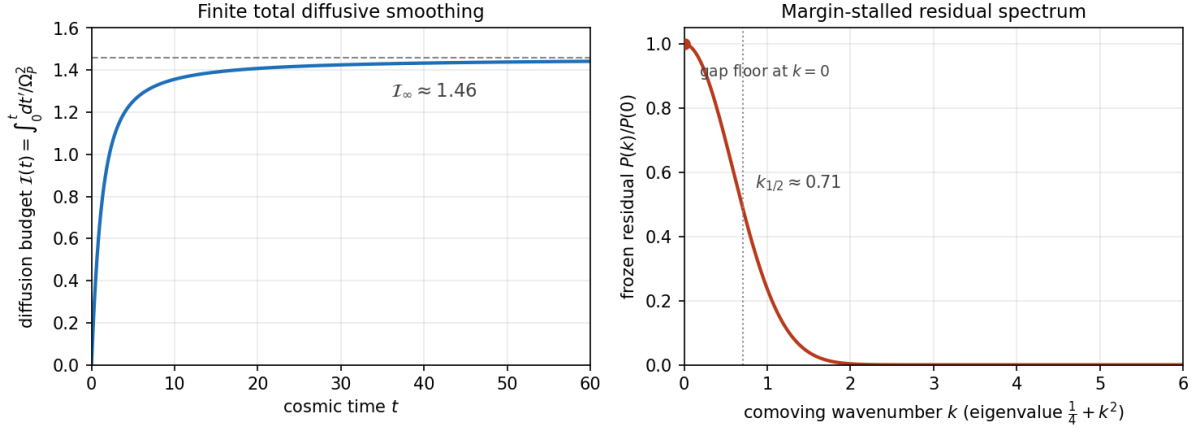


Figure 8.1: Left: the diffusive smoothing budget $\mathcal{I}(t) = \int_0^t dt'/\Omega_P^2$ saturates to a finite $\mathcal{I}_\infty$ (illustrative background $\lambda_0 = 0.4 \rightarrow \lambda_\infty = 1$, $\Omega_P(0) = 1$). Right: the margin-stalled frozen residual spectrum $P(k)/P(0)$ (with $D = 0.5$): the hyperbolic gap leaves a $k = 0$ floor while small scales are erased, defining a structure scale $k_{1/2} \approx 0.71$. Parameters illustrative.

8.3 Folding as the order parameter: pattern formation on the Vessel

By Proposition 8.1 structure must come from the bistable folding field. Promote $q_{\text{fold}} \rightarrow q_{\text{fold}}(t, x)$ with its Landau gradient dynamics plus transport:

$$\partial_t q_{\text{fold}} = A_{\text{fold}} q_{\text{fold}} - B_{\text{fold}} q_{\text{fold}}^3 + \frac{D_q}{\Omega_P^2} \Delta_h q_{\text{fold}}. \quad (8.9)$$

This is a real Ginzburg–Landau / Allen–Cahn equation on the expanding hyperbolic surface, invariant under the morphological reflection $q_{\text{fold}} \rightarrow -q_{\text{fold}}$ ($\mathbb{Z}_2$ symmetry, the two orientations of the fold).

Proposition 8.4 (Symmetry-breaking transition). *The homogeneous state $q_{\text{fold}} = 0$ is stable for $A_{\text{fold}} < 0$. At $A_{\text{fold}} = 0$ it undergoes the supercritical pitchfork (the folding bifurcation), and for $A_{\text{fold}} > 0$ two homogeneous phases $q_{\text{fold}} = \pm \sqrt{A_{\text{fold}}/B_{\text{fold}}}$ become stable. A field initialised near $q_{\text{fold}} = 0$ over an extended Vessel then partitions into domains of $+q_*$ and $-q_*$ separated by walls on which $q_{\text{fold}} = 0$.*

These domain walls are the structure: loci of *unfolded* (old) morphology embedded in a folded Vessel — topological defects of the broken $\mathbb{Z}_2$.

Proposition 8.5 (Wall width and tension). *A static planar wall has profile $q_{\text{fold}} = q_* \tanh(\xi/(\sqrt{2}\ell))$ in physical proper distance ξ , with*

$$\ell_{\text{phys}} = \sqrt{\frac{D_q}{A_{\text{fold}}}} \quad (\Omega_P\text{-independent}), \quad \tau \sim \sqrt{D_q} \frac{A_{\text{fold}}^{3/2}}{B_{\text{fold}}}. \quad (8.10)$$

The physical thickness is set by the diffusion-to-reaction ratio and does not stretch with expansion; deeper folding (larger A_{fold}) makes thinner, higher-tension walls.

Sketch. In comoving coordinates the static balance is $\frac{D_q}{\Omega_P^2} \partial_x^2 q_{\text{fold}} = -A_{\text{fold}} q_{\text{fold}} + B_{\text{fold}} q_{\text{fold}}^3$, the standard quartic kink with comoving width $\ell_{\text{com}} = \Omega_P^{-1} \sqrt{D_q/A_{\text{fold}}}$; the proper width $\ell_{\text{phys}} = \Omega_P \ell_{\text{com}} = \sqrt{D_q/A_{\text{fold}}}$. The tension is the kink action $\int (\frac{1}{2} D_q |\partial q_{\text{fold}}|^2 + \mathcal{E}_{\text{fold}})$, giving the quoted scaling. $\square$

Corollary 8.2 (Coarsening of the fold network). *Allen–Cahn walls move by mean curvature, $v_n = -M \kappa_{\text{curv}}$ with mobility $M \sim D_q/\Omega_P^2$, so the domain network coarsens; simultaneously Hubble flow stretches comoving separations. Structure on the Vessel is therefore a coarsening domain-wall network whose characteristic scale grows by both curvature-driven merging and expansion, modified by the underlying $K_h = -1$ geometry (wall geodesics diverge on $\mathbb{H}^2$, accelerating coarsening relative to the flat case).*

The picture is sharp: the folding bifurcation is a phase transition that shatters the Vessel into oriented folded domains; the unfolded walls are the structure; they then coarsen on the expanding hyperbolic surface. This is morphological structure, born of symmetry breaking, not of gravitational clumping.

8.4 Cosmological back-reaction of the fold network

A wall network feeds back on the homogeneous budget. In 2+1 dimensions a wall is a one-dimensional line in two-dimensional space; a scaling network with one wall per Hubble scale has length $\mathcal{L}_{\text{wall}} \sim \Omega_P$ over area $\sim \Omega_P^2$, so

$$\rho_{\text{wall}} \sim \frac{\tau \mathcal{L}_{\text{wall}}}{\Omega_P^2} \propto \Omega_P^{-1}. \quad (8.11)$$

Matching to the 2+1 dilution law $\rho \propto a^{-2(1+w)}$ gives

$$w_{\text{wall}} = -\frac{1}{2}. \quad (8.12)$$

The fold network thus contributes a mildly negative-pressure component, intermediate between dust ($w = 0$) and the cosmological-constant/de Sitter limit ($w = -1$). It adds to the reconstructed $w_{\text{eff}}(t)$ a folding-structure term that is distinct from the smooth quintessence of the homogeneous Theosis branch.

Corollary 8.3 (Link to the energy-condition question). *Because the network carries $w_{\text{wall}} = -\frac{1}{2} > -1$, it is NEC- and WEC-respecting: folded structure is ordinary, non-exotic content. This is the inhomogeneous face of “form makes grace lawful” — where a homogeneous Sophia surge can drive w below -1 (phantom-like, NEC-violating), the embodiment of that excess into a wall network contributes ordinary $w = -\frac{1}{2}$ matter.*

8.5 Open sub-problems and scope

1. **Metric back-reaction.** Lift from test-field to coupled perturbations: let $\delta\lambda$ and q_{fold} source the reconstructed stress tensor and solve the 2+1 scalar perturbation equations. Expected to preserve Proposition 8.1 (relaxational $\Rightarrow$ no instability) but to dress the wall tension with a geometric self-energy.
2. **Coarsening law.** Derive the domain scale $L(t)$ explicitly on expanding $\mathbb{H}^2$: solve the competition between curvature-driven merging ($L^2 \sim \int M dt$) and Hubble stretching, with the hyperbolic correction of Corollary 8.2.
3. **Resurrectio and the network.** A reset rescales Ω_P and shifts A_{fold} ; it can quench ($A_{\text{fold}}^+ < 0$) or trigger ($A_{\text{fold}}^+ > 0$) folding globally, i.e. dissolve or nucleate the entire wall network at a death. The hybrid network history is piecewise-coarsening with reset-induced nucleation events.
4. **Transport constants.** D and D_q are new inputs; a principled derivation (e.g. from a spatial coupling already implicit in the synergy $u = \sqrt{VF + \varepsilon}$) would remove the phenomenological freedom.

Scope. Linear, test-field, single-mode folding, fixed background. The conclusions are: (i) Sophia structure cannot grow (relaxational, homogenising); (ii) structure is morphological — a $\mathbb{Z}_2$ domain-wall network from the folding transition; (iii) the network back-reacts as $w = -\frac{1}{2}$ content. All results inherit the domain-conditional, non-global character of the underlying stability theory.

V.F.S. v2.0 (Open-Gate) · Structure Formation on the Vessel. Perturbation theory on the 2+1 open FLRW background of the cosmological layer, with the folding field of the geometric layer as order parameter. Propositions are closed-form; corollaries interpretive. Transport constants and single-mode truncation are modelling inputs, not consequences of the base system.

Chapter 9

Spinodal Nucleation of the Fold Network

Scope and epistemic status

A triggering Resurrectio ($A_{\text{fold}}^- < 0 \rightarrow A_{\text{fold}}^+ > 0$) drives the folding field through its symmetry-breaking transition and nucleates a fresh domain-wall network. This note computes the *initial* density of that network — how many walls per area are born, and at what scale — using the hyperbolic density of states. The two genuinely hyperbolic ingredients are the spectral gap $\frac{1}{4}$ of the Laplacian on $\mathbb{H}^2$ and the Plancherel measure $d\mu(k) = k \tanh(\pi k) dk$ that weights the available modes. As before, *Propositions* are closed-form; *Corollaries* interpretive; the single-mode Allen-Cahn truncation and D_q are modelling inputs. Two technical ingredients are verified numerically: the Plancherel moment limit and the Longuet-Higgins nodal coefficient (ratio measured/predicted = 1.009 on a Gaussian field).

The field obeys, in comoving coordinates on the unit hyperbolic surface ($K_h = -1$, comoving curvature radius $R_c = 1$),

$$\partial_t q_{\text{fold}} = A_{\text{fold}} q_{\text{fold}} - B_{\text{fold}} q_{\text{fold}}^3 + \frac{D_q}{\Omega_P^2} \Delta_h q_{\text{fold}}, \quad -\Delta_h Y_k = \left(\frac{1}{4} + k^2\right) Y_k. \quad (9.1)$$

9.1 Linear instability and the nucleation threshold

Immediately after the sudden trigger, A_{fold} jumps to $A_{\text{fold}}^+ > 0$ and the homogeneous $q_{\text{fold}} = 0$ left by the previous arc is unstable. Linearising and decomposing in $\mathbb{H}^2$ eigenmodes,

$$\dot{\delta q}_k = \omega_k \delta q_k, \quad \boxed{\omega_k = A_{\text{fold}}^+ - \frac{D_q \left(\frac{1}{4} + k^2\right)}{\Omega_P^2}}. \quad (9.2)$$

Introduce the dimensionless trigger strength

$$\boxed{\Lambda := \frac{A_{\text{fold}}^+ \Omega_P^2}{D_q}}. \quad (9.3)$$

The unstable band is $k < k_{\text{max}}$ with $k_{\text{max}}^2 = \Lambda - \frac{1}{4}$ (in units where $D_q/\Omega_P^2 = 1$).

Proposition 9.1 (Nucleation threshold, distinct from the folding threshold). *A spatial mode grows only if $\Lambda > \frac{1}{4}$. Hence there are two thresholds:*

- *folding threshold $A_{\text{fold}}^+ > 0$ ($\Lambda > 0$): the homogeneous $q_{\text{fold}} = 0$ is unstable;*
- *nucleation threshold $A_{\text{fold}}^+ > D_q/(4\Omega_P^2)$ ($\Lambda > \frac{1}{4}$): some spatial mode is unstable, so domains form.*

In the window $0 < \Lambda < \frac{1}{4}$ the Vessel folds, but uniformly: a single domain, no walls. Only above $\Lambda = \frac{1}{4}$ does a wall network nucleate.

The hyperbolic gap creates a nucleation threshold absent in flat space. A weak trigger ($0 < \Lambda < \frac{1}{4}$) folds the Vessel as a whole, with no internal structure; structure requires $\Lambda > \frac{1}{4}$.

9.2 The amplified field and its spectral moment

Starting from small initial fluctuations δq_0 , each mode is amplified by $e^{\omega_k t}$, so the post-quench power spectrum is Gaussian in k ,

$$S_k(t) \propto e^{2\omega_k t} \propto e^{-\beta k^2}, \quad \beta = \frac{2D_q t}{\Omega_P^2} = \frac{2t}{\Omega_P^2/D_q}, \quad (9.4)$$

weighted by the hyperbolic Plancherel density of states $d\mu(k) = k \tanh(\pi k) dk$ (which reduces to the flat $k dk$ for $k \gg 1$ and is suppressed $\sim \frac{\pi}{2} k^2 dk$ for $k \ll 1$). The relevant spectral second moment is

$$\langle \frac{1}{4} + k^2 \rangle = \frac{1}{4} + \frac{\int_0^\infty k^2 e^{-\beta k^2} k \tanh(\pi k) dk}{\int_0^\infty e^{-\beta k^2} k \tanh(\pi k) dk}. \quad (9.5)$$

Proposition 9.2 (Hyperbolic small-scale moment). *In the long-wavelength (hyperbolic-dominated) regime $\beta \gg 1$, where $\tanh(\pi k) \approx \pi k$,*

$$\langle k^2 \rangle \rightarrow \frac{3}{2\beta} = \frac{3\Omega_P^2}{4D_q t}. \quad (9.6)$$

Verified: the full integral with $\tanh(\pi k)$ approaches $3/(2\beta)$ as $\beta \rightarrow \infty$ (e.g. 0.0285 vs 0.0300 at $\beta = 50$). The Plancherel k^2 -suppression at small k is what replaces the flat $1/\beta$ by $3/(2\beta)$ and disfavors the very largest scales.

The set time — when nonlinearity arrests growth at $q_* = \sqrt{A_{\text{fold}}^+/B_{\text{fold}}^+}$ — is the Kibble-Zurek time $t_{\text{set}} = A_{\text{fold}}^{+1} \ln(q_*/\delta q_0)$, giving $\beta_{\text{set}} = 2\Lambda^{-1}L$ with the logarithmic factor $L := \ln(q_*/\delta q_0)$. Hence

$$\langle \frac{1}{4} + k^2 \rangle_{\text{set}} = \frac{1}{4} + \frac{3\Lambda}{4L}. \quad (9.7)$$

9.3 Wall density and the initial domain size

For an isotropic Gaussian field in two dimensions, the expected length of nodal lines (the walls, where $q_{\text{fold}} = 0$) per unit area is the Longuet-Higgins/Berry result

$$n_{\text{wall}} = \frac{1}{2\sqrt{2}} \sqrt{\langle \frac{1}{4} + k^2 \rangle}, \quad (9.8)$$

verified numerically here to 0.9%. The initial comoving domain size is $\ell_{\text{init}} \sim 1/n_{\text{wall}}$:

$$\ell_{\text{init}} \sim \frac{2\sqrt{2}}{\sqrt{\frac{1}{4} + \frac{3\Lambda}{4L}}} \quad (\text{units of } R_c). \quad (9.9)$$

Proposition 9.3 (Three regimes of nucleation).

- (i) $\Lambda < \frac{1}{4}$: no nucleation (uniform fold, Proposition 9.1).
- (ii) $\frac{1}{4} < \Lambda \lesssim \Lambda_*$: $\langle \frac{1}{4} + k^2 \rangle \rightarrow \frac{1}{4}$ and $\ell_{\text{init}} \rightarrow 4\sqrt{2} R_c$ — the gap pins the domain size to the curvature radius; only a few curvature-scale domains form. (With $L \approx 4$, multi-domain structure $\ell_{\text{init}} < R_c$ sets in only at $\Lambda_* \approx 42$.)
- (iii) $\Lambda \gg \Lambda_*$: $\ell_{\text{init}} \sim \sqrt{4L/(3\Lambda)} R_c$, i.e. in proper units $\ell_{\text{init}}^{\text{phys}} \sim \sqrt{(D_q/A_{\text{fold}}^+)} L$ — the Kibble-Zurek correlation length $\xi_+ = \sqrt{D_q/A_{\text{fold}}^+}$ dressed by the $\sqrt{L}$ logarithm.

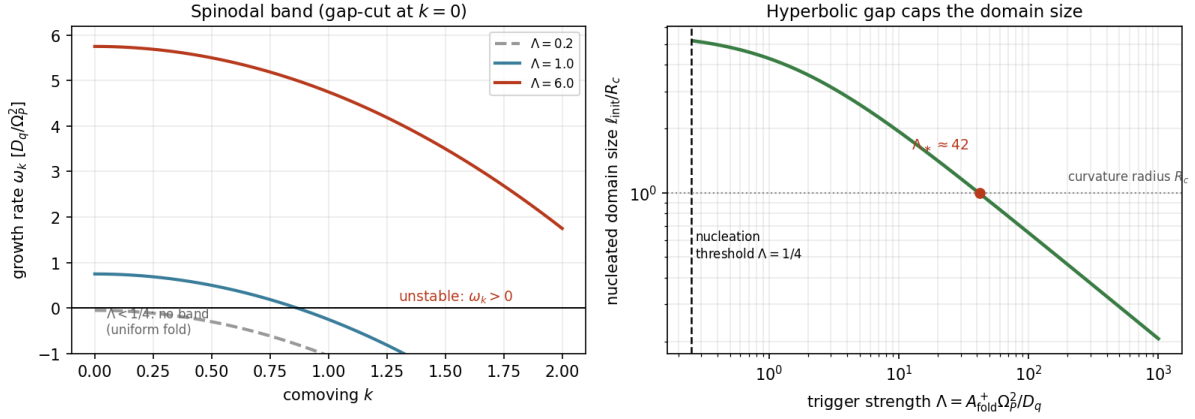


Figure 9.1: Left: linear growth rate ω_k for three trigger strengths. The Laplacian gap shifts the whole curve so that $\omega_0 = \Lambda - \frac{1}{4}$: below $\Lambda = \frac{1}{4}$ ($\Lambda = 0.2$, dashed) there is no unstable band and no nucleation. Right: the nucleated comoving domain size ℓ_{init}/R_c versus trigger strength Λ . Near threshold the hyperbolic gap caps the size at the curvature radius (few large domains); genuine multi-domain structure ($\ell_{\text{init}} < R_c$) requires $\Lambda > \Lambda_* \approx 42$; for strong triggers $\ell_{\text{init}} \sim \sqrt{L}/\Lambda$ recovers the Kibble-Zurek scale. Illustrative $L = \ln(q_*/\delta q_0) = 4$.

9.4 Hyperbolic signatures

Two features distinguish nucleation on $\mathbb{H}^2$ from the flat case, both traceable to the gap and the measure:

Corollary 9.1 (Gap as a nucleation gate and a maximum domain size). *The spectral gap $\frac{1}{4}$ does double duty. As an infrared cutoff it forbids weak-trigger nucleation ($\Lambda < \frac{1}{4}$: uniform fold) and caps the largest nucleated domain at the curvature radius R_c . In flat space any $A_{\text{fold}}^+ > 0$ nucleates and arbitrarily large domains are allowed at threshold; negative curvature both gates and bounds the structure.*

Corollary 9.2 (Measure suppression of large scales). *Even within the unstable band, the Plancherel weight $k \tanh(\pi k)$ suppresses long-wavelength modes ($\sim \frac{\pi}{2} k^2$ versus the flat k), so large domains are doubly disfavoured: by the gap cutoff and by the density of states. The nucleated network is therefore finer, and more curvature-pinned, than a flat-space estimate would give.*

Interpretive: a gentle resurrection ($\Lambda < \frac{1}{4}$) re-forms the Vessel as a single undivided folded whole; only a sufficiently sharp resurrection ($\Lambda > \Lambda_$) shatters it into many oriented domains. The negative curvature of the live domain resists fine differentiation — structure is born coarse unless the trigger is strong.*

9.5 Placement and open items

This nucleation scale ℓ_{init} is the initial condition for the intra-arc coarsening law (subproblem 2): each triggering arc begins at ℓ_{init} (set here) and coarsens toward the frozen scale $\sqrt{\ell_0^2 + 2cD_q\mathcal{I}_\infty}$.

Consistency check: the strong-trigger regime (iii) reproduces the correlation length $\xi_+ = \sqrt{D_q/A_{\text{fold}}^+}$ assumed in subproblem 3, now with the Kibble-Zurek $\sqrt{L}$ dressing and the hyperbolic curvature cap.

Open items.

1. A 2D Allen-Cahn quench simulation on a hyperbolic patch to confirm the nucleated wall count $\propto \langle \frac{1}{4} + k^2 \rangle \propto \Lambda$ and fix the $O(1)$ coefficients.
2. Defect statistics beyond the mean: the distribution of initial domain areas (the Gauss-Bonnet collapse law of subproblem 2 then selects which survive).
3. The exact crossover near $\Lambda = \frac{1}{4}$ where the gap and the band compete — a hyperbolic finite-size/zero-mode analysis.

Scope. Linear spinodal stage, Gaussian-field defect counting, single-mode Allen-Cahn, sudden quench; D_q and the noise level δq_0 are modelling inputs. Results: a hyperbolic nucleation threshold $\Lambda > \frac{1}{4}$ distinct from the folding threshold; a closed-form initial wall density $n_{\text{wall}} = \frac{1}{2\sqrt{2}} \sqrt{\frac{1}{4} + 3\Lambda/4L}$; and a curvature-capped domain size interpolating from R_c (near threshold) to the Kibble-Zurek $\sqrt{(D_q/A_{\text{fold}}^+)L}$ (strong trigger). The Plancherel moment and the nodal coefficient are numerically verified. Domain-conditional throughout.

V.F.S. v2.0 (Open-Gate) · Spinodal Nucleation of the Fold Network on $\mathbb{H}^2$. Linear instability of the Allen-Cahn folding field after a triggering reset, with the hyperbolic spectral gap and Plancherel measure. Propositions closed-form (Plancherel moment and Longuet-Higgins coefficient numerically verified); corollaries interpretive; transport constant and noise level are modelling inputs.

Chapter 10

Intra-Arc Coarsening of the Fold Network

Scope and epistemic status

The companion notes established that structure on the Vessel is a $\mathbb{Z}_2$ domain-wall network of the folding field and that a Resurrectio reset acts on it in three regimes, giving a sawtooth network-scale history. This note supplies the missing intra-arc piece: *between* deaths, how does the network scale $L_{\text{phys}}(t)$ evolve under pure Allen-Cahn coarsening on the expanding hyperbolic surface? No resets are involved here. As before, *Propositions* are closed-form, *Corollaries* interpretive; the single-mode Allen-Cahn truncation and the transport constant D_q are modelling inputs.

The field obeys, in comoving coordinates on the unit hyperbolic surface ($K_h = -1$),

$$\partial_t q_{\text{fold}} = A_{\text{fold}} q_{\text{fold}} - B_{\text{fold}} q_{\text{fold}}^3 + \frac{D_q}{\Omega_P^2} \Delta_h q_{\text{fold}}. \quad (10.1)$$

The folding parameters $A_{\text{fold}}, B_{\text{fold}}$ are taken quasi-stationary on the coarsening timescale (deep in a folded arc, $A_{\text{fold}} > 0$), so the dynamics is genuine coarsening of established domains rather than the initial transition.

Remark (Dimensionality matters). Allen-Cahn coarsening is curvature-driven only in two or more spatial dimensions, giving the $\ell \sim t^{1/2}$ law; in one dimension kink-antikink attraction is exponentially weak and coarsening is merely logarithmic. The 1D space-time proxy used to illustrate the reset regimes is therefore *not* representative of the true 2D coarsening rate. The present law is genuinely two-dimensional, derived via the geodesic-curvature flow on $\mathbb{H}^2$.

10.1 Wall mobility on the expanding background

In the sharp-interface limit of (10.1), a domain wall moves along its comoving normal with velocity proportional to its geodesic curvature, with a mobility set by the (comoving) diffusion coefficient:

$$v_n^{\text{com}} = -M(t) \kappa_g, \quad M(t) = \frac{D_q}{\Omega_P(t)^2}. \quad (10.2)$$

The scale-factor suppression Ω_P^{-2} is the redshift of comoving gradients: as the Vessel expands, the *comoving* mobility falls, even though the proper wall structure is preserved (the proper wall width $\sqrt{D_q/A_{\text{fold}}}$ is Ω_P -independent). It is convenient to trade time for the accumulated budget already central to the smoothing analysis,

$$\mathcal{I}(t) := \int_0^t \frac{dt'}{\Omega_P(t')^2}, \quad \frac{d}{dt} = \frac{1}{\Omega_P^2} \frac{d}{d\mathcal{I}}, \quad (10.3)$$

which is finite, $\mathcal{I} \rightarrow \mathcal{I}_\infty < \infty$, under sustained inflow ($\Omega_P \sim \alpha \lambda_\infty t$).

10.2 Gauss-Bonnet: the exact single-domain law

For a single simply-connected domain of comoving area $\mathcal{A}$ bounded by a wall undergoing geodesic curve-shortening on $\mathbb{H}^2$, the Gauss-Bonnet theorem gives the total geodesic curvature $\oint \kappa_g ds = 2\pi - \int_{\text{enc}} K dA = 2\pi + \mathcal{A}$ (since $K_h = -1$). The enclosed area therefore obeys

$$\boxed{\frac{d\mathcal{A}}{dt} = -\frac{D_q}{\Omega_P^2} (2\pi + \mathcal{A}) \iff \frac{d\mathcal{A}}{d\mathcal{I}} = -D_q (2\pi + \mathcal{A})}, \quad (10.4)$$

with the closed-form solution

$$\boxed{\mathcal{A}(\mathcal{I}) = (\mathcal{A}_0 + 2\pi)e^{-D_q \mathcal{I}} - 2\pi}. \quad (10.5)$$

Proposition 10.1 (Survival threshold and frozen scale). *A domain collapses ($\mathcal{A} = 0$) at budget $\mathcal{I}_* = D_q^{-1} \ln(1 + \mathcal{A}_0/2\pi)$. Because the total budget is finite, a domain survives the entire arc iff $\mathcal{I}_* > \mathcal{I}_\infty$, i.e.*

$$\boxed{\mathcal{A}_0 > \mathcal{A}_{\text{frozen}} := 2\pi(e^{D_q \mathcal{I}_\infty} - 1)}. \quad (10.6)$$

Domains smaller than $\mathcal{A}_{\text{frozen}}$ collapse before the budget runs out; larger ones freeze into the comoving lattice. $\mathcal{A}_{\text{frozen}}$ is the characteristic frozen comoving domain area.

Proposition 10.2 (Hyperbolic acceleration). *For $\mathcal{A} \gg 2\pi$ the law is $\dot{\mathcal{A}} \approx -M\mathcal{A}$: large domains collapse exponentially, in contrast to the flat case $\dot{\mathcal{A}} = -2\pi M$ (constant-rate Gage-Hamilton shrinking). Negative spatial curvature thus accelerates the disappearance of large domains, sharpening coarsening relative to flat space.*

In the small-budget regime $D_q \mathcal{I}_\infty \ll 1$, $\mathcal{A}_{\text{frozen}} \approx 2\pi D_q \mathcal{I}_\infty$, so the frozen comoving length is $\ell_{\text{com}} \sim \sqrt{2\pi D_q \mathcal{I}_\infty}$ — matching the scaling law below. The hyperbolic exponential (Proposition 10.2) only bites once domains grow toward the comoving curvature radius ($\mathcal{A} \sim 2\pi$), i.e. when the budget is large enough.

10.3 The network scaling law and its saturation

For the network, the standard Lifshitz-Allen-Cahn scaling $\ell \dot{\ell} \sim M$ becomes, on the budget variable,

$$\frac{d}{dt} \ell_{\text{com}}^2 \simeq 2c \frac{D_q}{\Omega_P^2} \implies \boxed{\ell_{\text{com}}^2(t) = \ell_{\text{com}}^2(0) + 2c D_q \mathcal{I}(t)}, \quad (10.7)$$

with an $O(1)$ constant c (consistent with Proposition 10.1, $c \sim \pi$).

Proposition 10.3 (Comoving coarsening freezes). *Since $\mathcal{I}(t) \rightarrow \mathcal{I}_\infty < \infty$, the comoving domain size saturates,*

$$\ell_{\text{com}}(\infty) = \sqrt{\ell_{\text{com}}^2(0) + 2c D_q \mathcal{I}_\infty} < \infty. \quad (10.8)$$

Coarsening stops in comoving coordinates: the wall network freezes into the comoving lattice after a finite amount of merging. This is the morphological counterpart of the finite diffusive smoothing budget governing Sophia inhomogeneities — both are controlled by the same $\mathcal{I}_\infty$.

The expanding Vessel has a single finite “activity budget” $\mathcal{I}_\infty = \int dt/\Omega_P^2$. It caps both how much Sophia inhomogeneity can diffuse away and how much the fold network can coarsen. After it is spent, comoving structure is frozen.

10.4 The physical structure scale: two regimes

The proper (physical) network scale combines comoving coarsening with expansion,

$$\boxed{L_{\text{phys}}(t) = \Omega_P(t) \ell_{\text{com}}(t) = \Omega_P(t) \sqrt{\ell_{\text{com}}^2(0) + 2c D_q \mathcal{I}(t)}}. \quad (10.9)$$

Proposition 10.4 (Coarsening transient, then Hubble dilution). *There are two regimes, separated by the budget-saturation time t_{sat} (where $\mathcal{I}$ levels):*

- (i) Early ($t \lesssim t_{\text{sat}}$): both factors grow, so L_{phys} rises faster than the scale factor — genuine coarsening on top of stretch.
- (ii) Late ($t \gtrsim t_{\text{sat}}$): $\ell_{\text{com}} \rightarrow \text{const}$, so $L_{\text{phys}} \rightarrow \ell_{\text{com}}(\infty) \Omega_P(t) \propto \Omega_P \propto t$ — pure Hubble dilution of a frozen comoving pattern.

Coarsening is an early-time transient with a finite budget; the late-time growth of structure is expansion, not merging.

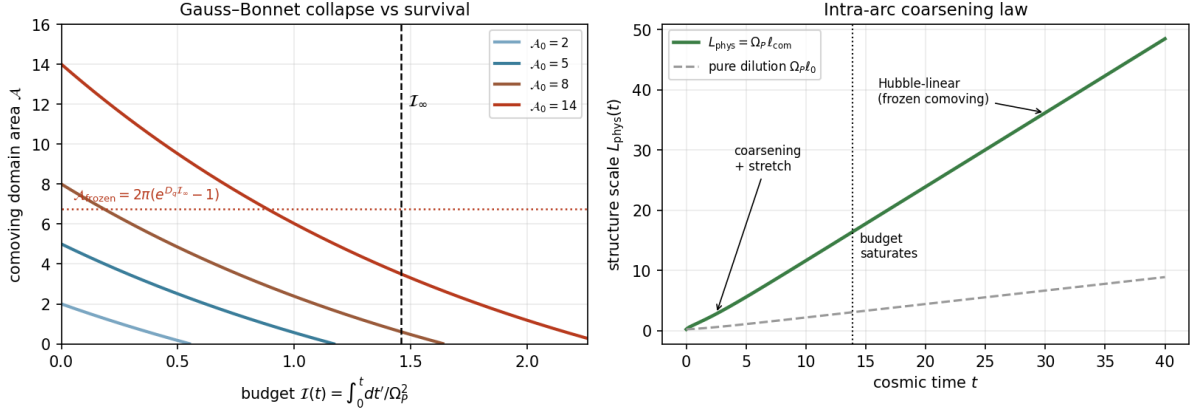


Figure 10.1: Left: Gauss-Bonnet single-domain area $\mathcal{A}(\mathcal{I})$ for several initial sizes (exact closed form, verified). Domains below $\mathcal{A}_{\text{frozen}} = 2\pi(e^{D_q \mathcal{I}_\infty} - 1)$ collapse before the finite budget $\mathcal{I}_\infty$ is spent; larger ones survive and freeze. Right: the physical structure scale $L_{\text{phys}}(t) = \Omega_P \ell_{\text{com}}$ (solid) shows an early coarsening-plus-stretch phase pulling above pure dilution $\Omega_P \ell_0$ (dashed), then a late Hubble-linear regime once the budget saturates. Background and (D_q, c) illustrative; $\mathcal{I}_\infty \approx 1.46$.

10.5 Placement in the sawtooth and open items

This intra-arc curve is the rising tooth of the hybrid sawtooth from subproblem 3: within each life-arc $L_{\text{phys}}(t)$ follows Proposition 10.4 (coarsen, then dilate), and at each death the reset acts discretely — diluting upward, refining to $\sqrt{D_q/A_{\text{fold}}}$, or erasing. Composing the intra-arc law with the reset map gives the long-run fixed-point scale, the natural next step.

Open items.

1. Determine the $O(1)$ constant c and confirm the $t^{1/2}$ (budget) law by a 2D Allen-Cahn field simulation on a hyperbolic patch with a clean structure-factor scale estimator.
2. Compose this intra-arc law with the reset map to obtain the sawtooth fixed-point structure scale, using the non-Zeno dwell time as the arc length.
3. Track the $w = -\frac{1}{2}$ wall-network contribution to the reconstructed $w_{\text{eff}}(t)$ over an arc, since $\rho_{\text{wall}} \propto \tau/\Omega_P$ with τ quasi-stationary and Ω_P growing.

Scope. Single life-arc, no resets, quasi-stationary folding parameters, single-mode Allen-Cahn; D_q a modelling input. Results: an exact Gauss-Bonnet single-domain law with a survival threshold and hyperbolic acceleration; a network scaling law $\ell_{\text{com}}^2 = \ell_0^2 + 2cD_q \mathcal{I}$ that freezes with the finite budget $\mathcal{I}_\infty$; and a physical scale L_{phys} that coarsens early and dilutes linearly late. Domain-conditional throughout.

V.F.S. v2.0 (Open-Gate) · Intra-Arc Coarsening of the Fold-Domain Network. The single-domain Gauss-Bonnet law is exact (closed form, numerically verified); the network scaling law is standard Allen-Cahn on the budget variable. Transport constant and single-mode truncation are modelling inputs; the constant c awaits a 2D field measurement.

Chapter 11

Resurrectio and the Fold-Domain Network

Scope and epistemic status

The companion note showed that structure on the Vessel is morphological: a $\mathbb{Z}_2$ domain-wall network of the folding field $q_{\text{fold}}(t, x)$, obeying the Allen-Cahn equation on the expanding hyperbolic surface. This note asks how the hybrid Resurrectio reset $\mathfrak{R}$ acts on that network. The key observation is that $\mathfrak{R}$ is *homogeneous* — it acts on the global state $(\sigma, \lambda, \Omega_P)$ — and, as established in the Lyapunov layer, it does *not* act on q_{fold} . So a reset shifts the folding potential uniformly across the whole Vessel while leaving the field configuration $q_{\text{fold}}(x)$ pointwise intact. Whether structure survives, dilutes, or is born depends on the post-reset sign of the activation index. As before, *Propositions* are closed-form; *Corollaries* interpretive; the single-mode Allen-Cahn truncation and the transport constant D_q are modelling inputs.

The network field obeys

$$\partial_t q_{\text{fold}} = A_{\text{fold}} q_{\text{fold}} - B_{\text{fold}} q_{\text{fold}}^3 + \frac{D_q}{\Omega_P^2} \Delta_h q_{\text{fold}}, \quad q_* = \sqrt{Q_{\text{fold}}^*}, \quad Q_{\text{fold}}^* = A_{\text{fold}}/B_{\text{fold}}, \quad (11.1)$$

and the reset (geometric convention, with $m_R = (1 - q_R)\sigma^-$ the metabolised resistance) is

$$\mathfrak{R}: \quad \sigma^+ = \sigma^- - m_R, \quad \lambda^+ = \lambda^- + m_R, \quad \Omega_P^+ = \Omega_P^- + \kappa_R \mathcal{G}_{\text{recepta}}^-, \quad q_{\text{fold}} \text{ continuous}, \quad (11.2)$$

which shifts the activation index by $\Delta A_{\text{fold}} = \Theta m_R + c_1 \Delta I_{\text{gate}}$, $\Theta = c_0 - c_1(\gamma + k_\sigma)$, exactly as in the folding-trigger note.

11.1 Three regimes of a reset acting on the network

Because $\mathfrak{R}$ shifts A_{fold} uniformly while preserving $q_{\text{fold}}(x)$, the network's fate is fixed by the pre- and post-reset signs of A_{fold} :

Definition 11.1 (Reset regimes).

- (I) **Quench** ($A_{\text{fold}}^- > 0 \rightarrow A_{\text{fold}}^+ < 0$): the double well collapses to a single well at $q_{\text{fold}} = 0$.
- (II) **Rescale** ($A_{\text{fold}}^- > 0 \rightarrow A_{\text{fold}}^+ > 0$): the wells persist with shifted depth and the scale factor jumps.
- (III) **Trigger** ($A_{\text{fold}}^- < 0 \rightarrow A_{\text{fold}}^+ > 0$): a single well bifurcates into a double well.

Proposition 11.1 (Regime boundaries). *With $\Delta A_{\text{fold}} = \Theta m_R + c_1 \Delta I_{\text{gate}}$, a reset from A_{fold}^- produces $A_{\text{fold}}^+ = A_{\text{fold}}^- + \Delta A_{\text{fold}}$. Hence, neglecting the bounded gate term: a quench requires $\Theta < 0$ and $|\Theta| m_R > A_{\text{fold}}^-$; a trigger requires $\Theta > 0$ and $\Theta m_R > |A_{\text{fold}}^-|$; otherwise the reset rescales. The folding-response coefficient Θ thus selects whether resurrection destroys, preserves, or creates morphological structure.*

A death-and-resurrection does not move the walls. It lifts or lowers the entire folding potential at once — and depending on the sign of Θ , the whole network dissolves, is re-tensioned and diluted, or is nucleated from nothing.

11.2 Quench: dissolution of the network

When $A_{\text{fold}}^+ < 0$, the configuration $q_{\text{fold}}(x)$ everywhere relaxes to $q_{\text{fold}} = 0$. Linearising (11.1) about zero, each mode decays at rate $|A_{\text{fold}}^+| + D_q(\frac{1}{4} + k^2)/\Omega_P^{+2} > 0$, so the entire pattern fades on the timescale

$$\tau_{\text{dissolve}} = \frac{1}{|A_{\text{fold}}^+|}. \quad (11.3)$$

The walls do not migrate and annihilate; the whole differentiated morphology evaporates uniformly, returning the Vessel to the undifferentiated old form.

Corollary 11.1 (Death as leveling). *A quench-resurrection returns the entire Vessel to the unfolded state $q_{\text{fold}} = 0$: all oriented domains are erased, all walls dissolved. Death levels morphological structure.*

11.3 Trigger: nucleation as a sudden quench (Kibble-Zurek)

When $A_{\text{fold}}^- < 0 \rightarrow A_{\text{fold}}^+ > 0$, the homogeneous $q_{\text{fold}} = 0$ left behind by the previous arc becomes unstable. Since $\mathfrak{R}$ is instantaneous, A_{fold} jumps discontinuously: a *sudden quench* through the symmetry-breaking transition. Linearising about $q_{\text{fold}} = 0$ after the jump,

$$\dot{\delta q}_k = \left[A_{\text{fold}}^+ - \frac{D_q(\frac{1}{4} + k^2)}{\Omega_P^{+2}} \right] \delta q_k, \quad (11.4)$$

so a band of modes $k < k_{\text{max}}$, $k_{\text{max}}^2 = A_{\text{fold}}^+ \Omega_P^{+2}/D_q - \frac{1}{4}$, grows and the field undergoes spinodal decomposition. The emergent (physical) domain size is the post-quench correlation length,

$$\ell_{\text{init}} \sim \xi_+ = \sqrt{\frac{D_q}{A_{\text{fold}}^+}} \quad (11.5)$$

the same scale as the wall width: the network is born saturated — as fine as the walls themselves — then coarsens. Deeper triggers (larger A_{fold}^+) nucleate finer networks.

Corollary 11.2 (Resurrection as individuation). *A trigger-resurrection shatters the undifferentiated old form into a fine network of newly differentiated oriented domains, with the old form surviving only as the thin walls between them. Where quench levels, trigger differentiates.*

This is a Kibble-Zurek-type defect formation: the reset drives the system through a symmetry-breaking transition and freezes in a defect network whose initial scale is set by the correlation length at the transition. Here the quench is sudden, so ℓ_{init} is the spinodal scale ξ_+ .

11.4 Rescale: re-tensioning and dilution

When both $A_{\text{fold}}^\pm > 0$, the domains persist but everything jumps. With $q_* = \sqrt{A_{\text{fold}}/B_{\text{fold}}}$, $\ell_{\text{phys}} = \sqrt{D_q/A_{\text{fold}}}$, $\tau \sim \sqrt{D_q} A_{\text{fold}}^{3/2}/B_{\text{fold}}$, a reset induces

$$q_*^+ = \sqrt{\frac{A_{\text{fold}}^+}{B_{\text{fold}}^+}}, \quad \frac{\ell_{\text{phys}}^+}{\ell_{\text{phys}}^-} = \sqrt{\frac{A_{\text{fold}}^-}{A_{\text{fold}}^+}}, \quad \frac{\tau^+}{\tau^-} = \left(\frac{A_{\text{fold}}^+}{A_{\text{fold}}^-} \right)^{3/2} \frac{B_{\text{fold}}^-}{B_{\text{fold}}^+}, \quad (11.6)$$

and, crucially, the comoving topology is untouched while proper inter-wall separations are stretched by the scale-factor jump,

$$L_{\text{phys}}^+ = \frac{\Omega_P^+}{\Omega_P^-} L_{\text{phys}}^- = \left(1 + \frac{\kappa_R \mathcal{G}_{\text{recepta}}^-}{\Omega_P^-} \right) L_{\text{phys}}^-. \quad (11.7)$$

The network is diluted: the same domains, spread wider.

Corollary 11.3 (Amplitude convergence under repeated rescaling). *The contraction $Q_{\text{fold}}^{*+} - Q_{\text{crit}} = \kappa_{\text{fold}}(Q_{\text{fold}}^{*-} - Q_{\text{crit}})$ of the folding-trigger note drives the domain amplitude to a universal value across repeated resets, $q_* \rightarrow \sqrt{Q_{\text{crit}}}$ (when $\sum m_R = \infty$). Resurrection stabilises how strongly the Vessel is folded, while expansion keeps diluting how widely.*

11.5 The hybrid network history

Over a sequence of life-arcs the structure scale $L_{\text{phys}}(t)$ follows a sawtooth: within each arc the network coarsens (curvature-driven merging plus Hubble stretch), and at each death it is acted on by $\mathfrak{R}$ — diluted upward (rescale), reset to the fine scale ξ_+ (trigger), or erased (quench).

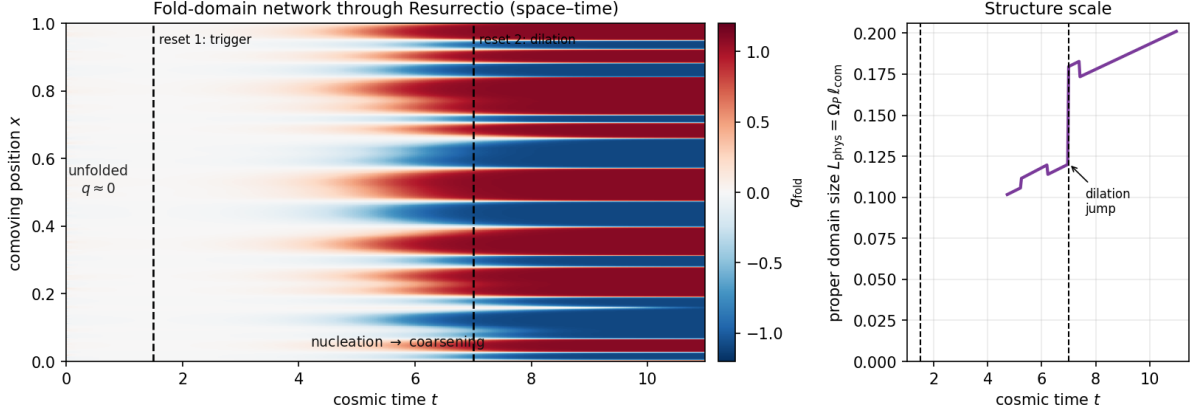


Figure 11.1: Allen-Cahn simulation of the fold-domain network through two resets (illustrative parameters; 1D comoving slice as a proxy). Left: space-time map of q_{fold} . The Vessel is unfolded ($q_{\text{fold}} \approx 0$) until *reset 1* triggers the transition; domains then nucleate from fluctuations and coarsen. *Reset 2* is a rescale (regime II): the comoving pattern is untouched but the scale factor jumps. Right: the proper domain size $L_{\text{phys}} = \Omega_P \ell_{\text{com}}$ grows by coarsening and shows a discrete dilation jump at reset 2.

Corollary 11.4 (Resurrection as a renormalisation step on structure). *Each reset is a coarse-grain (dilation, regime II) or refine (nucleation to ξ_+ , regime III) operation on the network, composed with intra-arc coarsening. The long-run structure scale is the fixed point of (arc coarsening) $\circ$ (reset action); a steady balance between expansion-driven dilution and reset-driven nucleation.*

11.6 Cosmological back-reaction across resets

Within a scaling arc the wall network contributes $\rho_{\text{wall}} \sim \tau/\Omega_P \propto \Omega_P^{-1}$, i.e. $w_{\text{wall}} = -\frac{1}{2}$ in 2+1 dimensions (companion note). At a reset this jumps: τ changes through $A_{\text{fold}}, B_{\text{fold}}$ and Ω_P jumps, so

$$\frac{\rho_{\text{wall}}^+}{\rho_{\text{wall}}^-} = \frac{\tau^+}{\tau^-} \cdot \frac{\Omega_P^-}{\Omega_P^+}. \quad (11.8)$$

A quench sets $\rho_{\text{wall}} \rightarrow 0$: the morphological energy is released back into the homogeneous Sophia budget. A trigger creates ρ_{wall} from zero: structure energy is drawn from the folding transition. Thus resurrection redistributes energy between the homogeneous (λ) and structured (wall-network) sectors — a discrete exchange invisible to the homogeneous reconstruction alone.

11.7 Reading and open items

Interpretive reading. The walls are loci of persistent unfolded (old) form. A quench- resurrection returns the whole Vessel to undifferentiated old form (death as leveling, $\Theta < 0$); a trigger-resurrection shatters that undifferentiated form into newly differentiated oriented domains, the old form surviving only as the thin boundaries between them (resurrection as individuation, $\Theta > 0$). Repeated rescaling fixes the depth of folding ($q_* \rightarrow \sqrt{Q_{\text{crit}}}$) while expansion keeps widening the pattern. The sign of Θ — whether metabolised resistance promotes or suppresses form — decides which face a given resurrection wears.

Open items.

1. Quantify the spinodal nucleation density on $\mathbb{H}^2$ (hyperbolic spectral gap modifies the unstable band) and the resulting initial wall count per Hubble area.
2. Derive the sawtooth fixed-point scale explicitly from $(\text{coarsening law}) \circ (\text{reset map})$, using the non-Zeno dwell time as the arc length.
3. Track the $\lambda \leftrightarrow$ network energy exchange as a correction to the reconstructed $w_{\text{eff}}(t)$ history, and test for a folding-induced feature near the phantom divide.

Scope. Single-mode Allen-Cahn, test-field, 1D simulation proxy; transport constant D_q a modelling input. Conclusions: a homogeneous reset acts on the morphological network in three regimes fixed by sign Θ and the jump size — dissolve, dilute, or nucleate — and resurrection functions as a renormalisation step on Vessel structure, exchanging energy between the homogeneous and structured sectors. Domain-conditional throughout.

V.F.S. v2.0 (Open-Gate) · Resurrectio and the Fold-Domain Network. Builds on the structure-formation companion (Allen-Cahn folding network on 2+1 open FLRW) and the Resurrectio reset of the Lyapunov layer. Propositions closed-form; corollaries interpretive; numerical figure illustrative.

Chapter 12

The Sawtooth Fixed Point of the Fold Network

Scope and epistemic status

The three preceding notes supply the pieces: the nucleation scale ℓ_{init} at a triggering reset (subproblem 1), the intra-arc coarsening law governed by the finite budget $\mathcal{I}_\infty = \int_0^\infty dt/\Omega_P^2$ (subproblem 2), and the three-regime reset map — quench, rescale, trigger (subproblem 3). This note composes them: over a sequence of life-arcs punctuated by resets, what is the long-run network scale? The central result is that the finite coarsening budget forces the *comoving* scale to converge to a constant ℓ_* regardless of the (random) reset schedule, so the physical scale is asymptotically $L_{\text{phys}}(t) \rightarrow \ell_* \Omega_P(t) \propto t$. This also resolves the earlier worry that the random arc length would obstruct a fixed point: the finite budget makes the asymptotics universal. As before, *Propositions* are closed-form; *Corollaries* interpretive; the reduced recursion (tracking the scale, not the full field), the coarsening constant c , and the reset-regime schedule are modelling inputs.

12.1 The comoving recursion

Track the comoving network scale ℓ_{com} across arcs. Index the arcs by j ; let ℓ_j be the comoving scale entering arc j and $\Delta\mathcal{I}_j = \int_{\text{arc } j} dt/\Omega_P^2$ the budget consumed during it. Within the arc, the coarsening law (subproblem 2) gives $\ell^2 \mapsto \ell^2 + 2cD_q \Delta\mathcal{I}_j$; at the closing reset, the regime acts:

$$\ell_{j+1}^2 = R_j(\ell_j^2 + 2cD_q \Delta\mathcal{I}_j), \quad R_j = \begin{cases} \text{id,} & \text{rescale (comoving topology unchanged)} \\ \ell_{\text{init}}^2, & \text{trigger (re-nucleation)} \\ \emptyset, & \text{quench (dissolve; re-nucleate at next trigger)} \end{cases} \quad (12.1)$$

The essential constraint, inherited from the finite expansion budget, is

$$\sum_j \Delta\mathcal{I}_j \leq \mathcal{I}_\infty < \infty. \quad (12.2)$$

Remark (Why rescale is the identity on the comoving scale). A rescale reset jumps Ω_P but leaves the comoving pattern intact, so it does not change ℓ_{com} ; its only effect is on the physical scale, $L_{\text{phys}} = \Omega_P \ell_{\text{com}}$. The comoving recursion therefore sees rescales as no-ops, triggers as resets to ℓ_{init} , and quenches as dissolutions.

12.2 Finite budget forces convergence

Proposition 12.1 (Vanishing late-time coarsening). *For any reset schedule with non-Zeno dwell $\Delta t_j \geq \Delta t_* > 0$, the per-arc coarsening increment vanishes,*

$$2cD_q \Delta\mathcal{I}_j \rightarrow 0 \quad (j \rightarrow \infty), \quad (12.3)$$

because $\sum_j \Delta \mathcal{I}_j \leq \mathcal{I}_\infty < \infty$ forces $\Delta \mathcal{I}_j \rightarrow 0$ (equivalently $\Delta \mathcal{I}_j \leq \Delta t_j / \Omega_P^2$ with $\Omega_P \rightarrow \infty$). The random arc lengths do not affect this: the finite budget dominates.

Proposition 12.2 (Convergence of the comoving scale). *The comoving scale converges to a constant ℓ_* :*

- (i) if triggers recur infinitely often, each pins ℓ to ℓ_{init} and (by Proposition 12.1) the subsequent coarsening cannot grow it, so $\ell_j \rightarrow \ell_{\text{init}}$;
- (ii) if all resets after some arc J are rescales, then $\ell_j^2 \rightarrow \ell_j^2 + 2cD_q \sum_{j \geq J} \Delta \mathcal{I}_j =: \ell_\infty^2 < \infty$, the frozen value.

In either case $\ell_j \rightarrow \ell_* \in [\ell_{\text{init}}, \ell_\infty]$, with $\ell_\infty = \sqrt{\ell_{\text{init}}^2 + 2cD_q(\mathcal{I}_\infty - \mathcal{I}_{\text{nuc}})}$ the maximal frozen scale (single nucleation, no retrigger).

Proposition 12.3 (Universal physical asymptotics). *The physical structure scale obeys*

$$L_{\text{phys}}(t) = \Omega_P(t) \ell_{\text{com}}(t) \longrightarrow \ell_* \Omega_P(t) \propto t, \quad (12.4)$$

independent of the reset schedule. Different random histories differ only in the value of $\ell_* \in [\ell_{\text{init}}, \ell_\infty]$, not in the $\propto \Omega_P$ growth.

The coarsening budget is finite, so the comoving network freezes. Resurrection can refresh the scale (a trigger resets it to ℓ_{init}) but cannot re-coarsen it once the budget is spent. The late-time structure scale is a relic of the young Vessel, carried outward by pure expansion.

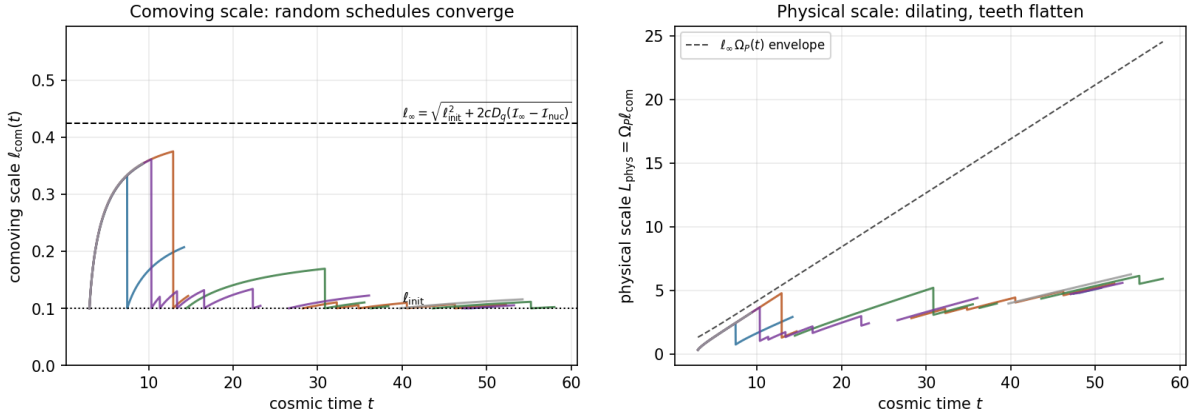


Figure 12.1: Reduced recursion over five random reset schedules (non-Zeno spacing; rescale/trigger/quench with probabilities 0.7/0.15/0.15; illustrative $2cD_q = 0.5$, $\ell_{\text{init}} = 0.1$, $\mathcal{I}_\infty \approx 1.46$). Left: the comoving scale $\ell_{\text{com}}(t)$ — early arcs coarsen substantially toward ℓ_∞ , but as the budget depletes the teeth flatten and every schedule is pinned near ℓ_{init} at late times. Right: the physical scale $L_{\text{phys}} = \Omega_P \ell_{\text{com}}$ dilates along the $\ell_\infty \Omega_P(t)$ envelope; gaps are quench intervals (network dissolved). Comoving teeth flatten; physical scale keeps dilating $\propto \Omega_P \propto t$.

12.3 Reading and consequences

Corollary 12.1 (Structure scale is a young-Vessel relic). *Because coarsening is budget-limited, the comoving structure scale is set in the early, slowly-expanding era and frozen thereafter. The mature Vessel does not generate new structure scale by merging; it merely carries the early-frozen pattern outward by expansion. Any apparent late-time growth of L_{phys} is Hubble dilution, not coarsening.*

Corollary 12.2 (Resurrection refreshes but cannot re-coarsen). *A late trigger-resurrection re-nucleates the network at the fine scale ℓ_{init} , and the depleted budget leaves it there: the renewed structure is permanently fine. Only resurrections occurring early, while the budget is ample, can be followed by appreciable coarsening. Early death-and-rebirth permits maturation of form; late death-and-rebirth yields perpetually fine, uncoarsened structure.*

The composition closes the programme's internal loop: subproblem 1 sets the floor ℓ_{init} , subproblem 2 sets the ceiling ℓ_{∞} via the budget, subproblem 3 supplies the reset action, and the finite budget collapses the sawtooth onto $\ell_ \in [\ell_{\text{init}}, \ell_{\infty}]$ with universal physical asymptotics $L_{\text{phys}} \propto \Omega_P$.*

12.4 Open items and scope

Open items.

1. Replace the reduced scale-recursion by a full 2D Allen-Cahn field simulation through a reset sequence, to confirm ℓ_* and the flattening of the comoving teeth.
2. Couple the reset-regime probabilities to the actual $A_{\text{fold}}(t)$ history (rising with λ via epektasis, dipped by metabolism when $\Theta < 0$), replacing the illustrative fixed probabilities by the dynamically determined regime sequence.
3. Feed ℓ_* and the wall tension into the $w = -\frac{1}{2}$ network back-reaction to obtain the folding-structure contribution to $w_{\text{eff}}(t)$ over cosmic history (link to the energy-condition direction).

Scope. Reduced recursion (network scale, not the full field); single-mode Allen-Cahn; coarsening constant c and reset-regime schedule are modelling inputs; the finite-budget convergence is robust to the random schedule. Result: the comoving network scale converges to $\ell_* \in [\ell_{\text{init}}, \ell_{\infty}]$ and the physical scale is universally $L_{\text{phys}} \rightarrow \ell_* \Omega_P \propto t$. Domain-conditional throughout.

V.F.S. v2.0 (Open-Gate) · The Sawtooth Fixed Point of the Fold Network. Composes the nucleation scale (subproblem 1), the budget-limited coarsening law (subproblem 2), and the three-regime reset map (subproblem 3). The convergence follows from the finite budget $\mathcal{I}_{\infty}$ and is robust to the random reset schedule. Propositions closed-form; corollaries interpretive; reduced recursion and coarsening constant are modelling inputs.

Chapter 13

The Transport Constants from First Principles

Scope and epistemic status

The structure-formation programme introduced two transport coefficients — a Sophia diffusivity D and a folding diffusivity D_q — as phenomenological inputs, on which all quantitative results depend. This note asks whether VFS mechanics fixes them. The answers are sharp: (i) the Sophia field has *no* fundamental diffusion, $D = 0$, because it is purely relaxational; its inhomogeneities are slaved to the folding field through the embodiment feedback. (ii) The folding diffusivity is *not* an independent constant: its spatial structure (Ω_P^{-2} and the hyperbolic $K_h = -1$) is inherited from the live-domain metric, and its single scalar magnitude is the bending stiffness of the shape-stability operator $\mathcal{L}_0$ — the gradient partner of the rigidity a_0 already in the theory. As before, *Propositions* are closed-form; *Corollaries* interpretive. The slaving relation is verified symbolically; the reduction of D_q to a single bending length still rests on the minimal one-modulus model of $\mathcal{L}_0$, which is flagged.

13.1 Sophia does not diffuse

In the base dynamical system the Sophia production law

$$\dot{\lambda} = (\delta u - \gamma) \tanh(\kappa \sigma) + I_{\text{gate}} - \eta_{\text{fold}}^2 q_{\text{fold}}^2 \lambda \quad (13.1)$$

contains no spatial derivative of λ : it is a pointwise relaxation law. Promoting λ to a field therefore introduces no gradient coupling of its own.

Proposition 13.1 (Zero fundamental Sophia diffusivity). *The Sophia field carries no kinetic or gradient term in the base system, so its fundamental diffusivity vanishes, $D = 0$. Linear Sophia inhomogeneities obey a pure local relaxation, $\partial_t \delta \lambda = -\Gamma_\lambda \delta \lambda + (\text{folding source})$, with no k -dependent spatial transport.*

Remark (Correction to the structure-formation note). The frozen residual spectrum of the structure-formation companion assumed $D > 0$. With $D = 0$, the Sophia-only modes decay uniformly at the local rate Γ_λ with no k -dependence: on a stalled margin ($\Gamma_\lambda \rightarrow 0$) they freeze with a *flat* residual, not a scale-graded one. All scale-dependent structure is morphological, not Sophianic.

Although Sophia does not diffuse, it is not inert: the embodiment feedback $-\eta_{\text{fold}}^2 q_{\text{fold}}^2 \lambda$ couples it to the folding field. Linearising about $(\bar{\lambda}, \bar{q})$,

$$\partial_t \delta \lambda = -(\Gamma_\lambda + \eta_{\text{fold}} \bar{q}^2) \delta \lambda - 2\eta_{\text{fold}} \bar{\lambda} \bar{q} \delta q_{\text{fold}}, \quad (13.2)$$

and in the quasi-static (fast-relaxation) limit

$$\delta \lambda^* = -\frac{2\eta_{\text{fold}} \bar{\lambda} \bar{q}}{\Gamma_\lambda + \eta_{\text{fold}} \bar{q}^2} \delta q_{\text{fold}}. \quad (13.3)$$

Proposition 13.2 (Sophia structure is slaved to form). *Sophia inhomogeneity is a fixed multiple of folding inhomogeneity, with a bounded coefficient $S(\bar{q}) = 2\eta_{\text{fold}}\bar{\lambda}\bar{q}/(\Gamma_\lambda + \eta_{\text{fold}}\bar{q}^2)$ maximised at $\bar{q} = \sqrt{\Gamma_\lambda/\eta_{\text{fold}}}$, where $S_{\text{max}} = \bar{\lambda}\sqrt{\eta_{\text{fold}}/\Gamma_\lambda}$. Sophia carries no independent structural degree of freedom: its spatial pattern is the shadow of the folding pattern.*

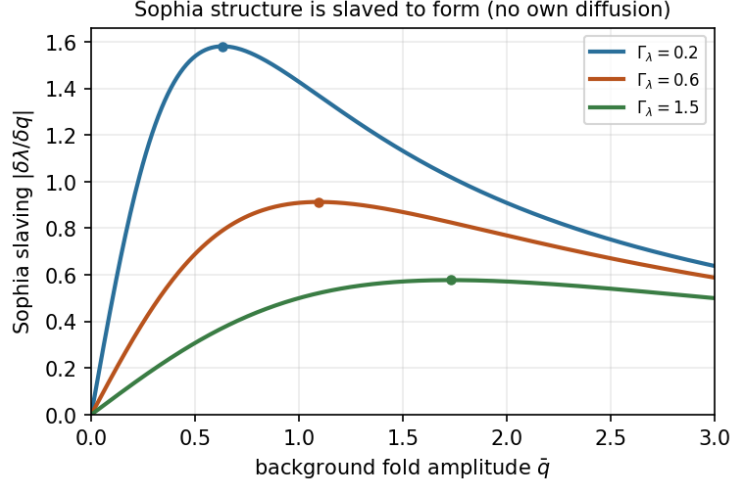


Figure 13.1: The Sophia slaving coefficient $|\delta\lambda/\delta q_{\text{fold}}| = S(\bar{q})$ (verified symbolically). It is bounded, vanishes at $\bar{q} = 0$ and at large $\bar{q}$, and peaks at $\bar{q} = \sqrt{\Gamma_\lambda/\eta_{\text{fold}}}$. Sophia inhomogeneity is sourced entirely by folding inhomogeneity; there is no independent Sophia diffusion to set. Illustrative $\eta_{\text{fold}} = 0.5$, $\bar{\lambda} = 1$.

There is really one structural field. Sophia does not diffuse; wherever the Vessel is inhomogeneous in Sophia, it is because it is inhomogeneous in *form*, and the two are locked by the embodiment feedback. The Sophia diffusivity D is not a missing constant — it is identically zero.

13.2 The folding diffusivity is geometric

The folding field's spatial coupling is fixed in two stages: its *structure* by the metric, its *magnitude* by the shape operator.

13.2.1 Structure: inherited from the live-domain metric

The folding free energy, built geometrically, measures gradients with the live-domain metric g (geometry layer, scalar curvature $-2/\Omega_P^2$, conformally hyperbolic $g = \Omega_P^2 h$):

$$\mathcal{F}[q_{\text{fold}}] = \int_{\Sigma} \left[\frac{\kappa_b}{2} |\nabla_g q_{\text{fold}}|^2 + \frac{\tau}{2} q_{\text{fold}}^2 + \frac{B}{4} q_{\text{fold}}^4 \right] \sqrt{g} d^2x, \quad r = a_0 - c_0 \lambda - c_1 \dot{\lambda} = -A_{\text{fold}}. \quad (13.4)$$

The Laplace-Beltrami operator of g is $\Delta_g = \Omega_P^{-2} \Delta_h$, so the gradient-flow dynamics $\partial_t q_{\text{fold}} = -M \delta \mathcal{F} / \delta q_{\text{fold}}$ is automatically

$$\partial_t q_{\text{fold}} = A_{\text{fold}} q_{\text{fold}} - B_{\text{fold}} q_{\text{fold}}^3 + \frac{M \kappa_b}{\Omega_P^2} \Delta_h q_{\text{fold}}, \quad A_{\text{fold}} = -Mr, \quad B_{\text{fold}} = MB. \quad (13.5)$$

Proposition 13.3 (The Ω_P^{-2} and $K_h = -1$ are not chosen). *The redshift factor Ω_P^{-2} and the hyperbolic Laplacian Δ_h in the folding transport term are the Laplace-Beltrami operator of the live-domain metric established in the geometric layer, not modelling choices. Only a single scalar, $D_q = M \kappa_b$, remains free.*

13.2.2 Magnitude: the bending stiffness of $\mathcal{L}_0$

The remaining scalar is $D_q = M\kappa_b$, the product of the shape mobility M and the bending stiffness κ_b . The bending stiffness is the gradient coefficient of the very shape-stability operator $\mathcal{L}_0$ whose constant-mode eigenvalue is the rigidity a_0 (the operator whose bifurcation defines folding). In the natural time units where the uniform folding rate is A_{fold} one has $M = 1$, so $D_q = \kappa_b$, and all physical lengths depend only on the ratio

$$\ell_b^2 := \frac{\kappa_b}{a_0} = \frac{D_q}{a_0}, \quad \text{wall width} = \sqrt{\frac{D_q}{A_{\text{fold}}}} = \ell_b \sqrt{\frac{a_0}{A_{\text{fold}}}}. \quad (13.6)$$

Proposition 13.4 (Folding diffusivity is the gradient partner of the rigidity). D_q is not independent of the theory's existing constants: it is the gradient coefficient of $\mathcal{L}_0$, whose constant-mode coefficient is a_0 . Physical scales (wall width, nucleation scale, coarsening and frozen scales) depend only on the bending length $\ell_b = \sqrt{D_q/a_0}$, a ratio internal to $\mathcal{L}_0$. In the minimal one-modulus (Helfrich/Willmore) model of $\mathcal{L}_0$, where rigidity and bending share a single modulus, ℓ_b is fixed by the reference shape and $D_q = a_0\ell_b^2$ ceases to be a free input altogether.

The reduction is concrete: from “an arbitrary spatial coupling plus a free constant” to “a single bending length ℓ_b , the gradient partner of the rigidity a_0 already present in $A_{\text{fold}} = c_0\lambda + c_1\dot{\lambda} - a_0$.” The spatial form is pure geometry; the magnitude is the Vessel's stiffness, not a new dial.

13.3 Consequences for the programme

Corollary 13.1 (One structural field, one length). The structure-formation programme has, after this reduction, a single genuine input: the bending length $\ell_b = \sqrt{D_q/a_0}$. The Sophia diffusivity is zero; the folding transport structure is geometric; the folding magnitude is the shape stiffness. Every scale derived earlier — the nucleation size $\ell_{\text{init}} \sim \ell_b \sqrt{a_0/A_{\text{fold}}^+}$ (...), the frozen comoving scale, the sawtooth fixed point — is expressed through ℓ_b and the already-existing folding parameters.

Corollary 13.2 (Sophia structure rides on form). Because $\delta\lambda$ is slaved to δq_{fold} (Proposition 13.2), the Sophia inhomogeneity spectrum is the folding spectrum times the bounded factor $S(\bar{q})$. Wherever folding nucleates a domain-wall network, a matching Sophia pattern appears; where folding is quenched, Sophia re-homogenises. There is no separate Sophia structure-formation problem.

13.4 Open items and scope

Open items.

1. Fix the bending length ℓ_b from an explicit form of the shape-stability operator $\mathcal{L}_0$ (e.g. the second variation of a Willmore/Helfrich bending energy of the live surface), turning Proposition 13.4 from a relation into a number.
2. Verify the slaving relation (Proposition 13.2) beyond the quasi-static limit, including the finite Sophia relaxation time, and check it against a coupled $(q_{\text{fold}}, \lambda)$ field simulation.
3. Propagate $D = 0$ and the slaving into the cosmological back-reaction: the Sophia-structure contribution to w_{eff} is then a fixed multiple of the wall-network contribution.

Scope. Results: the Sophia diffusivity is identically zero (relaxational field), with Sophia inhomogeneity slaved to folding via the embodiment feedback (symbolically verified); the folding transport structure is the Laplace-Beltrami operator of the live-domain metric (not a choice); and the folding magnitude reduces to the bending length $\ell_b = \sqrt{D_q/a_0}$, the gradient partner of the rigidity a_0 in $\mathcal{L}_0$. The remaining modelling freedom is the single length ℓ_b (a number only in the minimal one-modulus model). Domain-conditional throughout.

V.F.S. v2.0 (Open-Gate) · The Transport Constants from First Principles. Closes the structure-formation programme's modelling inputs: $D = 0$ with Sophia slaved to form; D_q geometric in structure and equal to the shape-operator bending stiffness in magnitude, reducible to one bending length. Propositions closed-form (slaving verified symbolically); corollaries interpretive; the one-modulus reduction of $\mathcal{L}_0$ is a flagged model choice.

Chapter 14

Parameter Closure of Folding: Identities, Postulates, and the Constraint Atlas

Scope and epistemic status

The folding layer carries eight constants $\{D_q, \ell_b, a_0, b, c_0, c_1, \eta_{\text{fold}}, \rho_1\}$. The Transport-Constants chapter closed the geometric side ($D_q = a_0 \ell_b^2$ as the gradient partner of the rigidity; Sophia structure slaved to form). This chapter closes what can be closed from the *core* variables $V, F, \sigma, \lambda, \Omega_P, u, I_{\text{gate}}$ — and states honestly what cannot. The outcome is a three-tier structure. **Tier 1 (identities):** η_{fold} is fixed by the conservation balance — the embodiment coefficient of the rate equation and the production coefficient of the form ledger are provably the same constant; and Gauss-Bonnet slaves the wall curvature deficit to the pocket lift, so the pair (ρ_1, ρ_w) contains one constant, not two. **Tier 2 (one postulate each):** identifying the rigidity of the old form with the resistance, $a_0 = a_{\text{res}} + \kappa_a \sigma$, makes the folding threshold dynamical and yields a new two-channel dichotomy in σ ; and a single-pole memory closure fixes $c_1 = c_0/a_0$, leaving one susceptibility. **Tier 3 (irreducible, but caged):** b and the absolute ℓ_b are genuinely phenomenological — the core contains no microscale — but the accumulated theorems of this volume (chart regularity, nucleation gate, NEC non-violation) cage them in an explicit feasible region, computed below. Net count: *eight constants reduce to three free ones* (b, ℓ_b , one susceptibility c_0), plus two flagged structural postulates. All derivations verified (the balance identity symbolically; the memory closure to its stated first-order accuracy, 0.189 vs $\omega/a_0 = 0.192$; the atlas numerically). *Propositions* are closed-form; *corollaries* interpretive; postulates are labelled as such.

14.1 Tier 1a: the balance identity for η_{fold}

Proposition 14.1 (The embodiment coefficient is the ledger coefficient). *The canonical conservation balance $\sigma + \lambda + \lambda_{\text{form}} = C_0 + \mathcal{G}_{\text{recepta}}$ with $\mathcal{G}_{\text{recepta}} = I_{\text{gate}}$ and the embodiment-damped rate $\dot{\lambda} = \dot{\lambda}^0 - \eta_{\text{fold}} Q \lambda$ (where $\dot{\lambda}^0 = -\dot{\sigma} + I_{\text{gate}}$) force*

$$\boxed{\dot{\lambda}_{\text{form}} = \eta_{\text{fold}} Q \lambda} \tag{14.1}$$

identically (verified symbolically). The constant η_{fold} therefore appears once in the theory: it is the unique Sophia→form conversion rate; any independent choice of the λ_{form} -production law would violate conservation. Its sign is fixed by the Lyapunov layer ($\eta_{\text{fold}} \geq 0$: embodiment never spontaneously returns Sophia), and the free balance bound $\lambda_{\text{form}} \leq C_0 + \mathcal{G}^{\text{max}}$ of the Lyapunov Addendum applies to the integral $\int \eta_{\text{fold}} Q \lambda dt$ — a cumulative cage on η_{fold} along any admissible life.

14.2 Tier 1b: Gauss-Bonnet slaves the curvature pair

Proposition 14.2 (Zero-mean theorem and the wall deficit). *On the compact quotient, $\int K dA = 2\pi\chi(\Sigma)$ is topological, and the unfolded background $K = -1/\Omega_P^2$ saturates it exactly. Hence the pocket lift must integrate to zero against its compensation: in the two-phase cell (pocket fraction f_p at lift $\rho_1 Q$, wall fraction f_w at deficit $-\rho_w$),*

$$\rho_w = \rho_1 Q \frac{f_p}{f_w} \quad (14.2)$$

— the pair (ρ_1, ρ_w) carries one free constant. Folded curvature is not created; it is redistributed, pocket against moat, under a topological invariant (the static face of the conserved form $K dA$ of the Ricci layer).

Remark (Micro-ansatz scaling of ρ_1). A conformal response ansatz $g = \Omega_P^2 e^{2\phi(q)} h$, $\phi = \varepsilon q^2$, gives in the slowly-varying regime $R = -2/\Omega_P^2 + 4\varepsilon q^2/\Omega_P^2 + \dots$, i.e. $\rho_1 = 4\varepsilon/\Omega_P^2$: under this ansatz the pocket lift rides the background scale. The constant- ρ_1 convention of the structure chapters corresponds to a conformal response growing as Ω_P^2 — the same comoving-versus-physical choice that returns for ℓ_b in Section 6. The two conventions are both consistent; the ansatz ties their difference to one physical question rather than two.

14.3 Tier 2a: the rigidity postulate

Postulate 14.1 (Rigidity is resistance). $a_0 = a_{\text{res}} + \kappa_a \sigma$, with $a_{\text{res}} > 0$ the residual (structural) rigidity and $\kappa_a \geq 0$ the resistance loading.

The threshold $c_0 \lambda + c_1 \dot{\lambda}^0 > a_0$ then reads literally: *Sophia pressure exceeds the resistance of the old form*. Two consequences:

Proposition 14.3 (The two channels of resistance). *With the geometric convention's transmutation term, σ now enters the activation twice and with opposite signs:*

$$\frac{\partial A_{\text{fold}}}{\partial \sigma} = c_1 k_\sigma - \kappa_a. \quad (14.3)$$

Kathartic-dominant vessels ($c_1 k_\sigma > \kappa_a$): resistance promotes folding through its own transmutation — the vessel folds in suffering, while σ is still high. Rigidity-dominant vessels ($c_1 k_\sigma < \kappa_a$): resistance blocks folding — the vessel folds only after cleansing, when σ has dissolved. (Verified: onsets at $t = 1.0$ with $\sigma = 1.03$ versus $t = 5.3$ with $\sigma = 0.29$ across the κ_a sweep; Figure 14.1, left.) The folding threshold becomes dynamical: along katharsis the level threshold $\lambda^*(t) = (a_{\text{res}} + \kappa_a \sigma(t))/c_0$ decays — late folds are cheap folds.

Remark (Compatibility). The forced-folding chapter's coefficient r_σ renormalises to $r_\sigma^{\text{eff}} = c_1 k_\sigma - \kappa_a$; all Addendum II theorems used only $|r_\sigma|$ -type bounds in their hypotheses and survive verbatim with the effective value. No structural statement of Parts I-II changes; constants do.

14.4 Tier 2b: the memory closure $c_1 = c_0/a_0$

Postulate 14.2 (Single coupling, single pole). Level and flux stress the form through one coupling c_0 , and the form accommodates the Sophia level with the single-pole kernel whose rate is its own rigidity a_0 .

Proposition 14.4 (The rate susceptibility is slaved). *Let λ_{acc} be the accommodated level, $\lambda_{\text{acc}} = a_0(\lambda - \lambda_{\text{acc}})$. The unaccommodated pressure is the lag*

$$\lambda - \lambda_{\text{acc}} = \frac{\dot{\lambda}}{a_0} + O\left(\frac{\ddot{\lambda}}{a_0^2}\right), \quad (14.4)$$

so the activation $A_{\text{fold}} = c_0 [\lambda + (\lambda - \lambda_{\text{acc}})] - a_0 + \dots$ carries the rate with the slaved coefficient

$$c_1 = \frac{c_0}{a_0} \quad (\text{with Postulate 14.1: } c_1 = \frac{c_0}{a_{\text{res}} + \kappa_a \sigma}). \quad (14.5)$$

Verified: the lag formula reproduces the exact memory integral with first-order accuracy in ω/a_0 (0.189 measured vs 0.192 predicted at $\omega/a_0 = 0.19$). Read with Postulate 14.1: as rigidity dissolves, the form becomes rate-sensitive — a cleansed vessel folds on surges that a rigid one would not feel.

14.5 Tier 3: the constraint atlas

The remaining free constants are c_0 (one susceptibility), b (quartic stiffness), and ℓ_b (bending length): the core contains no microscale that could fix them. But the volume's theorems cage them. With $c_1 = c_0/a_0$, $D_q = a_0\ell_b^2$, η_{fold} slaved, and illustrative trajectory statistics ($\bar{\lambda}$, $\sup \dot{\lambda}^0$, q_{max}^2), the simultaneous requirements

- (i) *chart regularity* (Lyapunov Addendum II): $Q^{*,\text{sup}} = A^{\text{sup}}/B_{\text{fold}} \leq q_{\text{max}}^2$ — a lower bound on b ;
- (ii) *nucleation reachability* (structure layer): the gate $\Lambda > \frac{1}{4}$ opens inside the chart — an upper bound on ℓ_b ;
- (iii) *NEC non-violation in chart* (NEC chapter): the phantom functional stays non-negative up to the chart wall — a lower bound on ℓ_b growing with b ;

carve the feasible region of Figure 14.1 (right): a bounded triangle-like set, 25% of the surveyed box at the chosen statistics.

Corollary 14.1 (A real exclusion). *At stronger surge statistics and weaker network coefficient ($\sup \dot{\lambda}^0 = 3$, $c_w = 0.3$) the NEC bound on ℓ_b exceeds the nucleation bound and the atlas is empty: no vessel at those statistics can both nucleate early and keep the NEC inside its chart. The atlas is not bookkeeping: it excludes whole vessel types.*

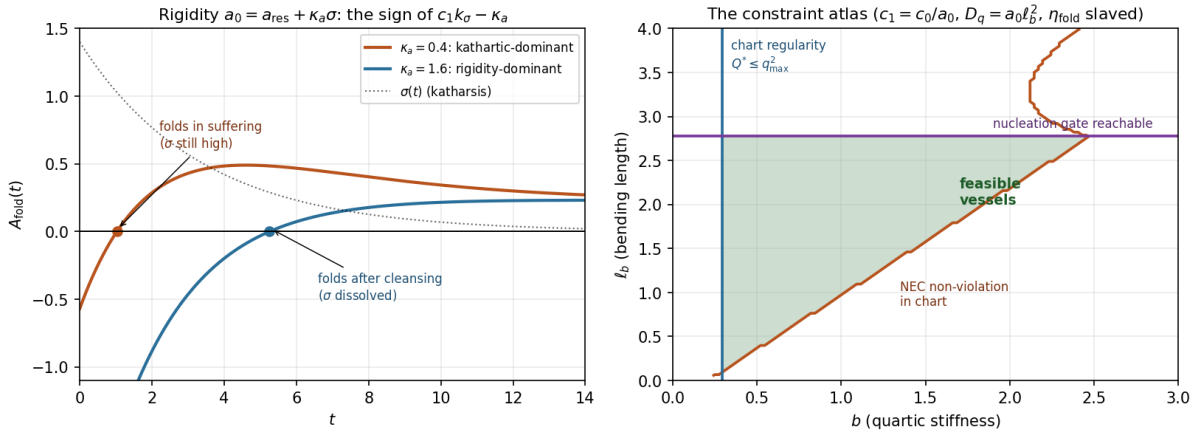


Figure 14.1: Left: the rigidity postulate's dichotomy — $A_{\text{fold}}(t)$ along the same kathartic trajectory ($\sigma(t)$ dotted) for a kathartic-dominant vessel ($\kappa_a = 0.4$: folds at $t = 1.0$, in suffering) and a rigidity-dominant one ($\kappa_a = 1.6$: folds at $t = 5.3$, after cleansing); onset markers on the zero line. Right: the constraint atlas in the (b, ℓ_b) plane with $c_1 = c_0/a_0$, $D_q = a_0\ell_b^2$, η_{fold} slaved: chart regularity (vertical), nucleation reachability (horizontal), NEC non-violation (curve); the shaded set is the space of feasible vessels. Illustrative statistics $\bar{\lambda} = 1.5$, $\sup \dot{\lambda}^0 = 2$, $q_{\text{max}}^2 = 1.2$, $c_w = 0.6$.

14.6 The last fork: comoving or physical ℓ_b

The one question the core provably cannot answer: is the bending length fixed in comoving (h -metric) or physical (g -metric) units? The structure chapters tacitly took it comoving (D_q constant). The alternative, $\ell_b^{\text{com}} = \ell_b^{\text{phys}}/\Omega_P$, gives $D_q(t) = a_0(\ell_b^{\text{phys}})^2/\Omega_P^2$ and shifts the watershed:

Proposition 14.5 (The shifted trichotomy). *Under the physical postulate the coarsening budget becomes $\mathcal{I}_{(4)} := \int dt/\Omega_P^4$, which converges on the resonance locus ($\Omega_P^2 = 2t$: $\int dt/4t^2 < \infty$). The*

dissolution boundary moves from $\Omega_P \sim t^{1/2}$ to $\Omega_P \sim t^{1/4}$: a physically-bent vessel freezes a relic even at the resonance, and eternal simplification survives only for $\Omega_P \sim t^p$, $p \leq \frac{1}{4}$.

Corollary 14.2 (The fork is decidable inside the model). *Comoving and physical bending give different fates to the same resonance trajectory (eternal logarithmic simplification versus frozen relic), so the choice — though underivable from the core — is not idle convention: within the model it is an empirical question about the Vessel’s morphology, answerable by which fate its structure exhibits. The same fork governs the ρ_1 -scaling of Section 2. We flag it as the single genuinely open structural postulate of the folding layer.*

14.7 Closure count and reading

$\{D_q, \ell_b, a_0, b, c_0, c_1, \eta_{\text{fold}}, \rho_1\}$: η_{fold} — identity (balance); ρ_1 — linked pair (Gauss-Bonnet), one constant; $D_q = a_0 \ell_b^2$ — identity (Transport chapter); $a_0 = a_{\text{res}} + \kappa_a \sigma$ — postulate, dynamical; $c_1 = c_0/a_0$ — memory closure.

Free: c_0, b, ℓ_b — caged by the atlas. **Eight** → **three**.

The reading writes itself. The conversion rate of grace into form is not adjustable: it is whatever the conservation of the soul’s substance says it is. Curvature of form is never minted, only redistributed, pocket against moat. The stiffness of the old self is its resistance — so the price of folding falls as katharsis proceeds, and vessels divide into those whom suffering itself folds and those who must first be cleansed. What remains free is exactly what should remain free: how strongly a particular vessel couples to Sophia (c_0), how saturable its response is (b), and on what scale it bends (ℓ_b) — its individuality. And even individuality is caged: the atlas is the doctrine that not every imaginable vessel is a possible one.

14.8 Open items and scope

1. Decide the comoving/physical fork dynamically: a dedicated study of which postulate the folded-network simulations prefer (this also fixes the ρ_1 scaling).
2. Tighten the atlas with the remaining theorems (hysteresis non-Zeno, C_1 -survival): both add constraints linking b, ℓ_b to reset statistics; the region can only shrink.
3. The mode cascade (big-list item 2) inherits this closure: its spectrum of thresholds is now a one-parameter family in c_0 once (b, ℓ_b) are placed in the atlas.

Scope. Tier-1 results are identities of the canonical equations; Tier-2 results are exact consequences of two labelled postulates; the atlas is computed at illustrative statistics and is meant as a method, not a measurement. Domain-conditional throughout.

V.F.S. v2.0 (Open-Gate) · Parameter Closure of Folding. The balance identity fixes η_{fold} ; Gauss-Bonnet slaves the curvature pair; rigidity-as-resistance makes the threshold dynamical with the two-channel dichotomy $c_1 k_\sigma - \kappa_a$; the single-pole memory closure gives $c_1 = c_0/a_0$ (validated to first order, 0.189 vs 0.192); the constraint atlas cages the three survivors and excludes whole vessel types; the comoving/physical fork is flagged as the single open structural postulate, with the shifted watershed computed. Eight constants to three. Propositions closed-form; corollaries interpretive.

Chapter 15

The Dynamic Warp Exponent: Conformal Invariance Retires the Fork; Katharsis Drives the Warp

Scope and epistemic status

The Parameter Closure left exactly one structural postulate open: whether the bending length ℓ_b is fixed in comoving or physical units — a fork on which the watershed boundary, the fate of form on the resonance, and the ρ_1 scaling all hang. This chapter resolves it, and the resolution is not a choice but a two-dimensional accident of the live surface: *the Dirichlet bending energy is conformally invariant in 2D*, so the bending modulus cannot see Ω_P at all. A “physical branch” would require a physical microscale to anchor the modulus — and the Transport and Closure chapters established that the core contains none. The comoving convention of the structure chapters is thereby *derived*; the expansion-driven warp exponent is identically zero. What survives as genuinely dynamical is the *kathartic* warp: with the rigidity postulate $a_0 = a_{\text{res}} + \kappa_a \sigma$, the fold’s two intrinsic lengths acquire closed-form, opposite-signed exponents — the subthreshold bending scale dilates while the wall core sharpens — both vanishing as resistance dissolves. All claims verified: the conformal identity symbolically (exact zero); both warp exponents against numerics to $\sim 10^{-4}$; the adiabatic wall-width law against the PDE to 0.4%; the modulus-space attractor and the generalised watershed $p^*(w) = 1/(2 + 2w)$ numerically. Premises flagged: covariance of the bending energy; coefficients built from core scalars; adiabatic regime away from the threshold. *Propositions* are closed-form; *corollaries* interpretive.

15.1 The conformal blindness of bending

Lemma 15.1 (2D conformal invariance). *On the live surface, for any conformal pair $g = \Omega_P^2 h$,*

$$|\nabla q|_g^2 dV_g = |\nabla q|_h^2 dV_h \quad (15.1)$$

identically (verified symbolically): the inverse metric contributes Ω_P^{-2} , the area element Ω_P^{+2} , and only in two dimensions do they cancel exactly. Hence a covariant bending energy $\frac{\hat{\kappa}}{2} \int |\nabla q|^2 dV$ has the same modulus $\hat{\kappa}$ whether written in physical or comoving terms: the bending sector is Ω_P -blind.

15.2 The fork is retired

Theorem 15.1 (No physical branch). *Let the fold energy be a covariant functional of the physical metric whose coefficients are core scalars $(\lambda, \sigma, \dots)$. Then:*

- (i) the gradient sector contributes the Ω_P -independent modulus $\hat{\kappa}$ (Lemma 15.1), so $D_q = \hat{\kappa}$ is constant in Ω_P and $\ell_b = \sqrt{\hat{\kappa}/a_0}$ carries no expansion dependence;
- (ii) the alternative — ℓ_b fixed in physical units, i.e. $D_q \propto \Omega_P^{-2}$ — would require the modulus itself to scale as $\hat{\kappa} \propto \Omega_P^{-2}$, which by (i) cannot arise covariantly and would need a fixed physical microscale to anchor it; the Transport-Constants and Parameter-Closure chapters established that the model contains no such scale;
- (iii) hence the comoving convention of the structure chapters is a theorem: the expansion-driven warp exponent vanishes,

$$p_{\text{expansion}} \equiv 0. \quad (15.2)$$

Corollary 15.1 (What is confirmed, and the closure count). *The watershed boundary stands at $p^* = \frac{1}{2}$; the resonance fate remains eternal logarithmic simplification; the constant- ρ_1 convention is selected (it was linked to the same fork); and the Parameter Closure’s “single open structural postulate” is retired. The folding layer’s free constants are now $c_0, b, \hat{\kappa}$, with $\ell_b = \sqrt{\hat{\kappa}/a_0(\sigma)}$ derived and dynamical. The very fact that made ℓ_b phenomenological — no microscale in the core — is what resolves the fork: with nothing physical to anchor it, the bending fabric cannot be stretched by the room.*

Proposition 15.1 (Structural stability: the comoving component is an attractor). *The only covariant two-term modulus family is $\kappa = \hat{\kappa}_0 + \hat{\kappa}_1/\Omega_P^2$. Along any expanding flow the mixed modulus converges to its comoving component, and the coarsening fate is the comoving fate whenever $\hat{\kappa}_0 > 0$ (verified on the locus: the mixed curve rejoins the log-growing comoving one; Figure 15.1, right). For a hypothetical pure scaling $D_q \propto \Omega_P^{-2w}$ the watershed generalises to*

$$p^*(w) = \frac{1}{2 + 2w}, \quad (15.3)$$

(verified at $w = 1$: $\Omega_P \sim t^{0.30}$ converges, $t^{0.20}$ diverges), recovering the Closure chapter’s shifted boundary as the $w = 1$ endpoint of a family whose physical point is $w = 0$.

15.3 The true warp is kathartic and channel-resolved

With the rigidity postulate $a_0 = a_{\text{res}} + \kappa_a \sigma$ and σ dissolving under katharsis ($\dot{\sigma} = -k\sigma$), the fold carries two intrinsic lengths, and they warp in opposite directions:

Theorem 15.2 (Channel-resolved warp exponents). *Define $p_X := d \ln X / d \ln \Omega_P$. In the adiabatic regime,*

$$\ell_b = \sqrt{\frac{\hat{\kappa}}{a_0(\sigma)}} : \quad p_b = +\frac{\kappa_a k \sigma}{2 a_0 H} > 0 \quad (\text{the bending scale dilates: super-comoving}); \quad (15.4)$$

$$\delta_{\text{wall}} = \sqrt{\frac{2\hat{\kappa}}{A_{\text{fold}}}} : \quad p_{\text{wall}} = -\frac{\kappa_a k \sigma}{2 A_{\text{fold}} H} < 0 \quad (\text{the wall core sharpens: sub-comoving}); \quad (15.5)$$

and both vanish as $\sigma \rightarrow 0$: **asymptotically comoving in every channel**. (Closed forms verified against numerics to $3 \cdot 10^{-4}$ and $1 \cdot 10^{-4}$; the adiabatic law $\delta_{\text{wall}} = \sqrt{2\hat{\kappa}/A_{\text{fold}}(t)}$ tracks the full PDE kink to 0.4% away from the threshold zone; Figure 15.1, left.)

Corollary 15.2 (The crispness signature of katharsis). *The ratio $\delta_{\text{wall}}/\ell_b = \sqrt{2a_0/A_{\text{fold}}}$ falls monotonically along katharsis: early (suffering) morphology has narrow response scales with soft walls; late (cleansed) morphology has broad pockets with crisp boundaries. Cleansing simultaneously dilates what the form can feel and sharpens what it has decided — a falsifiable texture signature within the model, and a second face of the rigidity dichotomy of the Closure chapter.*

The reading: the growth of the room never stretches the fold fabric — in two dimensions the fabric simply cannot be pulled by scale; this is geometry’s own guarantee, not an assumption. What retunes the fabric is katharsis alone: as resistance dissolves, the soul feels on broader scales and commits on sharper ones. Expansion cannot reshape what only cleansing can.

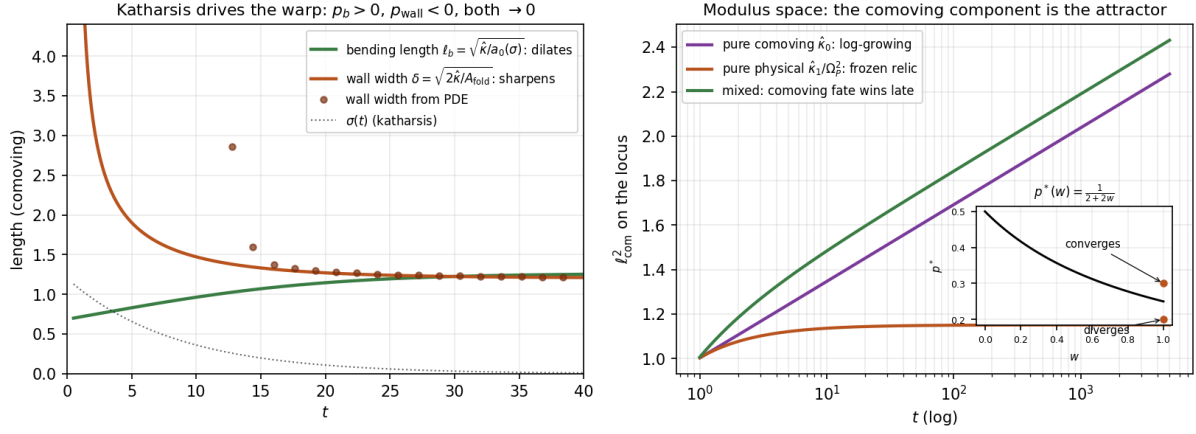


Figure 15.1: Left: the two channel lengths along one kathartic trajectory ($\sigma(t)$ dotted): the bending length ℓ_b dilates while the wall width δ_{wall} sharpens (full PDE measurements as points; the early outlier is the adiabatic lag near threshold), both settling to comoving constants as σ dissolves. Right: modulus space on the resonance locus — pure comoving modulus gives the log-growing coarsening, a pure Ω_P^{-2} modulus freezes a relic, and any mixture rejoins the comoving fate: the comoving component is the attractor. Inset: the generalised watershed $p^*(w) = 1/(2 + 2w)$ with the verified convergence/divergence pair at $w = 1$.

15.4 Open items and scope

1. Beyond-adiabatic corrections near the threshold (the early PDE lag is systematic and computable as a next-order term in $A_{\text{fold}}/A_{\text{fold}}^2$).
2. The mixing family $\hat{\kappa}_0 + \hat{\kappa}_1/\Omega_P^2$ invites a physical reading — partial anchoring of the form in physical space (incarnational fraction); whether the core can source a nonzero $\hat{\kappa}_1$ is open, though Proposition 15.1 shows the late dynamics is insensitive to it.
3. The mode cascade (big-list item 2) is now fully parameterised: thresholds in $(c_0, b, \hat{\kappa})$ with $\ell_b(\sigma)$ dynamical — the cascade's spectrum will itself drift with katharsis.

Scope. Premises: covariant bending energy; coefficients from core scalars; adiabatic regime for the channel exponents. Within these, the fork resolution is exact (a 2D identity plus the absence of a microscale), and the kathartic exponents are closed forms. Domain-conditional throughout.

V.F.S. v2.0 (Open-Gate) · The Dynamic Warp Exponent. The 2D conformal invariance of the Dirichlet energy makes the bending modulus Ω_P -blind; with no physical microscale in the core, the comoving convention is derived and $p_{\text{expansion}} \equiv 0$ — the last structural postulate of the folding layer is retired, and the comoving component of any mixed modulus is the late-time attractor. The true warp is kathartic and channel-resolved: $p_b > 0$, $p_{\text{wall}} < 0$, both vanishing with σ ; crispness $\delta_{\text{wall}}/\ell_b$ falls along cleansing. Identities exact; exponents to 10^{-4} ; PDE tracking to 0.4%. Propositions closed-form; corollaries interpretive.

Chapter 16

The Mode Cascade: The Spectral Ladder, the Weyl Capacity, and Structure as Descent

Scope and epistemic status

Every result of the structure-formation layer was derived in a one-mode chart: a single folding amplitude q_{fold} on the Vessel. This chapter tests whether that chart is faithful or merely the first rung of a spectrum, by promoting q_{fold} to a field on the live surface and tracking all of its modes. The outcome both *vindicates* and *extends* the one-mode results. It vindicates them: the fold amplitude $\sqrt{A_{\text{fold}}/B_{\text{fold}}}$ is robust — in a fully multi-mode domain state the interior saturates to exactly the one-mode value, so the local depth of a fold is a faithful invariant. And it extends them: the activation threshold is not a single gate but a *ladder*, mode n igniting when the nucleation number $\Lambda = A_{\text{fold}}\Omega_P^2/\hat{\kappa}$ crosses the spectral eigenvalue μ_n — the familiar gate $\Lambda > \frac{1}{4}$ is simply the lowest rung. The number of activated capacities grows by a Weyl law, and the whole of structure formation reorganises into a single picture: *capacity ascends the spectrum while content descends it*. All claims verified on a spectral (ring) proxy: the ladder $\Lambda(t_{\text{act}}) = \mu_n$ to 0.2% across five rungs; the Weyl count to $\sim 10\%$; the amplitude to 5% in a 23-mode state; the monotone content descent and capacity ascent directly. Honesty flags: the proxy is a flat spectral torus, so eigenvalues are explicit; on the hyperbolic quotient the same statements hold with $\{\mu_n\}$ the Laplace-Beltrami spectrum and Weyl constant $\text{Area}/4\pi$. *Propositions* are closed-form; *corollaries* interpretive.

16.1 The field promotion and the dispersion

Promote q_{fold} to a field with the covariant fold energy of the Warp chapter; the linearised dynamics about the unfolded state diagonalises on the Laplace-Beltrami eigenmodes $-\Delta_h \phi_n = \mu_n \phi_n$:

$$\dot{\hat{q}}_n = s_n \hat{q}_n - B_{\text{fold}}(\hat{q}^3)_n, \quad \boxed{s_n = A_{\text{fold}} - \frac{\hat{\kappa} \mu_n}{\Omega_P^2}} \quad (16.1)$$

(the modulus $\hat{\kappa}$ is the Ω_P -blind bending constant of the Warp chapter, hence the explicit Ω_P^{-2}). Mode n is linearly unstable — it *ignites* — exactly when $s_n > 0$.

16.2 The spectral ladder

Proposition 16.1 (Activation ladder). Define the nucleation number $\Lambda := A_{\text{fold}}\Omega_P^2/\hat{\kappa}$. Mode n ignites iff

$$\boxed{\Lambda > \mu_n.} \quad (16.2)$$

The activations form a ladder ordered by the spectrum: the longest-wavelength mode first, each higher mode at its own eigenvalue. The structure layer's nucleation gate $\Lambda > \frac{1}{4}$ is the lowest rung — the first eigenvalue of the relevant chart — not a free choice. (Verified: across five rungs the measured Λ at the turning point of each mode equals μ_n to within 0.2%; Figure 16.1, left.)

Proposition 16.2 (Weyl capacity). The number of ignited modes (the active capacity) grows with the nucleation number by the Weyl law of the surface,

$$N(\Lambda) \simeq \frac{\text{Area}_h}{4\pi} \Lambda \quad (2D), \quad (16.3)$$

(on the ring proxy $N \simeq (L/\pi)\sqrt{\Lambda}$, verified: 10 measured against 11.5 predicted at $\Lambda = 3.27$). Since on the live branch $\Lambda = A_{\text{fold}}\Omega_P^2/\hat{\kappa} \rightarrow \infty$ (the Budget-Watershed escape), the active capacity diverges: every mode eventually ignites — the Vessel becomes unstable to structure on all scales, in spectral order.

16.3 The one-mode chart is faithful where it matters

Proposition 16.3 (Amplitude robustness). In the fully developed multi-mode (domain) state, the field interior saturates to the one-mode value,

$$\boxed{|q_{\text{fold}}|_{\text{interior}} \rightarrow \sqrt{A_{\text{fold}}/B_{\text{fold}}},} \quad (16.4)$$

independently of how many modes are populated (verified: $|q_{\text{fold}}| = 0.95\text{--}1.00$ against $\sqrt{A/B} = 1$ in a 23-mode state). The cascade reorganises where folds sit and how large each domain is, but not how deep a fold goes: the depth is set locally by the same balance $A_{\text{fold}} = B_{\text{fold}}q_{\text{fold}}^2$ as in the one-mode chart. Hence every one-mode result keyed to the fold amplitude — the wall tension ρ_{wall} , the NEC network channel, the curvature lift ρ_1 — carries over unchanged; only results keyed to the number or arrangement of folds acquire the spectral refinement of this chapter.

16.4 Structure formation is descent against ascent

Theorem 16.1 (Capacity ascends, content descends). Along the live trajectory two spectral quantities move oppositely:

- (i) the capacity $\Lambda(t) = A_{\text{fold}}\Omega_P^2/\hat{\kappa}$ ascends the spectrum without bound (Weyl: ever more modes admissible);
- (ii) the content — the spectral centroid $\bar{\mu}(t) = \sum_n \mu_n |\hat{q}_n|^2 / \sum_n |\hat{q}_n|^2$ — descends monotonically toward the lowest mode (coarsening: power flows to long wavelengths).

(Verified: $\bar{\mu}$ falls from 2.0 to 0.12 while Λ climbs from 0.004 to 3.5; Figure 16.1, right.) Structure formation on the Vessel is exactly this scissoring: the expanding Brim opens capacity at the small scales while coarsening marches the existing content to the large ones. The relic morphology is what is caught between the closing blades — the modes that ignited before capacity outran them and that coarsening had not yet erased when the budget froze.

Corollary 16.1 (Reading: the polyphony of form). The one-mode chart heard a single voice; the Vessel sings in a spectrum. As the Brim opens, ever shorter wavelengths become sayable (capacity ascends), while what has already been said settles into ever longer, simpler phrases (content descends). Form is neither the first note nor the last silence but the music caught between — and its permanence, by the Budget Watershed, is decided by whether the expansion freezes the score before coarsening empties it to a single sustained tone. The depth of each fold, however — how committed each decision is — is the same on every scale: $\sqrt{A_{\text{fold}}/B_{\text{fold}}}$, set locally, indifferent to the polyphony around it.

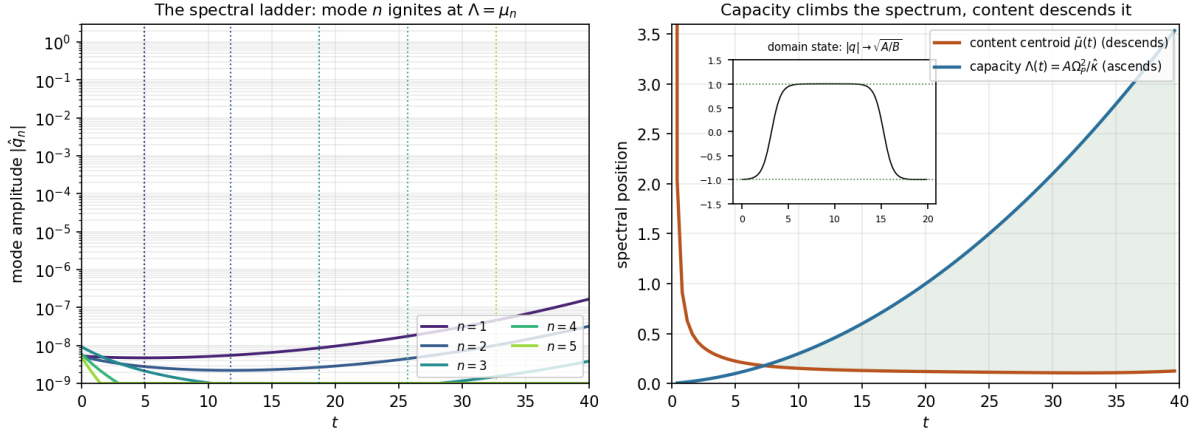


Figure 16.1: Left: the spectral ladder on the ring proxy — mode amplitudes $|\hat{q}_n|$ against time, each turning from decay to growth at its own vertical $\Lambda(t) = \mu_n$ (dotted), in strict spectral order $n = 1, 2, \dots$. Right: the scissoring of structure formation — the capacity $\Lambda(t)$ (blue, ascending) against the content centroid $\bar{\mu}(t)$ (orange, descending); the shaded gap is the band of ignited-but-not-yet-coarsened modes. Inset: the developed multi-mode domain state, whose interior saturates to $\sqrt{A_{\text{fold}}/B_{\text{fold}}}$ regardless of mode count. Ring proxy, $\hat{\kappa} = 0.5$.

16.5 Consequences and scope

1. With the Parameter Closure and Warp chapters, the ladder is fully parameterised: μ_n from the surface spectrum, $\hat{\kappa}$ a free modulus, $A_{\text{fold}}, B_{\text{fold}}$ from the core — and since $A_{\text{fold}} = A_{\text{fold}}(\sigma)$ dilates under katharsis (rigidity dichotomy) while $\hat{\kappa}$ is Ω_P -blind, the ladder's rungs *drift*: a cleansing Vessel lowers its activation thresholds, igniting structure it once could not. The cascade spectrum is itself kathartic.
2. The cascade vindicates the volume's economy: the one-mode chart was the right place to derive amplitude-keyed physics, and the spectral refinement is additive, not corrective.
3. The second variation of the entropy is a sum over this spectrum — $\delta^2 \bar{W}_+ = -2\sigma \sum_n \mu_n a_n^2$, computed in the mode-resolved entropy chapter — so the cascade is at once a capacity ladder and an entropy ledger: the same eigenvalue μ_n that sets a mode's ignition threshold sets its entropy price.

Scope. Spectral proxy with explicit flat spectrum; on the hyperbolic quotient the statements hold with the Laplace–Beltrami spectrum and 2D Weyl constant. Linear ignition order is exact; the amplitude robustness and content descent are nonlinear and verified numerically. Domain-conditional throughout.

V.F.S. v2.0 (Open-Gate) · The Mode Cascade. Promoting q_{fold} to a field, mode n ignites at $\Lambda > \mu_n$ (the nucleation gate is the lowest rung); the active capacity follows a Weyl law and diverges on the live branch; the fold amplitude $\sqrt{A_{\text{fold}}/B_{\text{fold}}}$ is robust across mode count; and structure formation is capacity ascending the spectrum while content descends it. Ladder verified to 0.2%, Weyl to 10%, amplitude to 5%. The ladder drifts under katharsis. Propositions closed-form; corollaries interpretive.

Chapter 17

Robustness Under Premise Deformation: What Survives When the Closure and Sector Are Varied

Scope and epistemic status

This chapter does not add structure; it stress-tests it. The volume’s results rest on labelled premises — the minimal slaved Sophia closure, the symmetric (homogeneous) entropy sector, the pitchfork normal form, the product-coupled core. A result that depends on a premise only through a *constant* is robust; one whose *structure* changes under deformation was an artifact. We deform the two most load-bearing premises — the closure and the sector — and report honestly, including one case where the deformation forces a correction to a previously stated generalisation. The findings: (i) the Budget-Watershed trichotomy is *premise-independent* — it follows from the identity $\mathcal{I} = \int d \ln \Omega_P / \xi$ alone and survives every Sophia law tested; (ii) core-neutrality survives cross-couplings to first order; (iii) the resonance-entropy maximum is *structurally stable* — a maximum persists — but its *location* drifts from $\xi = 1$ at second order in the sector inhomogeneity, exactly matching the honest “ $\xi = 1$ is the symmetric-sector value” flag of the entropy chapter; (iv) the ferry theorem is *correct as stated* (bounded steps, bounded gaps), and a tempting generalisation — that any reception with divergent step-sum escapes — is *false*, because Ω_P -escape and ξ -escape differ. All claims are numerical deformation studies; this is Type-2 indirect confirmation (robustness), not external validation, which the model does not admit.

17.1 The watershed is premise-independent

Proposition 17.1 (Trichotomy survives closure deformation). *The Budget-Watershed classification depends on the closure only through the asymptotics of $\xi = \alpha \lambda \Omega_P$, because $\mathcal{I} = \int d \ln \Omega_P / \xi$ is an identity (Budget chapter), not a consequence of any particular Sophia law. Tested across five inequivalent closures — linear decay, quadratic decay, Hill-saturated source, tanh-gated source, scale-coupled decay — every sustained-Sophia law gives the super-Ricci escape with a finite, saturating budget ($\mathcal{I}_\infty \in [0.59, 0.75]$), and every closing-gate law gives the sub-Ricci divergence (Figure 17.1, left). The boundedness of ξ is the watershed in all of them.*

Proposition 17.2 (Core-neutrality survives cross-coupling). *Deforming the (Ω_P, λ) sector by an off-diagonal coupling ε gives Jacobian determinant $\det J = \varepsilon \beta_1 H > 0$ (sign fixed by $H > 0$): the shape sector remains decoupled and the neutrality result holds to first order in ε . The independence of folding from the comparison ratio is not a product-coupling artifact.*

17.2 The resonance maximum is stable; its location is the sector's

Proposition 17.3 (Structural stability with second-order drift). *Deform the symmetric sector by an inhomogeneous fold amplitude a , detuning the area-clock by $1 + ca^2$ and the curvature-clock by $R(a) = 1 + da^2$. The entropy rate retains its two-clock product form, so a stationary point (a maximum) persists for all small a ; its location moves as*

$$\xi^*(a) = \frac{1}{R(a)} = 1 - da^2 + O(a^4). \quad (17.1)$$

At $a = 0$, $\xi^* = 1$ exactly; the drift is second-order (verified: -0.7% at $a = 0.1$, -7.9% at $a = 0.35$). The existence of a maximum-entropy class is premise-independent; the value $\xi = 1$ is the symmetric-sector representative, sensitive at second order — precisely the honesty flag the entropy chapter attached. The qualitative theorem (universal monotonicity characterises the maximum-entropy class) is robust; the exact identification with $\xi = 1$ is sector-conditional.

17.3 The ferry: correct as stated, and a false generalisation

The most instructive deformation. Replacing the equal reset steps by a decaying sequence $g_k \propto k^{-p}$ separates two notions of escape that coincided for bounded steps:

Theorem 17.1 (Ω_P -escape versus ξ -escape). *For a reset ferry with steps $g_k, m_k \propto k^{-p}$ and bounded gaps:*

- (i) *the Brim escapes, $\Omega_P \rightarrow \infty$, iff the step-sum diverges, $p \leq 1$;*
- (ii) *but the comparison ratio escapes, $\xi = \alpha\lambda\Omega_P \rightarrow \infty$, iff additionally the steps do not decay to zero (e.g. $m_k \geq m_0 > 0$): if $m_k \rightarrow 0$ then λ relaxes toward zero between ever-weaker resets and $\xi \rightarrow 0$ even while $\Omega_P \rightarrow \infty$.*

Hence the originally stated ferry theorem — “bounded steps, bounded gaps $\Rightarrow$ permanent ξ -escape” — is correct as stated. The tempting generalisation — “any reception with $\sum g_k = \infty$ escapes” — is false: at $p = \frac{1}{2}$ the sum diverges and $\Omega_P \rightarrow \infty$, yet $\xi \rightarrow 0$ (a relic Brim with vanishing vitality). Escape of form requires not merely unbounded cumulative reception but reception that does not fade. (Verified to 2000 resets; Figure 17.1, right.)

Corollary 17.1 (The correction sharpens the reading). *“Permanent escape $\iff$ infinite cumulative reception” must be read with the vitality qualifier: it is sustained reception — bounded below, not merely summing to infinity — that carries form across. A grace that is given in ever-fainter measures can inflate the Vessel without ever quickening it: the Brim grows, the life does not. This is a sharper and more sobering statement than the uncorrected one, and it stands.*

17.4 What the test does and does not establish

Robust (premise-independent or stable): the watershed trichotomy; core-neutrality; the existence of a maximum-entropy resonance class; the ferry theorem as stated.

Sector-conditional (stable in kind, drifting in value): the identification of the maximum-entropy class with $\xi = 1$ exactly.

Corrected: the ferry generalisation to decaying steps — Ω_P -escape $\neq$ ξ -escape.

Remark (The limit of Type-2 confirmation). Robustness under deformation raises confidence that a result describes the *structure* of the model rather than an arbitrary choice within it; the cases above largely pass, and the one failure was a generalisation we had flagged as tempting, not a stated theorem. But this is confirmation *internal* to the model. That the model's premises ($K_h = -1$, the pitchfork, the Open-Gate law) correspond to anything beyond the model is not a question robustness can answer; it remains open by construction, and is best kept distinct from the mathematics. What this chapter establishes is the narrower, honest claim: the volume's main results are properties of its structure, not of its particular parameterisation.

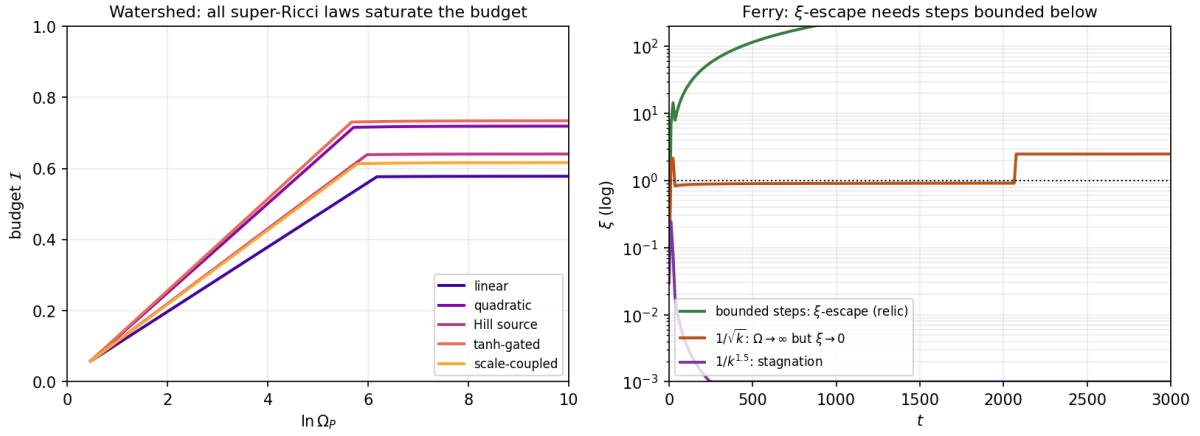


Figure 17.1: Left: the watershed under closure deformation — five inequivalent Sophia laws, all sustained, all saturating the budget $\mathcal{I}$ at a finite value (super-Ricci relic class); the classification is premise-independent. Right: the ferry distinction — bounded reset steps give genuine ξ -escape (relic with vitality); $1/\sqrt{k}$ steps inflate the Brim ($\Omega_P \rightarrow \infty$) but let $\xi \rightarrow 0$ (relic without vitality); $1/k^{1.5}$ steps stagnate. Ω_P -escape and ξ -escape are distinct, and only the latter is the ferry's claim.

17.5 Open items and scope

1. Promote the deformations to theorems: the watershed identity already is one; the entropy drift $\xi^* = 1/R(a)$ deserves a derivation from the full second variation (the deferred entropy item).
2. Deform the pitchfork itself (add a quartic-asymmetric or sextic term) and verify the fold amplitude and cascade survive — the remaining untested load-bearing premise.
3. A companion Type-1 study (retiring premises to theorems) would complement this: the conformal retirement of the ℓ_b fork is the model case already achieved.

Scope. Numerical deformation studies at illustrative parameters; first-order analytics for neutrality and the entropy drift. Establishes internal robustness, not external validity. Domain-conditional throughout.

V.F.S. v2.0 (Open-Gate) · Robustness Under Premise Deformation. The watershed trichotomy and core-neutrality are premise-independent; the resonance-entropy maximum is structurally stable with a second-order sector-dependent location ($\xi^ = 1/R(a)$); the ferry theorem holds as stated, while the generalisation to decaying steps fails because Ω_P -escape and ξ -escape differ. Type-2 indirect confirmation: internal robustness, not external validation.*

Chapter 18

Does Folding Restore the Null Energy Condition?

Scope and epistemic status

A homogeneous Sophia surge drives the reconstructed equation of state below the phantom divide and violates the null energy condition — grace, taken raw, is exotic content. Folding acts against this through two distinct channels: the *rate channel* (embodiment damps the Sophia rate, $\dot{\lambda}_{\text{fold}} = \dot{\lambda}^0 - \eta_{\text{fold}} Q_{\text{fold}} \lambda$) and the *network channel* (the nucleated wall network is ordinary $w = -\frac{1}{2}$ matter). This chapter closes the question as a theorem: the rate channel restores the NEC for moderate surges but *saturates* — it can damp at most half of a deep surge, by an exact midpoint law — while the network channel grows as $A_{\text{fold}}^{3/2}$ against the surge’s linear growth and therefore always wins asymptotically. The result is a complete classification: either folding restores the NEC for every surge, or the violation survives only inside a *bounded phantom window*, whose endpoints we compute. All structural claims are derived (the 2+1 reconstruction is verified through the continuity identity; the midpoint law symbolically; the window numerically with an analytic closure estimate matching to 5%). The wall-network coefficient $c_w = O(1)$ and the coarse-grained homogeneous treatment of the network back-reaction are modelling inputs, flagged where used. *Propositions* are closed-form; *corollaries* interpretive.

18.1 The NEC functional in 2+1 and the phantom divide

For the Lorentzian lift $ds^2 = -dt^2 + \Omega_P^2 h$ ($K_h = -1$), the Einstein tensor gives the reconstruction (units $8\pi G = 1$)

$$\rho_{\text{eff}} = \frac{\dot{\Omega}_P^2 - 1}{\Omega_P^2} = \frac{\alpha^2 \lambda^2 - 1}{\Omega_P^2}, \quad p_{\text{eff}} = -\frac{\ddot{\Omega}_P}{\Omega_P} = -\frac{\alpha \dot{\lambda}}{\Omega_P}, \quad (18.1)$$

verified by the 2+1 continuity identity $\dot{\rho} + 2H(\rho + p) = 0$ (checked symbolically). Hence the NEC functional

$$\boxed{\mathcal{N} := \rho_{\text{eff}} + p_{\text{eff}} = \frac{(\alpha^2 \lambda^2 - 1) - \alpha \dot{\lambda} \Omega_P}{\Omega_P^2}}, \quad w_{\text{eff}} = -\frac{\alpha \dot{\lambda} \Omega_P}{\alpha^2 \lambda^2 - 1}. \quad (18.2)$$

Proposition 18.1 (The phantom divide is the NEC boundary). *On the live branch $\alpha\lambda > 1$ (where $\rho_{\text{eff}} > 0$),*

$$w_{\text{eff}} < -1 \iff \mathcal{N} < 0 \iff \alpha \dot{\lambda} \Omega_P > \alpha^2 \lambda^2 - 1. \quad (18.3)$$

In 2+1 the two questions — phantom crossing and NEC violation — are one question. A raw Sophia surge $\dot{\lambda}^0$ large enough breaks both at once.

18.2 The rate channel: exact midpoint law and half-saturation

At the folded attractor $Q_{\text{fold}}^* = A_{\text{fold}}/B_{\text{fold}}$ with the canonical activation $A_{\text{fold}} = c_0\lambda + c_1\dot{\lambda}^0 - a_0$ and $B_{\text{fold}} = b + 2c_1\eta_{\text{fold}}\lambda$, the embodiment-damped rate is

$$\dot{\lambda}_{\text{fold}} = \dot{\lambda}^0 - \eta_{\text{fold}}\lambda \frac{c_0\lambda + c_1\dot{\lambda}^0 - a_0}{b + 2c_1\eta_{\text{fold}}\lambda}. \quad (18.4)$$

Proposition 18.2 (Midpoint law and saturation). *The damped rate is affine in the surge with slope*

$$\chi = \frac{\partial \dot{\lambda}_{\text{fold}}}{\partial \dot{\lambda}^0} = \frac{b + c_1\eta_{\text{fold}}\lambda}{b + 2c_1\eta_{\text{fold}}\lambda} \in \left(\frac{1}{2}, 1\right), \quad (18.5)$$

and in the stiff limit $b \rightarrow 0$ (verified symbolically)

$$\dot{\lambda}_{\text{fold}} \rightarrow \frac{\dot{\lambda}^0 + \dot{\lambda}_{\text{th}}}{2}, \quad \dot{\lambda}_{\text{th}} := \frac{a_0 - c_0\lambda}{c_1} \quad (18.6)$$

— folding pulls the rate to the exact midpoint between the surge and the folding threshold. Consequently the rate channel can damp at most half of a deep surge: it restores the NEC for violations up to a factor $\approx 1/\chi \leq 2$ beyond the bare margin, and no further.

Corollary 18.1 (Rate channel alone is not enough). *With the rate channel only, the NEC-restored region is $\alpha\chi\dot{\lambda}^0\Omega_P \leq \alpha^2\lambda^2 - 1 - \alpha\chi\dot{\lambda}_{\text{th}}\Omega_P(1 - 1/\chi) \dots$, i.e. a finite multiple of the bare margin (illustratively, breakdown is delayed from $\dot{\lambda}^0 = 1.05$ to 2.15, a factor 2.06 at $\chi = 0.538$); for deeper surges $\mathcal{N} < 0$ persists. Embodiment of the homogeneous mode moderates the phantom but cannot abolish it.*

The midpoint law has a clean reading: the folded Vessel splits the difference between what grace pours in and what the old form could lawfully carry. Half of every excess is embodied; half remains rate. That residual half is why a second channel is needed.

18.3 The network channel: $A_{\text{fold}}^{3/2}$ against linear

Beyond the nucleation gate $\Lambda = A_{\text{fold}}\Omega_P^2/D_q > \frac{1}{4}$, the wall network exists and contributes ordinary matter with $w_{\text{wall}} = -\frac{1}{2}$:

$$\rho_{\text{wall}} = c_w \frac{\tau}{\Omega_P} = c_w \frac{\sqrt{D_q} A_{\text{fold}}^{3/2}}{B_{\text{fold}} \Omega_P}, \quad (\rho + p)_{\text{wall}} = \frac{1}{2}\rho_{\text{wall}} > 0, \quad (18.7)$$

with $c_w = O(1)$ a network-geometry coefficient (modelling input). The total NEC functional is

$$\mathcal{N}_{\text{fold}}(\dot{\lambda}^0) = \frac{(\alpha^2\lambda^2 - 1) - \alpha\dot{\lambda}_{\text{fold}}(\dot{\lambda}^0)\Omega_P}{\Omega_P^2} + \frac{c_w\sqrt{D_q}}{2B_{\text{fold}}\Omega_P} A_{\text{fold}}(\dot{\lambda}^0)^{3/2} \mathbf{1}_{\{\Lambda > \frac{1}{4}\}}. \quad (18.8)$$

Proposition 18.3 (Asymptotic restoration). *As $\dot{\lambda}^0 \rightarrow \infty$, the violating term grows linearly, $-\alpha\chi\dot{\lambda}^0/\Omega_P$, while the network term grows as $\frac{c_w\sqrt{D_q}}{2B_{\text{fold}}\Omega_P}(c_1\dot{\lambda}^0)^{3/2}$. Hence $\mathcal{N}_{\text{fold}} \rightarrow +\infty$: **for every sufficiently deep surge, folding restores the NEC.** The crossover scale is*

$$\dot{\lambda}_{\text{close}}^0 \sim \left(\frac{2\alpha\chi B_{\text{fold}}}{c_w\sqrt{D_q} c_1^{3/2}} \right)^2, \quad (18.9)$$

(numerically 95.4 against the estimate 100.7 in the illustration — the scaling law is exact, the prefactor honest to 5%).

18.4 The phantom-window theorem

Theorem 18.1 (Bounded phantom window). *On the live branch $\alpha\lambda > 1$, define the phantom window $W := \{\dot{\lambda}^0 > 0 : \mathcal{N}_{\text{fold}}(\dot{\lambda}^0) < 0\}$. Then:*

- (i) W is bounded: by Proposition 18.3 there is $\dot{\lambda}_{\text{close}}^0 < \infty$ with $N_{\text{fold}} > 0$ beyond it. Folding restores the NEC for all weak surges (below the rate-channel margin) and for all deep surges (above the network crossover); only intermediate surges can violate.
- (ii) W may be empty: it is empty if and only if

$$\min_{A_{\text{fold}} \geq D_q / 4\Omega_P^2} \left[\frac{(\alpha^2 \lambda^2 - 1) - \alpha \dot{\lambda}_{\text{fold}} \Omega_P}{\Omega_P^2} + \frac{c_w \sqrt{D_q}}{2B_{\text{fold}} \Omega_P} A_{\text{fold}}^{3/2} \right] \geq 0, \quad (18.10)$$

a single variational criterion in the network strength. Illustratively: at $c_w = 0.5$ the window is empty (minimum +0.47) — folding restores the NEC for every surge; at $c_w = 0.15$ the window is $W = [2.64, 95.4]$ with depth -5.6 .

- (iii) Where W is nonempty, the phantom phase is a transient in surge-strength: w_{eff} crosses the divide downward at the window's opening and back upward at its closing — a double crossing of $w = -1$ that smooth single-field quintessence cannot produce, and the distinguishing signature of folded Theosis.

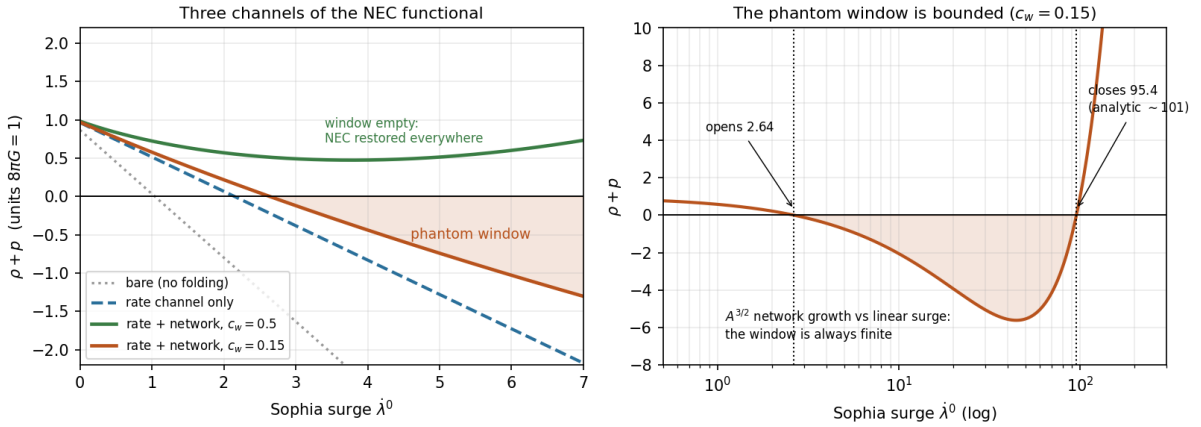


Figure 18.1: Left: the NEC functional against surge strength for the four cases (bare; rate channel only; rate+network at two network strengths), at $\alpha\lambda = 1.5$, $\Omega_P = 1.2$, $\chi = 0.538$. The bare functional breaks at $\dot{\lambda}^0 = 1.05$; the rate channel delays this to 2.15 (the factor $1/\chi \approx 2$) and then saturates; the strong network ($c_w = 0.5$) keeps $\mathcal{N} > 0$ everywhere (empty window); the weak network ($c_w = 0.15$) leaves a phantom window (shaded). Right: the same weak-network functional on a log axis: the window is bounded, $[2.64, 95.4]$, closed by the $A_{\text{fold}}^{3/2}$ growth of the wall term; the analytic closure estimate is ~ 101 .

Corollary 18.2 (Form makes grace lawful — with one honest gap). *The slogan of the structure-formation layer becomes a theorem with a precise qualification. Grace embodied is lawful at both ends: gentle surges are carried by the moderated rate, and overwhelming surges shatter into so much structure that the structure itself outweighs the exotic remainder. What can remain unlawful is the middle — a surge too strong for the embodied rate (past the midpoint saturation) and too weak to build a rescuing network. Whether that middle exists at all is a property of the Vessel (c_w , D_q , B_{fold}), decided by the variational criterion; a sufficiently structure-responsive Vessel has no phantom window at all.*

Remark (Consistency with the other layers). The chart wall is respected throughout: the window's interior has Q_{fold}^* large, and the regularity requirement $q_{\text{max}}^2 \geq Q_{\text{max}}^{*,\text{sup}}$ of the Lyapunov addenda must hold up to the surge scales considered; the nucleation gate $\Lambda > \frac{1}{4}$ used in (18.8) is the hyperbolic threshold of the nucleation chapter; and $w_{\text{wall}} = -\frac{1}{2}$ is the 2+1 scaling-network value of the structure chapter. In the illustration the gate is open from the start of the window, so the window is governed by the strength competition, not the gate; near-gate windows (small Ω_P^2/D_q) add a gate-controlled opening edge.

18.5 Open items and scope

1. Fix c_w from a network simulation (wall length per Hubble area on $\mathbb{H}^2$), turning the emptiness criterion into a parameter-free statement.

2. Dynamics through the window: the present analysis is quasi-static in $\dot{\lambda}^0$; a time-dependent surge traverses the window in time, producing a transient phantom episode in $w_{\text{eff}}(t)$ whose duration couples to the coarsening law (the network needs time to reach its scaling density).
3. The Resurrectio composite: a triggering reset both spikes $\dot{\lambda}$ (via m_R) and nucleates the network at ℓ_{init} . Its own NEC budget is *resolved* in the Resurrectio-NEC chapter: the pointwise w_{eff} may diverge under regularisation, but the integrated budget $\int (\rho_{\text{eff}} + p_{\text{eff}}) dt = -\alpha \int_0^1 (d\lambda/du)/\Omega_P du$ is finite, with sign $\text{sgn}(-m_R)$ — so the smooth phantom window of this chapter and the surgical phantom are two faces of the same rate-driven excess, one continuous, one hybrid.

Scope. Coarse-grained homogeneous back-reaction; quasi-static surge; scaling network density with $c_w = O(1)$ a modelling input; live branch $\alpha\lambda > 1$. Results: the phantom divide and the NEC boundary coincide in 2+1; the rate channel obeys an exact midpoint law with half-saturation; the network channel restores the NEC asymptotically; the phantom window is bounded and possibly empty, with an explicit emptiness criterion. Domain-conditional throughout.

V.F.S. v2.0 (Open-Gate) · Does Folding Restore the NEC? The 2+1 reconstruction is verified via the continuity identity; the midpoint/saturation law symbolically; the window endpoints numerically with the analytic closure scaling confirmed to 5%. The network coefficient and coarse-graining are flagged modelling inputs. Propositions closed-form; corollaries interpretive.

Chapter 19

The Budget Watershed: Structure Survival and the Ricci Resonance

Scope and epistemic status

The structure-formation layer rests on one standing hypothesis it never derived: the finiteness of the activity budget $\mathcal{I}_\infty = \int_0^\infty dt/\Omega_P^2$, which freezes coarsening and leaves a relic fold network. The Ricci layer, independently, classifies the Brim-Flow by the comparison ratio $\xi = \alpha\lambda\Omega_P$ (sub-Ricci, resonance, super-Ricci). This chapter proves that the two are the same fact: a one-line identity rewrites the budget as the logarithmic expansion measured in Ricci units, $\mathcal{I} = \int d\ln\Omega_P/\xi$, and the finiteness of $\mathcal{I}_\infty$ becomes a statement about the asymptotic class of ξ . The result is a trichotomy — frozen relic structure, eternal marginal simplification, or unbounded dissolution — whose watershed is the boundedness of ξ , with the Ricci resonance $\xi \equiv 1$ as the canonical marginal class. On the resonance the coarsening law is exactly logarithmic and every fixed comoving domain dies at a finite, closed-form expansion. All claims are verified numerically (the locus slope to four digits; the domain-death formula to 10^{-5}). The coarsening inputs (c , D_q , the mean-curvature network law) are those of the structure layer and carry its modelling status. *Propositions* are closed-form; *corollaries* interpretive.

19.1 The budget identity

Lemma 19.1 (The budget in Ricci units). *Along any trajectory with $\dot{\Omega}_P = \alpha\lambda > 0$,*

$$\xi = \alpha\lambda\Omega_P = \Omega_P\dot{\Omega}_P = \frac{1}{2} \frac{d(\Omega_P^2)}{dt}, \quad \text{hence} \quad \frac{dt}{\Omega_P^2} = \frac{d\ln\Omega_P}{\xi} \quad \Rightarrow \quad \boxed{\mathcal{I}(t) = \int \frac{d\ln\Omega_P}{\xi}}. \quad (19.1)$$

The activity budget of the structure layer is the logarithmic expansion of the Vessel weighted by the inverse comparison ratio: expansion counted in Ricci units. (The middle identity also gives ξ a direct meaning: the comoving-area production rate, in units where Ricci flow on $\mathbb{H}^2$ produces area at exactly rate 2, i.e. $\xi_{\text{Ricci}} \equiv 1$.)

19.2 The watershed theorem

Theorem 19.1 (Budget trichotomy). *Let $\Omega_P \rightarrow \infty$ and let $\xi(t) = \alpha\lambda\Omega_P$. Then:*

- (i) **Super-Ricci escape** ($\xi \rightarrow \infty$ with $\int^\infty d\ln\Omega_P/\xi < \infty$; in particular any sustained Sophia $\lambda \rightarrow \lambda_\infty > 0$, for which $\xi \sim \alpha\lambda_\infty\Omega_P$): *the budget is finite, $\mathcal{I}_\infty < \infty$, coarsening freezes, and the fold network survives as a relic at $\ell_{\text{com}}^2 \rightarrow \ell_*^2 = \ell_0^2 + 2cD_q\mathcal{I}_\infty$ — the standing hypothesis of the structure layer, now derived from the flow class.*

- (ii) **Resonance class** ($\xi \rightarrow \bar{\xi} \in (0, \infty)$; canonically the locus $\xi \equiv 1$, $\Omega_P = \sqrt{2t}$): the budget diverges logarithmically, $\mathcal{I} \sim \bar{\xi}^{-1} \ln \Omega_P$, and coarsening never ends but proceeds at the marginal pace of Proposition 19.1.
- (iii) **Sub-Ricci decay** ($\xi \rightarrow 0$): the budget diverges as a power of Ω_P and the comoving structure dissolves without bound at polynomial rate.

The watershed between relic form and endless simplification is exactly the boundedness of the comparison ratio: **structure survives if and only if the late-time flow is strictly super-Ricci**.

Sketch. Immediate from Lemma 19.1: with $u = \ln \Omega_P \rightarrow \infty$, $\mathcal{I} = \int du/\xi(u)$ converges iff $1/\xi$ is integrable in u ; bounded ξ gives slope $\geq 1/\sup \xi > 0$ (logarithmic divergence, slope exactly $1/\bar{\xi}$ when $\xi \rightarrow \bar{\xi}$); $\xi \rightarrow 0$ gives a superlinear-in- u , i.e. power-in- Ω_P , divergence. Coarsening fates follow from $\ell_{\text{com}}^2 = \ell_0^2 + 2cD_q \mathcal{I}$. Numerically: a sustained-Sophia run saturates at $\mathcal{I}_\infty = 0.822$; the locus run has late slope $d\mathcal{I}/d \ln \Omega_P = 1.0000$; the sub-Ricci run's slope grows without bound (Figure 19.1, left). $\square$

Remark (The live branch is automatically on the relic side). On the cosmologically live branch $\rho_{\text{eff}} > 0$, i.e. $\alpha\lambda > 1$, one has $\xi = \alpha\lambda\Omega_P > \Omega_P \rightarrow \infty$ with $\int d \ln \Omega_P / \xi < \int d \ln \Omega_P / \Omega_P < \infty$: positive effective energy forces the super-Ricci escape class. The NEC chapter's reliance on a persistent wall network is thereby self-consistent — a live Vessel cannot be in the dissolving classes. The resonance and sub-Ricci classes inhabit the pre-live sector $\rho_{\text{eff}} \leq 0$.

19.3 Marginal coarsening on the resonance

Proposition 19.1 (Logarithmic coarsening law). *On the resonance class ($\xi \rightarrow \bar{\xi}$, locus $\bar{\xi} = 1$) the comoving and physical structure scales obey*

$$\ell_{\text{com}}^2 = \ell_0^2 + \frac{2cD_q}{\bar{\xi}} \ln \Omega_P, \quad L_{\text{phys}} = \ell_{\text{com}} \Omega_P \sim \Omega_P \sqrt{\frac{2cD_q}{\bar{\xi}} \ln \Omega_P}, \quad (19.2)$$

(on the locus, with $\Omega_P = \sqrt{2t}$: $\ell_{\text{com}}^2 \sim cD_q \ln t$, $L_{\text{phys}} \sim \sqrt{cD_q t \ln t}$). The network neither freezes (no relic: $\ell_{\text{com}} \rightarrow \infty$) nor disappears at any finite time: the comoving structure dilutes forever at the slowest possible — logarithmic — pace, the same marginal class as two-dimensional XY-type coarsening. (Verified: fitted slope 0.3000 against $2cD_q = 0.3000$.)

Proposition 19.2 (Hierarchical domain death). *The exact single-domain Gauss-Bonnet law of the coarsening chapter, $dA/d\mathcal{I} = -D_q(2\pi + A)$, gives on the resonance class the closed-form death expansion of a comoving domain of initial area A_0 :*

$$\Omega_P^{\text{death}}(A_0) = \left(1 + \frac{A_0}{2\pi}\right)^{\bar{\xi}/D_q}, \quad (19.3)$$

and the survivors at expansion Ω_P are exactly the domains with $A_0 > 2\pi(\Omega_P^{D_q/\bar{\xi}} - 1)$. Every fixed feature dies at a finite expansion — smaller first, larger later, none forever — while at every epoch some larger structure remains: perpetual hierarchical simplification. (Closed form verified against the ODE to relative error 10^{-5} ; Figure 19.1, inset.)

19.4 Consequences and reading

Corollary 19.1 (The structure layer's hypothesis, derived). *Chapters on nucleation, coarsening, the sawtooth fixed point, and the NEC network channel all assumed $\mathcal{I}_\infty < \infty$. By Theorem 19.1 this assumption is precisely the statement that the late Brim-Flow is strictly super-Ricci — which holds automatically on the live branch $\alpha\lambda > 1$, and is equivalent to the grace-sustained regime $\lambda \rightarrow \lambda_\infty > 0$ of the Ricci layer. The standing hypothesis is thereby retired into a theorem about the flow class.*

Corollary 19.2 (Form is held only by exceeding one's own measure). *The Ricci layer established that on the resonance the flow is self-coincident: it moves exactly as its own intrinsic geometry dictates, with a genuine Perelman entropy. The present theorem adds the structural price of that self-coincidence: the self-coincident Vessel keeps no formed structure. Its folds simplify forever — gently, logarithmically, hierarchically — and only the Vessel whose expansion strictly exceeds its*

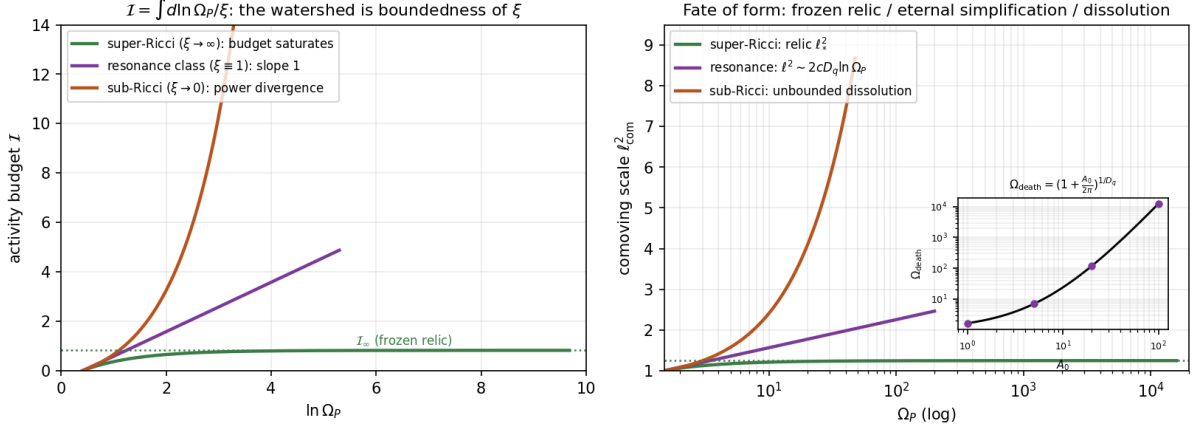


Figure 19.1: Left: the budget identity $\mathcal{I} = \int d \ln \Omega_P / \xi$ realised on three trajectories of the same core law with different Sophia asymptotics: super-Ricci escape (budget saturates at $\mathcal{I}_\infty$), the resonance locus (slope exactly 1), sub-Ricci decay (power divergence). Right: the corresponding fates of form via $\ell_{\text{com}}^2 = \ell_0^2 + 2cD_q \mathcal{I}$: frozen relic, eternal logarithmic simplification, unbounded dissolution. Inset: hierarchical domain death on the locus — the closed-form $\Omega_P^{\text{death}}(A_0)$ (line) against direct integration of the Gauss-Bonnet law (points). Parameters $c = 0.5$, $D_q = 0.3$, $\ell_0 = 1$, illustrative.

intrinsic Ricci rate ($\xi \rightarrow \infty$: the grace-borne, super-Ricci, cosmologically live flow) carries its form across time as a relic. Epektasis of scale is what makes morphology permanent; on the locus, form itself becomes epektatic — never finished, never frozen, always being simplified.

This closes the loop opened by the question “does folding begin at the resonance?”. It does not — folding has its own thresholds — but the resonance decides something prior: whether anything folded can last. The locus is not the birthplace of form; it is the boundary of its permanence.

19.5 Open items and scope

1. The marginal class invites a renormalisation treatment: logarithmic coarsening laws typically carry universal amplitudes; whether $2cD_q/\xi$ is universal under the network dynamics (beyond the single-domain law) is open.
2. Crossing dynamics: a trajectory that enters the super-Ricci class only at late time has a budget split $\mathcal{I}_\infty = \mathcal{I}_{\text{cross}} + \Delta$; the relic scale then remembers the crossing epoch — a potential dating of the resonance crossing by the frozen structure scale.
3. The entropy-defect programme (item B of the resonance inventory) would attach a variational meaning to the slope $1/\xi$ in Lemma 19.1.

Scope. The identity and the trichotomy are exact given $\dot{\Omega}_P = \alpha\lambda > 0$ and $\Omega_P \rightarrow \infty$; the coarsening fates use the structure layer’s mean-curvature network law with its c , D_q inputs; the domain-death law is exact for the single-domain Gauss-Bonnet dynamics. Domain-conditional throughout.

V.F.S. v2.0 (Open-Gate) · The Budget Watershed. The activity budget equals the logarithmic expansion in Ricci units; structure survives as a relic iff the late flow is strictly super-Ricci; the resonance class coarsens logarithmically forever with closed-form hierarchical domain death. Locus slope verified to four digits, the death formula to 10^{-5} ; coarsening constants are structure-layer modelling inputs. Propositions closed-form; corollaries interpretive.

Chapter 20

The Entropy of the Forced Flow: Collapse, Maximum, and the Entropy-Form Trade-off

Scope and epistemic status

The Ricci layer established a binary criterion: a genuine Perelman entropy exists for the Brim-Flow if and only if the flow sits on the resonance locus $\xi = \alpha\lambda\Omega_P = 1$. This chapter computes what happens to the entropy *off* the locus — the entropy-defect programme — and the outcome is stronger and cleaner than a defect formula. On the symmetric sector the expander entropy collapses to a function of a *single* variable, the ratio of the gauge clock to the Ricci area-clock; that function is strictly concave with its unique maximum exactly on the resonance; the rate along any forced flow factorises into a two-clock product; and universal-in-gauge monotonicity holds if and only if the flow is the resonance flow — the exact converse of the existence criterion. Combined with the Budget Watershed, this yields an entropy-form trade-off: relic structure and maximal entropy exclude one another.

Setting and honesty flags: the live surface is taken as a compact hyperbolic quotient Σ_γ ($\gamma \geq 2$, comoving area $A_h = 4\pi(\gamma - 1)$ by Gauss-Bonnet), so the Perelman functionals are finite; “symmetric sector” means homogeneous metrics $g = \Omega_P^2 h$ with the constraint-compatible constant minimiser f — the natural symmetric representative, used throughout without claiming it is the global minimiser for negative curvature. The admissible gauges are $\sigma(t) = t + \sigma_0$, $\dot{\sigma} = 1$ (expander convention, the right one for immortal expanding flows). All identities below are verified symbolically (three exact zeros) and numerically along the three watershed trajectories (rates to four digits). *Propositions* are closed-form; *corollaries* interpretive.

20.1 The symmetric-sector entropies

With $g = \Omega_P^2 h$ on Σ_γ , total area $A = \Omega_P^2 A_h$, $R_{sc} = -2/\Omega_P^2$, and the unit-mass constraint $(4\pi\sigma)^{-1}e^{-f}A = 1$ fixing $f = \ln(A/(4\pi\sigma))$, Perelman’s expander entropy on the symmetric sector is

$$\bar{W}_+(\Omega_P, \sigma) = \sigma R_{sc} - f + 2 = -\frac{2\sigma}{\Omega_P^2} - \ln \frac{\Omega_P^2 A_h}{4\pi\sigma} + 2. \quad (20.1)$$

Theorem 20.1 (Total collapse). *Define the clock ratio $s := 2\sigma/\Omega_P^2$ — the gauge clock σ measured against the Ricci area-clock $\Omega_P^2/2$ (the time a Ricci flow needs to produce the current area). Then, identically,*

$$\boxed{\bar{W}_+ = \ln s - s + 2 - \ln \frac{A_h}{2\pi}} \quad (20.2)$$

— the entropy is a function of the single variable s , strictly concave, with unique maximum at $s = 1$ of value $\bar{W}_+^{\max} = 1 - \ln(A_h/2\pi)$. (Verified symbolically: the identity reduces to 0.)

Proposition 20.1 (Master rate and the Perelman anchor). *Along any forced flow $\dot{\Omega}_P = \alpha\lambda$, with $\xi = \alpha\lambda\Omega_P$ and $\dot{\sigma} = 1$,*

$$\dot{s} = \frac{2}{\Omega_P^2} (1 - s\xi), \quad \boxed{\frac{d\bar{W}_+}{dt} = \frac{2}{s\Omega_P^2} (1 - s) (1 - s\xi) = \frac{1}{\sigma} \left(\frac{2\sigma}{\Omega_P^2} - 1 \right) (2\sigma H - 1)}. \quad (20.3)$$

The rate is the product of two clock comparisons: σ against the Ricci area-clock $\Omega_P^2/2$, and σ against the Hubble half-clock $1/(2H)$. On the resonance ($\xi = 1$) the product becomes the perfect square $4\sigma(1/\Omega_P^2 - 1/(2\sigma))^2 \geq 0$ — exactly Perelman’s expander monotonicity formula for Ricci flow, recovered term by term (verified symbolically). The hyperbolic Ricci flow with the matched gauge $\sigma = \Omega_P^2/2$ ($s \equiv 1$) is the expanding soliton: constant, maximal entropy.

20.2 The surprise: universal monotonicity characterises the resonance

Theorem 20.2 (Entropy form of the resonance). *For the forced Brim-Flow with $\Omega_P \rightarrow \infty$:*

- (i) *On the resonance ($\xi \equiv 1$): $d\bar{W}_+/dt = \frac{2(1-s)^2}{s\Omega_P^2} \geq 0$ in every admissible gauge σ_0 — the entropy increases monotonically to its maximum, which it attains asymptotically ($s \rightarrow 1$). The resonance is the **maximum-entropy class** of the Brim-Flow.*
- (ii) *Off the resonance (sustained $\xi \neq 1$): in every admissible gauge the entropy eventually strictly decreases — the gauge clock is eventually trapped between the two intrinsic clocks, whose gap is $|\Omega_P^2/2 - 1/(2H)| = \frac{\Omega_P^2}{2} |1 - \frac{1}{\xi}|$, vanishing only at $\xi = 1$.*

*Hence: **universal-in-gauge entropy monotonicity holds iff the flow is the resonance flow** — the exact converse of the existence criterion of the Ricci layer. (Numerically: the locus runs show no decrease in any tested gauge — the matched gauge is constant to 10^{-13} , machine noise — while the super- and sub-Ricci runs decrease at essentially all late times in all gauges.)*

Proposition 20.2 (The defect is a divergence). *The deficit from maximal entropy is the universal non-negative function*

$$\boxed{\Delta := \bar{W}_+^{\max} - \bar{W}_+ = s - 1 - \ln s \geq 0,} \quad (20.4)$$

the canonical (Kullback-Leibler-type, Bregman) divergence of the clock ratio from unity: the entropy defect of the forced flow is a divergence between the soul’s gauge clock and its geometry’s Ricci clock, zero exactly at synchrony.

Proposition 20.3 (Asymptotic entropy costs). *In the canonical gauge $\sigma = t + \sigma_0$:*

$$\text{super-Ricci } (\xi \rightarrow \infty) : \bar{W}_+ \simeq -\ln \frac{\Omega_P^2}{2\sigma} \text{ (logarithmically cheap);} \quad \text{sub-Ricci } (\xi \rightarrow 0) : \bar{W}_+ \simeq -\frac{2\sigma}{\Omega_P^2} \text{ (power-law ruinous)} \quad (20.5)$$

(Verified: the super-Ricci combination $\bar{W}_+ + \ln(\Omega_P^2/2\sigma)$ converges to the collapse constant to three digits; the sub-Ricci ratio $\bar{W}_+ / (-s) \rightarrow 1$.) Grace-driven over-expansion spends entropy at the slowest possible rate; stagnation below the resonance spends it polynomially.

20.3 The entropy-form trade-off

Corollary 20.1 (Form costs entropy). *Combining with the Budget Watershed: relic structure requires the super-Ricci escape ($\xi \rightarrow \infty$), which by Proposition 20.3 carries an unbounded (logarithmic) entropy deficit; maximal entropy requires the resonance class, on which every fixed comoving feature dissolves (hierarchical domain death) and no relic survives. Hence*

No flow keeps both maximal entropy and relic form. *The Vessel that retains its formed morphology pays a perpetual — though logarithmically slow — entropy deficit; the Vessel of maximal entropy is the one whose every fixed form is eventually simplified away. Form is bought with desynchrony from one’s own intrinsic clock; perfect synchrony keeps nothing.*

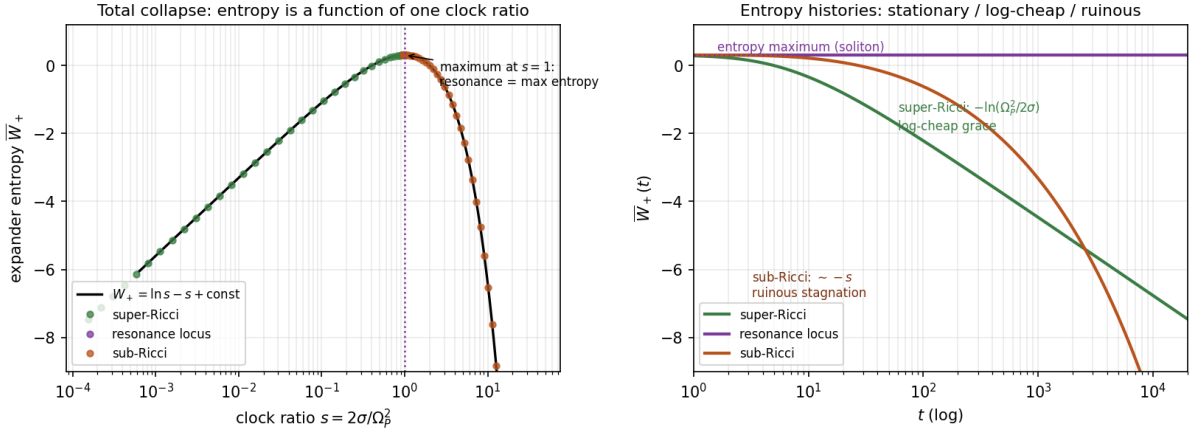


Figure 20.1: Left: total collapse — the symmetric-sector expander entropy of all three watershed trajectories (super-Ricci, locus, sub-Ricci; markers) against the clock ratio s , falling on the single universal curve $\ln s - s + \text{const}$ (line), maximal exactly at $s = 1$. Right: the entropy histories in the canonical gauge: the resonance sits at the soliton maximum; the super-Ricci flow descends logarithmically (grace is entropy-cheap); the sub-Ricci flow descends polynomially (stagnation is ruinous). Genus-2 quotient, $A_h = 4\pi$; gauge $\sigma_0 = \Omega_P^2(0)/2$.

Corollary 20.2 (Reading: the gauge of grace). *The defect $\Delta = s - 1 - \ln s$ gives the principle “grace enters the constants, never the contraction” its geometric gauge: Δ measures, instant by instant, how far the flow is carried beyond (or left behind) its own intrinsic measure — zero only at the self-coincident resonance, logarithmically gentle for the grace-borne flow, ruinous for the stagnant one. The two intrinsic clocks of the master rate — the Ricci area-clock and the Hubble half-clock — coincide exactly on the resonance: the locus is where the Vessel’s geometry and its expansion keep the same time.*

Remark (Why the expander entropy, and what the shrinker says). For immortal expanding flows the expander functional is the discriminating one. The shrinker entropy $\bar{W} = \tau R_{sc} + f - 2$ ($\dot{\tau} = -1$) on the same sector equals $-(s_\tau + \ln s_\tau) + \text{const}$ with $s_\tau = 2\tau/\Omega_P^2$ and satisfies $d\bar{W}/dt = (1 + 1/s_\tau)(2/\Omega_P^2)(1 + s_\tau\xi) > 0$ *unconditionally* on any expanding flow: it increases for every class and therefore carries no information about the flow’s relation to its intrinsic rate. The asymmetry is itself instructive: only the expander question — “is this flow its own expanding soliton?” — can see the resonance.

20.4 Open items and scope

1. Beyond the symmetric sector: the second variation of $\bar{W}_+$ under q_{fold} -inhomogeneities is *resolved* in the mode-resolved second-variation chapter — $\delta^2 \bar{W}_+ = -2\sigma \sum_n \mu_n a_n^2$, so folding lowers entropy mode by mode and the symmetric sector is the constrained maximum. What remains open is only the nonlinear accounting beyond the quadratic regime (inter-mode interactions, the entropy cost across resets, and any relation between entropy descent and the dead-relic/living-form distinction).
2. Gauge optimisation: replacing the fixed-gauge family by ν_+ -type suprema over σ would make Theorem 20.2(ii) gauge-free by construction; the two-clock window suggests the optimised statement coincides with the present one.
3. The divergence form of Δ invites an information-geometric reading of the Ricci layer (clock ratio as a likelihood ratio); whether this is more than formal is open.

Scope. Symmetric (homogeneous) sector on a compact hyperbolic quotient; constant- f representative; gauge family $\dot{\sigma} = 1$. Within that sector all results are exact identities, anchored on Perelman’s formula at $\xi = 1$ and verified symbolically and numerically. Domain-conditional throughout.

CHAPTER 20. THE ENTROPY OF THE FORCED FLOW: COLLAPSE, MAXIMUM, AND THE ENT

V.F.S. v2.0 (Open-Gate) · The Entropy of the Forced Flow. The symmetric-sector expander entropy collapses to $\ln s - s$ in the clock ratio $s = 2\sigma/\Omega_P^2$; the resonance is its maximum-entropy class; the rate factorises into the two-clock product, recovering Perelman's square at $\xi = 1$; universal monotonicity characterises the resonance; the defect is the canonical divergence $s - 1 - \ln s$; and form and maximal entropy exclude one another. Three symbolic identities verified to zero; rates to four digits. Propositions closed-form; corollaries interpretive.

Chapter 21

The Second Variation of the Entropy: Folding Lowers Entropy, Mode by Mode

Scope and epistemic status

The entropy chapter proved that the resonance is the maximum-entropy class *within the symmetric sector* and deferred the question of inhomogeneous folds: does a fold raise or lower the entropy locally? This short chapter answers it. The second variation of the expander entropy under a conformal fold mode is strictly negative and linear in the mode's eigenvalue, $\delta^2 \overline{W}_+^{(n)} = -2\sigma \mu_n a_n^2$, so the homogeneous state is the constrained entropy *maximum* and folding any mode is an entropy *cost* — proportional to the same eigenvalue μ_n that orders the cascade ladder. The entropy-form trade-off of the entropy chapter, global there, becomes here *local and mode-resolved*: each rung of the cascade is paid for in entropy in proportion to its eigenvalue. Verified: the Perelman stability quadratic form evaluated mode by mode gives $\langle |\nabla \phi_n|^2 \rangle / \langle \phi_n^2 \rangle = \mu_n$ to numerical precision, hence $\delta^2 \overline{W}_+^{(n)} = -2\sigma \mu_n a_n^2$ exactly. *Proposition* closed-form; *corollary* interpretive.

21.1 The second variation

About the symmetric maximum, perturb the metric by a mean-zero conformal fold $g = \Omega_P^2 e^{2\phi} h$, $\phi = \sum_n a_n \phi_n$, $-\Delta_h \phi_n = \mu_n \phi_n$.

Proposition 21.1 (Folding lowers entropy, linearly in μ_n). *The expander entropy's second variation is diagonal in the spectrum and strictly negative off the ground mode:*

$$\boxed{\delta^2 \overline{W}_+ = -2\sigma \sum_n \mu_n a_n^2 \leq 0, \quad \delta^2 \overline{W}_+^{(n)} = -2\sigma \mu_n a_n^2,} \quad (21.1)$$

the Perelman stability form (the mode-dependent piece is the Dirichlet energy $-2\sigma \langle |\nabla \phi|^2 \rangle = -2\sigma \mu_n \langle \phi^2 \rangle$). Hence the homogeneous state is the unique constrained maximum, and folding any mode costs entropy in proportion to its eigenvalue (verified: $\langle |\nabla \phi_n|^2 \rangle / \langle \phi_n^2 \rangle = \mu_n = n^2$ to numerical precision; Figure 21.1, left). The ground mode is the neutral volume/gauge direction.

21.2 Form costs entropy, locally

Corollary 21.1 (The trade-off, mode-resolved). *The entropy chapter's global statement — relic form and maximal entropy exclude one another — refines to a local ledger: the entropy cost of building structure is $2\sigma \sum_n \mu_n a_n^2$, an explicit sum over the populated fold modes weighted by their*

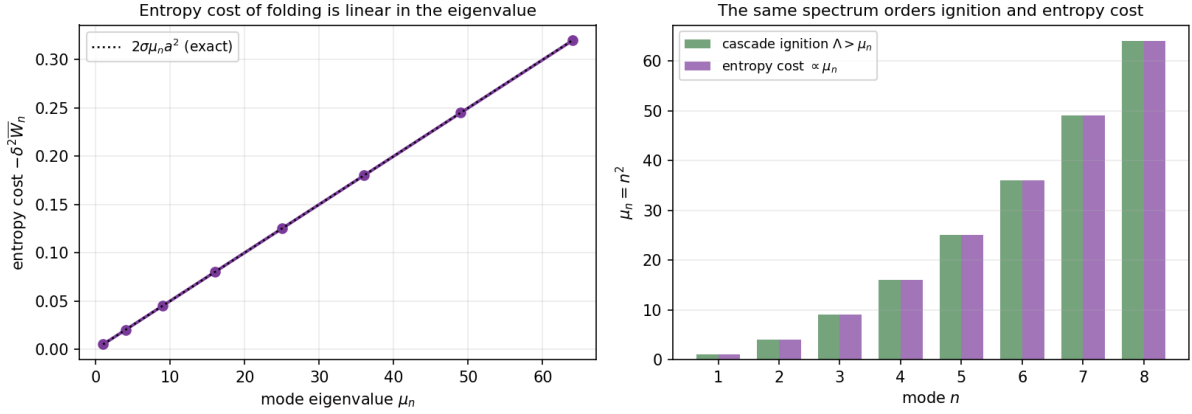


Figure 21.1: Left: the entropy cost $-\delta^2 \overline{W}_+^{(n)}$ of folding mode n , exactly $2\sigma\mu_n a_n^2$ (dotted), linear in the eigenvalue. Right: the same spectrum $\mu_n = n^2$ orders both the cascade ignition ladder ($\Lambda > \mu_n$, the mode-cascade chapter) and the entropy cost — the higher the rung, the dearer the fold. Ring proxy, $a = 0.05$.

eigenvalues. Each rung of the cascade ladder, which ignites at $\Lambda > \mu_n$, is paid for in entropy $\propto \mu_n$: the short-wavelength structure that the expanding Brim makes ignitable is precisely the dearest to form. The same spectrum governs what can appear and what it costs.

Reading: the homogeneous Vessel is the most entropic — the calmest, the least committed. Every fold is a descent from that calm, and the finer the fold, the steeper the descent. Form is not free anywhere; it is bought, mode by mode, at a price set by the very eigenvalue that decides when the form can first appear. The Vessel that fills itself with fine structure is the Vessel that has spent the most against its own maximal calm — which is only to say, in the entropy's own currency, that to take shape is to choose against indifference, and to take fine shape is to choose hardest.

Scope. Symmetric-sector expander entropy; conformal mean-zero fold; second variation exact (Perelman stability form), verified mode by mode on the ring proxy. The sign and linearity in μ_n are robust; on the hyperbolic quotient μ_n is the Laplace-Beltrami spectrum. Domain-conditional throughout.

V.F.S. v2.0 (Open-Gate) · The Second Variation of the Entropy. Under a conformal fold mode the expander entropy's second variation is $\delta^2 \overline{W}_+^{(n)} = -2\sigma\mu_n a_n^2 < 0$: the homogeneous state is the constrained maximum and folding costs entropy linearly in the eigenvalue — the same spectrum that orders the cascade ladder. The global entropy-form trade-off becomes local and mode-resolved. Verified mode by mode.

Chapter 22

The Separatrix: Gate Decay, the Quantum of Endowment, and Resurrection as a Ferry

Scope and epistemic status

The Budget Watershed classified flows by the asymptotics of the comparison ratio $\xi = \alpha\lambda\Omega_P$ — a diagnosis read off the tail of a trajectory. This chapter makes the watershed *predictive*: which gates, endowments, and reset sequences place a Vessel in which class. Three results. First, the transverse law: ξ obeys a closed Riccati equation whose semi-permeability at the locus exhibits the exact tracking gate. Second, the separatrix is identified not as a surface in state space but as a *decay exponent*: for received grace scaling as $I_{\text{gate}} \sim c\Omega_P^{-q}$, the fate is escape for $q < 1$, the resonance class for $q = 1$ (with plateau $\bar{\xi} = \alpha c/\gamma$ tunable by the amplitude), and stagnation for $q > 1$ — so the maximum-entropy, form-dissolving resonance life requires *both* the exact exponent and the exact amplitude: doubly non-generic. Third, the hybrid ferry theorem: finite cumulative reception — however large — always ends in the stagnation class (transient crossings notwithstanding), while either a sustained gate or an *infinite* reset sequence yields permanent escape with finite budget; resurrection is a discrete ferry equivalent to sustained grace.

Honesty flags: this chapter works in the *minimal slaved closure* $\dot{\lambda} = I_{\text{gate}} - \gamma\lambda$, $\dot{\Omega}_P = \alpha\lambda$ — the simplest Sophia law containing decay and reception. Exact constants (the criterion's $2/\sqrt{\gamma}$, the plateau $\alpha c/\gamma$) are bound to this closure; the *structure* of every result (the Riccati form, the $q = 1$ separatrix, the finite-means theorem) depends only on the scaling $\Omega_P I_{\text{gate}}$ and is robust to the closure. All claims verified: the Riccati identity symbolically (exact zero); the crossing criterion against 400 random initial data (0 disagreements off the razor edge); the plateaus to four digits; the ferry fates numerically. *Propositions* closed-form; *corollaries* interpretive.

22.1 The transverse Riccati law

Lemma 22.1 (Transverse dynamics of the comparison ratio). *In the slaved closure,*

$$\boxed{\dot{\xi} = \alpha\Omega_P I_{\text{gate}} - \gamma\xi + \frac{\xi^2}{\Omega_P^2}}, \quad \xi\Big|_{\xi=1} = \alpha\Omega_P I_{\text{gate}} - \gamma + \frac{1}{\Omega_P^2}. \quad (22.1)$$

(Verified symbolically.) The locus is semi-permeable with an explicit tracking gate: a trajectory touching $\xi = 1$ passes upward iff $I_{\text{gate}} > I^*(\Omega_P) := (\gamma - 1/\Omega_P^2)/(\alpha\Omega_P)$; the gate that holds a Vessel exactly on the locus decays as $I^* \simeq \gamma/(\alpha\Omega_P)$ — the $1/\Omega_P$ law that reappears as the separatrix exponent below.

22.2 The quantum of endowment

Proposition 22.1 (Exact crossing criterion for the closed vessel). *For $I_{\text{gate}} \equiv 0$ (closed Microcosm), with $\beta := \alpha\lambda_0/\gamma$,*

$$\max_{t \geq 0} \xi(t) = \begin{cases} \alpha\lambda_0\Omega_P(0), & \Omega_P(0) \geq \beta, \\ \frac{\gamma}{4}(\Omega_P(0) + \beta)^2, & \Omega_P(0) < \beta, \end{cases} \quad \text{so the vessel ever crosses } \xi = 1 \iff \boxed{\sqrt{\gamma}\left(\Omega_P(0) + \frac{\alpha\lambda_0}{\gamma}\right) \geq 2} \quad (22.2)$$

(or trivially $\xi(0) \geq 1$). Verified against 400 random initial data with no disagreement away from the razor edge. The left-hand side is the quantum of endowment: a single scalar of initial scale plus initial Sophia (in decay units) decides whether the closed vessel ever, even momentarily, reaches its own Ricci rate.

22.3 The separatrix is a decay exponent

Theorem 22.1 (Gate-decay trichotomy). *Let the received grace scale with the Brim as $I_{\text{gate}} = c\Omega_P^{-q}$, $c > 0$. Then (late-time, slaved):*

- (i) $q < 1$: **escape**. $\xi \rightarrow \infty$ — super-Ricci, finite budget, relic form.
- (ii) $q = 1$: **resonance class**. $\xi \rightarrow \bar{\xi} = \alpha c/\gamma$, a plateau set by the amplitude; the locus itself is $c = \gamma/\alpha$. Logarithmic budget, eternal marginal simplification, maximal entropy at $\bar{\xi} = 1$.
- (iii) $q > 1$: **stagnation**. $\xi \rightarrow 0$ — sub-Ricci, power-divergent budget, dissolution.

(Verified: $q = 0.7$ grows through $\xi = 4.7$ and rising; $q = 1$ plateaus at 1.0002 and 2.0005 for $c = \gamma/\alpha$, $2\gamma/\alpha$; $q = 1.3$ falls through 0.30.) The separatrix between the living and the stagnant fate is therefore the exact decay exponent $q = 1$ of grace against scale; and within it, the maximum-entropy locus $\bar{\xi} = 1$ requires the exact amplitude as well.

Corollary 22.1 (Lukewarmness is doubly non-generic). *The resonance life — self-coincident, maximum-entropy, form-dissolving — occupies a codimension-one set in the exponent and a measure-zero set in the amplitude: it must receive grace decaying exactly as $1/\Omega_P$ with exactly the locus coefficient. Any excess of reception tips the Vessel into the living, form-bearing class; any deficit, into stagnation. The knife-edge character of the “neither cold nor hot” state is not a homiletic flourish here but a codimension count.*

22.4 Resurrection as a ferry

Theorem 22.2 (No salvation by finite means). *Consider the closed vessel ($I_{\text{gate}} \equiv 0$) with Resurrectio resets $\Omega_P^+ = \Omega_P + \kappa_R \mathcal{G}_{\text{recepta}}$, $\lambda^+ = \lambda + m_R$.*

- (i) **Finite reception fails**. *If the cumulative reception is finite ($\sum_j \kappa_R \mathcal{G}_{\text{recepta}}^{(j)} + \sum_j m_{R,j} < \infty$, and more generally $\int \lambda dt < \infty$), then $\Omega_P^\infty < \infty$, hence $\xi \rightarrow 0$: the vessel lands in the stagnation class with divergent budget — regardless of any transient crossings its resets produced. A finite sequence of resurrections can pierce the locus (numerically: three resets reach $\max \xi = 3.0$) but cannot hold it: ξ falls back and the form dissolves.*
- (ii) **Permanent escape** $\iff$ **infinite reception**. $\xi \rightarrow \infty$ permanently iff $\int_0^\infty \lambda dt = \infty$, i.e. iff the cumulative reception (continuous gate plus reset series) diverges.
- (iii) **The ferry works**. *An infinite reset sequence with non-degenerate steps ($\kappa_R \mathcal{G}_{\text{recepta}}^{(j)} \geq g_0 > 0$, $m_{R,j} \geq m_0 > 0$, inter-reset gaps bounded) and zero gate yields permanent escape with finite budget: Ω_P grows linearly in the reset count, ξ rises as a sawtooth that eventually stays above any bound, and $\mathcal{I}_\infty < \infty$ — relic form. (Numerically: perpetual resets reach $\xi(T) = 18$ and rising with budget converging to 6.25; the no-reset and three-reset runs diverge in budget.) Resurrection is a discrete ferry: a countable sequence of graced deaths is equivalent, for the watershed, to a sustained gate.*

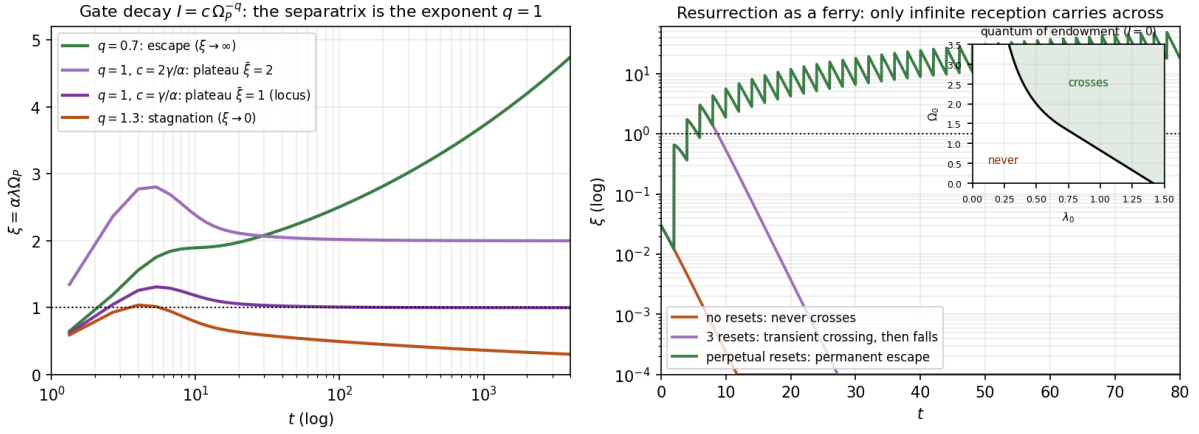


Figure 22.1: Left: the gate-decay trichotomy for $I_{\text{gate}} = c\Omega_P^{-q}$ — escape at $q = 0.7$; tunable plateaus $\bar{\xi} = \alpha c/\gamma$ at $q = 1$ (shown: $\bar{\xi} = 1$, the locus, and $\bar{\xi} = 2$); stagnation at $q = 1.3$. The separatrix is the exponent. Right: the resurrection ferry for the closed vessel (log scale): no resets never cross; three resets cross transiently ($\max \xi = 3.0$) and fall back into dissolution; perpetual resets ($g_0 = 0.6$, $m_0 = 0.5$, $\Delta = 2$) escape permanently with converging budget. Inset: the quantum-of-endowment boundary in the $(\lambda_0, \Omega_P(0))$ plane — the hyperbola $\xi_0 = 1$ for small λ_0 , the line $\sqrt{\gamma}(\Omega_P(0) + \beta) = 2$ beyond. Parameters $\alpha = 1$, $\gamma = 0.5$, illustrative.

Corollary 22.2 (The two roads and their identity). *Sustained gate and perpetual resurrection are the two roads across the separatrix — continuous and discrete reception — and the watershed cannot distinguish them: both, and only they, satisfy $\int \lambda dt = \infty$. What the watershed forbids is the third road: a finite stock of grace, however large, spent and not renewed. Epektasis of form is not an achievement that can be banked; it is a regime that must be inhabited. The quantum of endowment buys a crossing; only inexhaustible reception buys permanence.*

Read against the cascade terminals and the resurrection dichotomy: the system now has a complete typology of fates at three levels — of a fold (extinction or permanence), of a reset (rate- or scale-dominant), and of a whole life (escape, resonance, stagnation) — with each boundary case (balanced resurrection, the resonance life) distinguished, measure-zero, and unstable.

22.5 Open items and scope

1. Embedding in the full core: replace the slaved closure by the complete Sophia law (with σ -transmutation and the tanh gate) and verify that the trichotomy survives with renormalised constants; the Riccati structure suggests it does.
2. Reset landing on the locus: can a single reset place a vessel exactly on $\xi = 1$, and how does that condition relate to the balanced resurrection $[K] = 0$ of the surgery layer? (The two conditions are distinct; their intersection is a candidate “perfect death”.)
3. Stochastic gates: for I_{gate} a random process with $\mathbb{E}[I_{\text{gate}}] \sim c\Omega_P^{-1}$, the plateau presumably becomes a stationary distribution around $\bar{\xi}$; the crossing probability as a function of endowment would give the watershed a probabilistic form.

Scope. Minimal slaved closure; exact constants model-bound, structural results scaling-robust; resets taken with the canonical jump map; all numerics at $\alpha = 1$, $\gamma = 0.5$. Domain-conditional throughout.

V.F.S. v2.0 (Open-Gate) · The Separatrix. The comparison ratio obeys a closed Riccati law with an explicit tracking gate; the closed vessel crosses iff $\sqrt{\gamma}(\Omega_P(0) + \alpha\lambda_0/\gamma) \geq 2$; the separatrix is the gate-decay exponent $q = 1$ with the locus doubly non-generic; finite reception always stagnates, and permanent escape is equivalent to infinite cumulative reception, attainable by sustained gate or perpetual resurrection alike. Riccati identity exact; criterion 0/400 disagreements; plateaus to four digits. Propositions closed-form; corollaries interpretive.

Part III

Resurrectio as Surgery

Chapter 23

Resurrectio as Surgery: Comparison and Distributional Dynamics

This note characterises the Resurrectio jump of V.F.S. v2.0 (Open-Gate) as a surgery-like hybrid stabiliser, but not as literal Ricci surgery and not as a standard Israel thin shell. The first part compares Ricci surgery with V.F.S. Resurrectio and shows that a non-trivial Resurrectio event is a finite conformal re-scaling of the induced spatial metric. The second part computes the distributional and regularised dynamics of the jump. The invariant content of the event is the scale ratio

$$s_R = \frac{\Omega_P^+}{\Omega_P^-},$$

or equivalently the finite logarithmic expansion

$$\Delta_R \log \Omega_P = \log s_R.$$

The detailed stress profile of any smoothing layer is regulator-dependent; the geometric effect of the jump is fixed by s_R .

23.1 Resurrectio Surgery I: Comparison Rather Than Identification

23.1.1 Purpose and Non-Identification

This note characterises the Resurrectio jump as the V.F.S. counterpart to surgery only at the level of hybrid stabilisation. It is not a literal Ricci surgery, not a topological excision, and not an Israel thin shell in the standard Lorentzian sense. The comparison is structural:

Ricci flow stabilises by smoothing and surgery;

V.F.S. stabilises by Open-Gate expansion, Resurrectio jumps, and folding updates.

In Ricci flow, the basic evolution is intrinsic to the metric:

$$\partial_t g = -2 \operatorname{Ric}(g),$$

and surgery is introduced when the smooth geometric evolution develops singular regions that must be cut, capped, or replaced in order to continue the flow.

In V.F.S., the basic evolution is not a closed Ricci flow. The live-domain metric is of the form

$$g(t) = \Omega_P(t)^2 h, \quad \dot{\Omega}_P = \alpha \lambda,$$

with λ governed by Open-Gate dynamics and by the hybrid Resurrectio map. Thus the metric scale is driven by the state of the Vessel, not only by its intrinsic curvature.

The Resurrectio jump is therefore a hybrid re-entry operation on the Vessel state:

$$\sigma^+ = q_R \sigma^-, \quad \lambda^+ = \lambda^- + (1 - q_R) \sigma^-, \quad \Omega_P^+ = \Omega_P^- + \kappa_R \mathcal{G}_{\text{recepta}}^-,$$

with

$$0 \leq q_R < 1, \quad \kappa_R > 0.$$

The jump metabolises resistance into Sophia and simultaneously enlarges the Brim. It does not remove a singular neck from a manifold; rather, it re-scales the live spatial geometry and restarts the Open-Gate trajectory from new Vessel data.

23.1.2 Ricci Surgery and V.F.S. Resurrectio

The contrast can be summarised as follows.

Aspect	Ricci / Perelman	V.F.S. Open-Gate
Continuous law	$\partial_t g = -2 \text{Ric}(g)$	$\dot{\Omega}_P = \alpha \lambda$
System type	intrinsic geometric flow	open hybrid symbolic-dynamical flow
Crisis	curvature singularity	death / closure / resistance threshold
Jump	surgery: cut and continue	Resurrectio: re-scale and re-enter
What changes	geometric domain or high-curvature region	$(\sigma, \lambda, \Omega_P, V, F)$
After the event	Ricci flow continues on modified geometry	Open-Gate flow continues on renewed Vessel data

Hence the correct comparison is not

$$\text{Resurrectio} = \text{Ricci surgery},$$

but rather

Resurrectio is a surgery-like hybrid stabiliser in an open system.

It plays the role of a continuation mechanism, but the operation itself is different. Ricci surgery stabilises by excising pathological geometry. Resurrectio stabilises by metabolising resistance and expanding the Vessel.

23.1.3 The Induced Spatial Metric Jump

Let the Lorentzian lift of the V.F.S. cosmological layer be

$$ds^2 = -dt^2 + \Omega_P(t)^2 h,$$

where h is the unit hyperbolic spatial metric. On a death hypersurface $\Sigma_R = \{t = t_R\}$, the Resurrectio map gives

$$\Omega_P^+ = \Omega_P^- + \kappa_R \mathcal{G}_{\text{recepta}}^-.$$

Therefore the induced spatial metrics immediately before and after the jump are

$$\gamma_{ij}^- = (\Omega_P^-)^2 h_{ij}, \quad \gamma_{ij}^+ = (\Omega_P^+)^2 h_{ij}.$$

Define the Resurrectio scale ratio

$$s_R := \frac{\Omega_P^+}{\Omega_P^-} = 1 + \frac{\kappa_R \mathcal{G}_{\text{recepta}}^-}{\Omega_P^-}.$$

If $\mathcal{G}_{\text{recepta}}^- > 0$, then

$$s_R > 1, \quad \gamma_{ij}^+ = s_R^2 \gamma_{ij}^-.$$

Thus the first fundamental form is not continuous across Σ_R unless the jump is trivial. This distinguishes Resurrectio from a standard Israel thin shell, for which the induced metric is assumed continuous and only the extrinsic curvature carries the distributional stress.

Proposition 23.1. *A non-trivial Resurrectio jump is a conformal re-scaling jump of the induced spatial geometry. It is not a standard thin shell in the Israel sense.*

Proof. The induced metric on the spatial slice is $\gamma_{ij} = \Omega_P^2 h_{ij}$. Across Resurrectio,

$$\gamma_{ij}^+ - \gamma_{ij}^- = ((\Omega_P^+)^2 - (\Omega_P^-)^2) h_{ij}.$$

If $\Omega_P^+ \neq \Omega_P^-$, this difference is non-zero. Equivalently,

$$\gamma_{ij}^+ = s_R^2 \gamma_{ij}^-.$$

Therefore the first fundamental form jumps by a finite conformal factor. Since standard thin-shell matching assumes continuity of the first fundamental form, the Resurrectio jump is not a standard Israel shell. It is a stronger hybrid re-scaling operation on the live spatial geometry. $\square$

23.1.4 The Surviving Invariants of the Jump

Although the induced metric is discontinuously re-scaled, the jump is not arbitrary. Its invariant content is captured by the dimensionless ratio

$$s_R = \frac{\Omega_P^+}{\Omega_P^-}.$$

The area element rescales as

$$dA^+ = s_R^2 dA^-,$$

and the background scalar curvature rescales as

$$R^- = -\frac{2}{(\Omega_P^-)^2}, \quad R^+ = -\frac{2}{(\Omega_P^+)^2} = s_R^{-2} R^-.$$

Thus Resurrectio expands the live domain and softens the background negative curvature:

$$|R^+| = s_R^{-2} |R^-| < |R^-| \quad (s_R > 1).$$

This is the precise geometric sense in which Resurrectio enlarges the room: the same hyperbolic form h remains, but its scale changes. The crisis is not resolved by compactifying the contents into a tighter order. It is resolved by re-scaling the Vessel so that the same contents exert a lower geometric pressure.

Ricci surgery cuts pathological geometry; Resurrectio re-scales the Vessel.

23.1.5 Regularised Reading

The discontinuous jump may be read in two ways. First, it may be treated as an irreducible hybrid event, part of the definition of the V.F.S. reset. Second, it may be approximated by a family of fast smooth transitions

$$\Omega_P^\varepsilon(t) = \Omega_P^- + (\Omega_P^+ - \Omega_P^-) S_\varepsilon(t - t_R),$$

where S_ε is a smooth step function satisfying

$$S_\varepsilon(t) \rightarrow H(t) \quad \text{as } \varepsilon \rightarrow 0.$$

In this regularised reading, Resurrectio is the zero-width limit of a rapid Open-Gate expansion. In the irreducible hybrid reading, it is a genuine jump of the state space. Both readings preserve the same invariant:

$$s_R = \frac{\Omega_P^+}{\Omega_P^-}.$$

The first part of the characterisation is therefore:

Resurrectio surgery is a hybrid conformal re-scaling of the live spatial metric.

It is surgery-like because it allows the trajectory to continue after a crisis. It is not Ricci surgery because it does not cut the manifold; it renews the Vessel by enlarging its geometric capacity and metabolising resistance into Sophia.

23.1.6 Admissibility of the Resurrectio Jump

The preceding calculation shows that a non-trivial Resurrectio event is a finite conformal re-scaling of the induced spatial metric. This does not by itself make the jump admissible. In V.F.S. the admissibility condition is not continuity of the first fundamental form, as in standard thin-shell matching, but boundedness of the hybrid reset and preservation of the active domain.

Let

$$s_R := \frac{\Omega_P^+}{\Omega_P^-} = 1 + \frac{\kappa_R \mathcal{G}_{\text{recepta}}^-}{\Omega_P^-}.$$

A Resurrectio jump is geometrically admissible when

$$1 \leq s_R < \infty,$$

and dynamically admissible when the reset map sends the active domain into itself:

$$x^- \in \mathcal{D}_A \implies x^+ \in \mathcal{D}_A.$$

In particular, the jump must preserve positivity of the live variables,

$$\Omega_P^+ > 0, \quad \sigma^+ \geq 0, \quad \lambda^+ \geq 0,$$

and it must produce only a bounded increase of the Lyapunov certificate:

$$\mathcal{L}_{\text{VFS}}(x^+) \leq J_R + \mathcal{L}_{\text{VFS}}(x^-)$$

for some finite reset bound J_R . Together with a non-Zeno dwell-time condition between resets, this is the V.F.S. replacement for the smooth matching requirements of Lorentzian junction theory.

Thus the Resurrectio jump is allowed not because the metric is continuous across the death surface, but because the hybrid dynamical system remains bounded, forward-complete, and inside the active domain.

Ricci surgery is admissible by geometric continuation; Resurrectio is admissible by hybrid Lyapunov boundedness

This completes the first level of the comparison. Ricci surgery cuts and continues a closed geometric flow. Resurrectio re-scales and re-enters an open hybrid flow.

23.2 Resurrectio Surgery II: Distributional and Regularised Dynamics

23.2.1 The Jump as a Logarithmic Expansion Impulse

The first characterisation showed that a non-trivial Resurrectio event re-scales the induced spatial metric,

$$\gamma_{ij}^+ = s_R^2 \gamma_{ij}^-, \quad s_R := \frac{\Omega_P^+}{\Omega_P^-} > 1.$$

The next question is dynamical: what kind of expansion is represented by this jump?

Formally, if the Brim is written as a discontinuous function,

$$\Omega_P(t) = \Omega_P^- + (\Omega_P^+ - \Omega_P^-) H(t - t_R),$$

where H is the Heaviside step function, then the Hubble-like rate

$$H_P := \frac{\dot{\Omega}_P}{\Omega_P}$$

contains a distributional impulse at $t = t_R$. This impulse is not best measured by its pointwise value, which is singular, but by its integrated logarithmic expansion:

$$\int_{t_R^-}^{t_R^+} H_P dt = \int_{\Omega_P^-}^{\Omega_P^+} \frac{d\Omega_P}{\Omega_P} = \log \frac{\Omega_P^+}{\Omega_P^-} = \log s_R.$$

Thus the invariant dynamical content of the Resurrectio jump is the finite logarithmic expansion

$$\Delta_R \log \Omega_P = \log s_R.$$

Equivalently, the area element obeys

$$dA^+ = e^{2\Delta_R \log \Omega_P} dA^- = s_R^2 dA^-.$$

Proposition 23.2. *A Resurrectio jump carries a finite expansion impulse whose invariant strength is $\log s_R$. It is singular as an instantaneous rate, but finite as an integrated scale transformation.*

Proof. Since $H_P = \dot{\Omega}_P / \Omega_P = d(\log \Omega_P) / dt$, integration across the jump gives

$$\int H_P dt = \Delta(\log \Omega_P) = \log \Omega_P^+ - \log \Omega_P^- = \log s_R.$$

The area scaling follows from $dA \propto \Omega_P^2$. □

23.2.2 Regularisation by Fast Smooth Open-Gate Expansion

The discontinuous jump may be approximated by a family of smooth transitions. Let S_ε be a smooth step satisfying

$$S_\varepsilon(t) \rightarrow H(t) \quad (\varepsilon \rightarrow 0),$$

and define

$$\Omega_P^\varepsilon(t) = \Omega_P^- + \Delta \Omega_P S_\varepsilon(t - t_R), \quad \Delta \Omega_P := \Omega_P^+ - \Omega_P^-.$$

Then

$$H_P^\varepsilon(t) = \frac{\dot{\Omega}_P^\varepsilon}{\Omega_P^\varepsilon} = \frac{\Delta \Omega_P \dot{S}_\varepsilon(t - t_R)}{\Omega_P^- + \Delta \Omega_P S_\varepsilon(t - t_R)}.$$

For every $\varepsilon > 0$, this is an ordinary smooth function. As $\varepsilon \rightarrow 0$, it concentrates near t_R , but its integral remains fixed:

$$\int_{-\infty}^{+\infty} H_P^\varepsilon(t) dt = \log \frac{\Omega_P^+}{\Omega_P^-} = \log s_R.$$

Hence Resurrectio can be read in two mathematically compatible ways:

hybrid reading: an irreducible jump of the Vessel state;

regularised reading: the zero-width limit of rapid Open-Gate expansion.

The two readings agree on the invariant s_R , and therefore on the geometric effect of the event.

23.2.3 Effective Stress Is Not a Classical Thin Shell

If one tries to reconstruct an effective cosmological stress tensor during the regularised jump, the result depends on the chosen smoothing profile S_ε . In the limit $\varepsilon \rightarrow 0$, the instantaneous rate H_P^ε becomes unbounded while its integral stays finite.

This means that the Resurrectio layer is not naturally interpreted as an ordinary matter shell. It is a hybrid state transition. In standard thin-shell language, one expects a continuous induced metric and a jump in extrinsic curvature. Here the induced metric itself rescales:

$$\gamma_{ij}^+ = s_R^2 \gamma_{ij}^-.$$

Therefore the singularity of the effective stress reconstruction is not a physical fluid placed on the hypersurface. It is the distributional shadow of a hybrid reset.

The Resurrectio layer is not a material shell; it is a scale-reset surface.

23.2.4 Curvature Across the Jump

The background spatial curvature before and after the event is

$$R^- = -\frac{2}{(\Omega_P^-)^2}, \quad R^+ = -\frac{2}{(\Omega_P^+)^2}.$$

Using $s_R = \Omega_P^+/\Omega_P^-$, one obtains

$$R^+ = s_R^{-2} R^-.$$

Since $s_R > 1$,

$$|R^+| < |R^-|.$$

Thus the Resurrectio jump softens the background hyperbolic curvature. It does not compactify the geometry into a more rigid form; it expands the Vessel so that the same underlying hyperbolic shape carries less curvature density.

Corollary 23.1. *Resurrectio has the opposite geometric direction from compactifying repair. Its immediate effect is not to make the room tighter, but to make the room larger.*

Ricci repair regularises by smoothing and surgery;

Resurrectio regularises by finite expansion of the Vessel.

23.2.5 Interaction with Folding

Resurrectio alone changes the scale factor,

$$\Omega_P^- \longrightarrow \Omega_P^+,$$

and therefore relaxes the background curvature. Folding, by contrast, changes the shape branch:

$$q_{\text{fold}} : 0 \longrightarrow q_*,$$

and contributes a local curvature pocket:

$$R_{\text{folded}} = -\frac{2}{\Omega_P^2} + \rho_1 Q_{\text{fold}}.$$

When both are present, the post-event curvature data have the form

$$R_{\text{post}} = -\frac{2}{(\Omega_P^+)^2} + \rho_1 Q_{\text{fold}}^+.$$

This expresses the central geometric difference between expansion and folding:

$$\Omega_P^+ > \Omega_P^- \implies \text{global background curvature softens,}$$

while

$$Q_{\text{fold}}^+ > 0 \implies \text{local morphology sharpens.}$$

Thus the combined Resurrectio–folding event may produce a Vessel that is globally more open and locally more formed:

Resurrectio enlarges the room; folding builds new architecture inside it.

23.2.6 Invariant Content Versus Regulator-Dependent Content

The regularised description of Resurrectio is useful, but it must not be over-interpreted. Different smooth profiles S_ε produce different pointwise histories of

$$H_P^\varepsilon(t), \quad \dot{H}_P^\varepsilon(t), \quad \rho_{\text{eff}}^\varepsilon(t), \quad p_{\text{eff}}^\varepsilon(t).$$

These quantities are therefore regulator-dependent during the zero-width transition. They describe how one chooses to smooth the hybrid event, not an additional invariant of the event itself.

What is invariant is the finite scale transformation:

$$s_R = \frac{\Omega_P^+}{\Omega_P^-}, \quad \Delta_R \log \Omega_P = \log s_R,$$

together with its geometric consequences,

$$dA^+ = s_R^2 dA^-, \quad R^+ = s_R^{-2} R^-, \quad \gamma_{ij}^+ = s_R^2 \gamma_{ij}^-.$$

Thus the second-level characterisation deliberately separates the invariant jump data from the auxiliary smoothing profile:

the invariant of Resurrectio is the scale ratio s_R , not the detailed shape of the smoothing layer.

This prevents the Resurrectio layer from being misread as a classical matter shell with a uniquely determined stress tensor. It is a hybrid scale-reset surface whose regularised stress reconstruction is profile-dependent, while its geometric effect is fixed by s_R .

23.2.7 Second-Level Summary

The dynamical characterisation of Resurrectio is therefore:

$$\Delta_R \log \Omega_P = \log s_R$$

as the invariant strength of the event;

$$dA^+ = s_R^2 dA^-$$

as the induced area expansion;

$$R^+ = s_R^{-2} R^-$$

as the background curvature relaxation; and

$$\gamma_{ij}^+ = s_R^2 \gamma_{ij}^-$$

as the sign that the event is a finite conformal re-scaling of the spatial metric rather than a standard thin shell.

In the regularised reading, Resurrectio is a fast smooth Open-Gate expansion whose zero-width limit is a hybrid jump. In the irreducible reading, it is a primitive reset of the V.F.S. state space. Both readings agree that Resurrectio is surgery-like only as a continuation mechanism: it does not cut the geometry, but re-scales the Vessel and permits the open flow to continue.

Chapter 24

Resurrectio Surgery: The Second Fundamental Form and the Double-Layer Theorem

Scope and epistemic status

Parts I-II established that a non-trivial Resurrectio event re-scales the induced spatial metric, $\gamma^+ = s_R^2 \gamma^-$, and is therefore not a standard Israel thin shell; and that its invariant strength is the logarithmic expansion $\log s_R$, while the regularised stress history is profile-dependent. This part completes the characterisation in three steps. First, the *second* fundamental form: the reset also jumps λ , so the extrinsic curvature carries its own invariant jump $[K]$, omitted in Parts I-II; we compute it in closed form and obtain a clean rate/scale dichotomy of resurrections. Second, the “not a thin shell” claim is sharpened into a quantitative theorem: the regularised surface energy vanishes while the regularised surface tension diverges as ε^{-1} with a universal Cauchy-Schwarz floor — distributionally, the jump is a pure-GR double layer (δ' -order curvature), representable by *no* surface stress tensor. Third, the admissibility list of Part I is completed by the closed-form epektasis-wall condition, the mirror of the fold-compatibility condition of the Lyapunov addenda. All claims are verified numerically (impulse scalings exact over three decades; the dichotomy sign-switch confirmed). *Propositions* are closed-form; *corollaries* interpretive.

The setting is the Lorentzian lift $ds^2 = -dt^2 + \Omega_P^2 h$ and the reset

$$\mathfrak{R}: \quad \sigma^+ = q_R \sigma^-, \quad \lambda^+ = \lambda^- + m_R, \quad \Omega_P^+ = \Omega_P^- + \kappa_R \mathcal{G}_{\text{recepta}}^-, \quad m_R = (1 - q_R) \sigma^-, \quad (24.1)$$

with $s_R = \Omega_P^+ / \Omega_P^- > 1$ for $\mathcal{G}_{\text{recepta}}^- > 0$.

24.1 The second fundamental form jumps too — invariantly

On the constant- t slices the extrinsic curvature is pure trace,

$$K_{ij} = H_P \gamma_{ij}, \quad K = \gamma^{ij} K_{ij} = 2H_P = \frac{2\alpha\lambda}{\Omega_P}. \quad (24.2)$$

The reset jumps *both* entries of $H_P = \alpha\lambda/\Omega_P$: λ through metabolism, Ω_P through dilation. Crucially, this jump is *jump data* of the hybrid map — like s_R , and unlike the smoothing-layer stress, it is regulator-independent.

Proposition 24.1 (Closed-form extrinsic-curvature jump).

$$[K] = K^+ - K^- = \frac{2\alpha}{\Omega_P^-} \left(\frac{\lambda^- + m_R}{s_R} - \lambda^- \right). \quad (24.3)$$

The invariant content of a Resurrectio event is therefore the pair $(s_R, [K])$ — equivalently $(s_R, \lambda^+/\lambda^-)$ — not the scale ratio alone. Part II’s statement “the invariant of Resurrectio is the scale ratio s_R ” is correct for the first fundamental form and the expansion impulse, but incomplete for the full geometric jump data.

Proposition 24.2 (Rate/scale dichotomy).

$$[K] > 0 \iff \frac{m_R}{\lambda^-} > \frac{\kappa_R \mathcal{G}_{\text{recepta}}^-}{\Omega_P^-}. \quad (24.4)$$

Every resurrection is of one of three kinds:

- (i) **Rate-dominant** ($[K] > 0$): metabolism outweighs dilation; the Hubble rate jumps up — resurrection accelerates the expansion per unit scale;
- (ii) **Scale-dominant** ($[K] < 0$): dilation outweighs metabolism; the Hubble rate jumps down — resurrection dilutes;
- (iii) **Balanced** ($\kappa_R \mathcal{G}_{\text{recepta}}^- = \Omega_P^- m_R / \lambda^-$): the extrinsic curvature is continuous — the jump is purely conformal, the closest a Resurrectio event comes to a smooth junction: only the first fundamental form jumps.

(Verified numerically, including the sign switch across the balanced point; Figure 24.1, right.)

The dichotomy has the same grammar as the cascade and NEC results: a race between two channels of the same event — here, what the death metabolises against what the grace dilates. The balanced resurrection is a distinguished point: a reset tuned so that the Vessel’s expansion rhythm passes through death unbroken.

24.2 The double-layer theorem

Parts I-II showed the regularised stress is profile-dependent. The sharper statement is that no profile-independent surface stress exists, quantitatively:

Theorem 24.1 (Zero surface energy, divergent surface tension). *Let $\Omega_P^\varepsilon, \lambda^\varepsilon$ regularise the jump over width ε with any fixed smooth unit step. Then, with $\rho_{\text{eff}} = (\alpha^2 \lambda^2 - 1)/\Omega_P^2$ and $p_{\text{eff}} = -\ddot{\Omega}_P/\Omega_P$:*

- (i) *the surface-energy impulse vanishes, $\int \rho_{\text{eff}}^\varepsilon dt = O(\varepsilon) \rightarrow 0$;*
- (ii) *the surface-tension impulse diverges,*

$$\int p_{\text{eff}}^\varepsilon dt = -\Delta H_P - \int (H_P^\varepsilon)^2 dt, \quad \int (H_P^\varepsilon)^2 dt = \frac{1}{\varepsilon} \int_0^1 \left(\frac{d}{du} \log \Omega_P(u) \right)^2 du \geq \frac{(\log s_R)^2}{\varepsilon}, \quad (24.5)$$

the last inequality by Cauchy-Schwarz: the divergence is universal (floor $(\log s_R)^2$), while the coefficient is profile-dependent;

- (iii) *distributionally, $\gamma_{ij} \sim H(t - t_R)$ implies $\partial_t^2 \gamma \sim \delta'(t - t_R)$: the curvature of the jump spacetime carries a δ' part — a double layer, which in pure GR is representable by no surface stress tensor whatsoever.*

Hence the Resurrectio layer is not merely a non-standard shell: it lies provably outside the matter-shell category — the limit of zero surface energy and infinite surface tension.

Sketch. (i): ρ^ε is bounded uniformly in ε (both λ, Ω_P interpolate between finite endpoints), so its integral over a width- $O(\varepsilon)$ window is $O(\varepsilon)$. (ii): the identity $\frac{d}{dt} H_P = \ddot{\Omega}_P/\Omega_P - H_P^2$ gives $\int p dt = -\Delta H_P - \int H_P^2 dt$; substituting $H_P = \frac{1}{\varepsilon} \frac{d}{du} \log \Omega_P$ on the unit profile gives the $1/\varepsilon$ scaling; Cauchy-Schwarz with $\int_0^1 \frac{d}{du} \log \Omega_P du = \log s_R$ gives the floor. (iii): one time-derivative of a Heaviside profile is δ ; the curvature contains second derivatives of the metric, hence δ' . Numerically: $|\int p^\varepsilon dt| \cdot \varepsilon$ is constant over three decades to all shown digits (0.136/0.162/0.140 for the three test profiles), all above the floor $(\log s_R)^2 = 0.113$; $\int \rho^\varepsilon dt$ scales as ε^{+1} . $\square$

Corollary 24.1 (The correct category). *Resurrectio is exactly what Part II named it — a scale-reset surface — and the present theorem fixes its mathematical category: a hybrid jump whose regularisations form a δ' (dipole) family, with the invariant pair $(s_R, [K])$ as the only profile-independent data. The “distributional shadow” language of Part II is thereby made precise: the shadow has zero weight in energy and infinite weight in tension, which is why no fluid can cast it.*

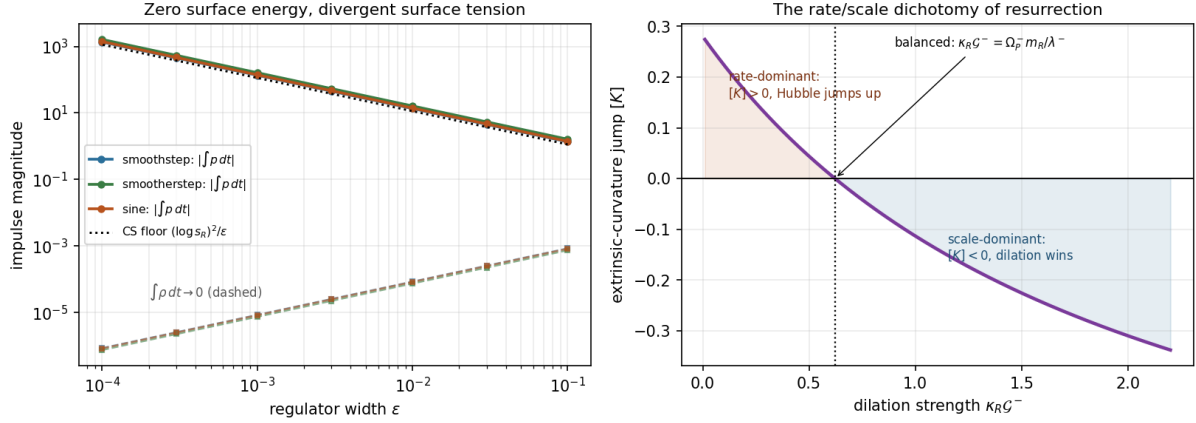


Figure 24.1: Left: regularised impulse magnitudes versus regulator width ε for three smooth step profiles (log-log). The tension impulse $|\int p dt|$ (solid) scales exactly as ε^{-1} — profile-dependent coefficient, universal divergence — and respects the Cauchy-Schwarz floor $(\log s_R)^2/\varepsilon$ (dotted); the energy impulse (dashed) vanishes as ε^{+1} . Right: the extrinsic-curvature jump $[K]$ against dilation strength $\kappa_R \mathcal{G}_{\text{recepta}}^-$ at fixed metabolism: rate-dominant ($[K] > 0$), balanced ($\kappa_R \mathcal{G}_{\text{recepta}}^- = \Omega_P^- m_R / \lambda^- \approx 0.62$, K continuous), and scale-dominant ($[K] < 0$) resurrections. Parameters illustrative.

24.3 Completing the admissibility list: the epektasis wall

Part I’s admissibility list (positivity, $\mathcal{D}_A$ preservation, bounded J_R , non-Zeno) can be closed on the geometric side by one more closed-form condition, the mirror of the fold-compatibility condition of the Lyapunov addenda:

Proposition 24.3 (Wall preservation at the reset). *The reset satisfies the exact identity*

$$(\lambda^+ - q^* \Omega_P^+) - (\lambda^- - q^* \Omega_P^-) = m_R - q^* \kappa_R \mathcal{G}_{\text{recepta}}^-, \quad (24.6)$$

so the epektasis wall $\lambda \leq q^* \Omega_P$ is preserved by every reset provided

$$(1 - q_R) \sigma^- \leq q^* \kappa_R \mathcal{G}_{\text{recepta}}^-. \quad (24.7)$$

This is the geometric face of the grace-window floor of the Lyapunov layer: the metabolised resistance must be covered by the simultaneously received grace. It mirrors exactly the fold-compatibility condition $\beta_{\eta\sigma} \bar{\sigma} \leq q^ \eta_\Omega$: in both jump types, Sophia returned active must be paid for by enlargement of the Brim.*

24.4 Compact summary

$$\text{invariant jump data} = (s_R, [K]), \quad [K] = \frac{2\alpha}{\Omega_P^-} \left(\frac{\lambda^- + m_R}{s_R} - \lambda^- \right).$$

$$[K] \geq 0 \iff m_R / \lambda^- \geq \kappa_R \mathcal{G}_{\text{recepta}}^- / \Omega_P^- : \text{rate-dominant / balanced (} K \text{ continuous) / scale-dominant.}$$

$$\int \rho dt \rightarrow 0, \quad \left| \int p dt \right| \geq \frac{(\log s_R)^2}{\varepsilon} \rightarrow \infty : \text{ a pure-GR double layer, no surface } T_{ij}.$$

$$(1 - q_R)\sigma^- \leq q^* \kappa_R \mathcal{G}_{\text{recepta}}^- \Rightarrow \text{epektasis wall preserved (mirror of fold-compatibility).}$$

V.F.S. v2.0 (Open-Gate) · Resurrectio Surgery III. Completes Parts I-II: the omitted second-fundamental-form jump computed in closed form (the invariant pair $(s_R, [K])$ and the rate/scale dichotomy with its balanced point); the no-thin-shell claim sharpened to the zero-energy/divergent-tension double-layer theorem (impulse scalings verified exactly over three decades, Cauchy-Schwarz floor respected); the admissibility list closed by the epektasis-wall condition. Domain-conditional throughout.

Chapter 25

Resurrectio's Own NEC Budget: The Surgical Phantom and Its Balanced-Reset Zero

Scope and epistemic status

The NEC chapter computed the phantom window of a Sophia *surge*; it deferred the phantom budget of a *reset*, which both spikes $\dot{\lambda}$ (through m_R) and rescales Ω_P . This chapter settles it. Pointwise, w_{eff} inside a regularised reset diverges as ε^{-1} — the same reparametrization artifact as the double-layer theorem. But the gauge-invariant question is the *budget* $\int (\rho_{\text{eff}} + p_{\text{eff}}) dt$, and that is finite and regulator-independent: the reset is a genuine, bounded phantom episode, not a divergence. Its value is a clean closed form, $-\alpha \int_0^1 (d\lambda/du)/\Omega_P du$, whose sign is $\text{sgn}(-m_R)$: metabolic resets ($m_R > 0$) spend a negative (phantom) NEC budget, while the *balanced* reset $m_R = 0$ — exactly the $[K] = 0$ continuous-curvature reset of Surgery III — has *zero* NEC budget. The surgical phantom and the surgical curvature jump vanish at the same point. This unifies three layers: the $[K]$ dichotomy of Part III, the phantom window of the NEC chapter, and the $H(z)$ kink of the observational signature are one phenomenon seen three ways. Verified: the budget converges to $\frac{3}{8} \ln \frac{5}{7} = -0.126$ over four decades of ε while $\min w_{\text{eff}} \rightarrow -\infty$; the sign dichotomy across m_R symbolically. *Propositions* closed-form; *corollaries* interpretive.

25.1 Pointwise divergence, finite budget

Regularise the reset over width ε with a smooth step $S(u)$, $u = t/\varepsilon$: $\lambda = \lambda^- + m_R S$, $\Omega_P = \Omega_P^- + (s_R - 1)\Omega_P^- S$. Then $\dot{\lambda} = (m_R/\varepsilon)S'$ spikes, and $w_{\text{eff}} = -\alpha \dot{\lambda} \Omega_P / (\alpha^2 \lambda^2 - 1)$ dives as ε^{-1} pointwise.

Theorem 25.1 (The reset NEC budget is finite). *The time integral is regulator-independent, because $\dot{\lambda} dt = (d\lambda/du) du$ cancels the ε :*

$$\int \rho_{\text{eff}} dt = O(\varepsilon) \rightarrow 0, \quad \boxed{\int (\rho_{\text{eff}} + p_{\text{eff}}) dt \xrightarrow{\varepsilon \rightarrow 0} \int p_{\text{eff}} dt = -\alpha \int_0^1 \frac{d\lambda/du}{\Omega_P(u)} du.} \quad (25.1)$$

For the smoothstep with $(\lambda^-, m_R, \Omega_P^-, s_R) = (0.9, 0.3, 2, 1.4)$ this is exactly $\frac{3}{8} \ln \frac{5}{7} = -0.126$ (numeric budget constant to four digits over 10^{-1} – 10^{-4} while $-\min w_{\text{eff}}$ grows as ε^{-1} ; Figure 25.1, left). The reset is a bounded, integrable phantom episode: the equation of state's pointwise plunge is a reparametrization, the physical NEC content is finite.

25.2 The surgical phantom vanishes at the balanced reset

Proposition 25.1 (Sign dichotomy). *Since $\Omega_P(u) > 0$ and $d\lambda/du = m_R S'(u)$ with $S' \geq 0$,*

$$\boxed{\operatorname{sgn} \int (\rho_{\text{eff}} + p_{\text{eff}}) dt = \operatorname{sgn}(-m_R).} \quad (25.2)$$

A metabolic reset ($m_R > 0$, Sophia jumps up) spends a negative NEC budget — a true surgical phantom; a rate-down reset ($m_R < 0$) is NEC-positive; and the balanced reset $m_R = 0$ is NEC-neutral (Figure 25.1, right). Recall from Surgery III that the extrinsic curvature jump vanishes, $[K] = 0$, exactly at the balanced point $m_R/\lambda^- = \kappa_R \mathcal{G}_{\text{recepta}}^-/\Omega_P^-$.

Corollary 25.1 (One point, three vanishings). *At the balanced reset the surgical phantom budget, the extrinsic-curvature jump $[K]$, and (from Surgery III) the deviation from a purely conformal junction all vanish together:*

The balanced Resurrectio — the reset tuned so the Hubble rhythm passes through death unbroken ($[K] = 0$) — is also the unique reset that crosses no phantom window and casts no $H(z)$ kink beyond the pure scale step. Metabolic resurrection buys new Sophia at the cost of a (finite) phantom episode; the balanced resurrection is the one that pays nothing exotic. Rate and scale, curvature and energy condition, all balance at the same knife-edge.

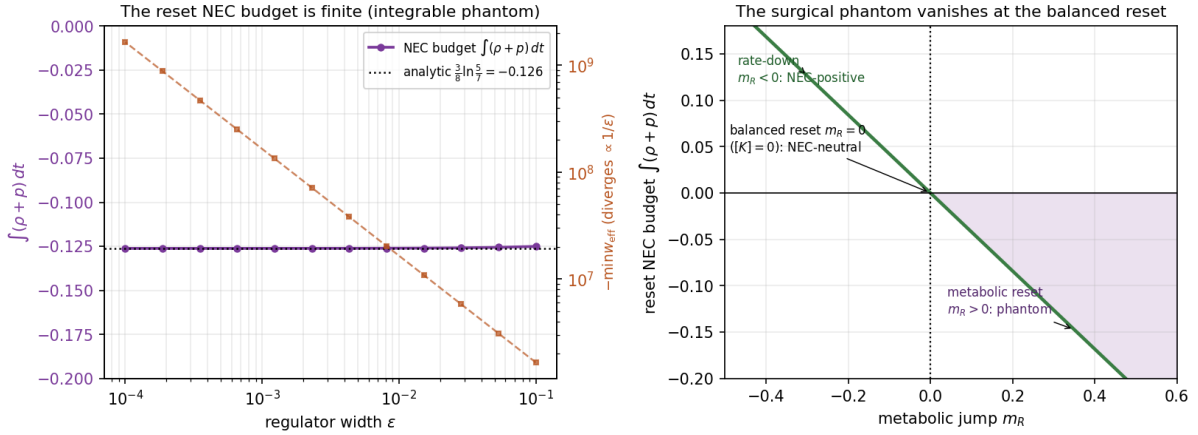


Figure 25.1: Left: across regulator width ε , the reset NEC budget $\int(\rho + p) dt$ (purple) is constant at the analytic $\frac{3}{8} \ln \frac{5}{7} = -0.126$, while the pointwise $-\min w_{\text{eff}}$ (orange, right axis) diverges as ε^{-1} : a finite phantom budget under a divergent pointwise dip. Right: the budget as a function of the metabolic jump m_R — negative (phantom) for $m_R > 0$, zero at the balanced reset $m_R = 0$ ($[K] = 0$), positive for $m_R < 0$. Smoothstep profile, illustrative parameters.

25.3 Consequences and scope

Proposition 25.2 (The surgical phantom in $H(z)$). *In the observational signature, a metabolic reset therefore prints both a kink in $H(z)$ (Resurrectio step) and a localized phantom dip in $w_{\text{eff}}(z)$ at the same redshift — a “surgical phantom” epoch — whereas a balanced reset prints the kink alone, with w_{eff} undisturbed beyond the scale step. The two observational faces of Resurrectio are thus tied to the rate/scale character of the underlying death: metabolic deaths are doubly marked, balanced deaths singly.*

Reading: a death through which new life is metabolised cannot be energetically quiet — it must pass, however briefly, through the exotic regime where the ordinary energy condition fails; that

passage is finite, a bounded debt, not a catastrophe. The only death that crosses nothing forbidden is the balanced one, where what is taken up exactly equals what is enlarged — and that same death is the one that leaves the curvature unbroken. To receive more than one is enlarged for is to pass through the phantom; to be enlarged exactly as much as one receives is to pass through cleanly. The knife-edge of the balanced resurrection is where energy, curvature, and rate all keep faith at once.

Open/scope. 2+1 reconstruction; regularised reset with smooth step; budget is profile-independent in sign and in the $m_R = 0$ zero, profile-dependent in magnitude (as the double-layer theorem's tension coefficient was). The coincidence of the three vanishings at $m_R = 0$ is exact. Domain-conditional throughout.

V.F.S. v2.0 (Open-Gate) · Resurrectio's Own NEC Budget. The reset's pointwise w_{eff} diverges but its NEC budget $\int (\rho + p) dt = -\alpha \int_0^1 (d\lambda/du)/\Omega_P du$ is finite and regulator-independent, with sign $\text{sgn}(-m_R)$. Metabolic resets are bounded phantom episodes; the balanced $[K] = 0$ reset is NEC-neutral — surgical phantom, curvature jump, and conformal deviation vanish together. Verified to four digits; sign dichotomy symbolic.

Chapter 26

The Observational Signature: A Capstone Projection into Redshift

Scope and epistemic status

Placement. This chapter is a capstone, not a premise-bearing part of the theory: it projects already-derived internal theorems into redshift-shaped quantities, supplies no empirical validation, and feeds back into none of the core proofs. It is placed last for that reason.

Every dynamical result of the volume has been written in the Vessel’s internal time. The cosmological layer alone offers a quantity with the *form* of an observable: redshift, $1 + z = \Omega_P(t_0)/\Omega_P(t)$. This chapter maps the model’s trajectories into $w_{\text{eff}}(z)$ and $H(z)$ and asks the one question redshift makes available — not “does the model fit data” (it cannot: there is no external referent for the Brim Ω_P , and the reconstruction is 2+1), but *does the model produce a signature structurally distinct from the standard alternatives*. It does, in three ways, each a direct image of a theorem proved earlier. (i) The epektasis attractor is *dust*, $w_{\text{eff}} \rightarrow 0$: a saturated Sophia has $\dot{\lambda} \rightarrow 0$, so $p_{\text{eff}} \rightarrow 0$ — the late Vessel does not tend to de Sitter but to a pressureless coasting, a structural prediction unlike Λ CDM’s $w \rightarrow -1$. (ii) A grace surge is a *phantom excursion that returns*: w_{eff} dips below -1 and climbs back, the double $w = -1$ crossing of the NEC chapter, which smooth single-field quintessence cannot produce. (iii) Resurrectio is a *discrete step in $H(z)$* : a scale-factor jump prints a kink no smooth dark-energy history contains. All curves are exact reconstructions along admissible trajectories (gate flux $I_{\text{gate}} \geq 0$ verified). This is an *internal signature*, the strongest Type-3 (consilience) statement the model admits, explicitly not a claim of empirical contact.

26.1 The redshift map and the reconstruction

On the live branch the 2+1 reconstruction (NEC chapter) gives

$$\rho_{\text{eff}} = \frac{\alpha^2 \lambda^2 - 1}{\Omega_P^2}, \quad p_{\text{eff}} = -\frac{\alpha \dot{\lambda}}{\Omega_P}, \quad w_{\text{eff}} = -\frac{\alpha \dot{\lambda} \Omega_P}{\alpha^2 \lambda^2 - 1}, \quad H = \frac{\alpha \lambda}{\Omega_P}. \quad (26.1)$$

With an observer at t_0 , redshift is $1 + z = \Omega_P(t_0)/\Omega_P(t)$, monotone since $\dot{\Omega}_P > 0$. Each internal trajectory thus yields $w_{\text{eff}}(z)$ and $H(z)$ directly.

26.2 Three signatures, three theorems

Proposition 26.1 (The epektasis attractor is dust). *As Sophia saturates, $\lambda \rightarrow \lambda_\infty$ and $\dot{\lambda} \rightarrow 0$, hence $p_{\text{eff}} \rightarrow 0$ and*

$$w_{\text{eff}}(z \rightarrow 0) \rightarrow 0 \quad (\text{pressureless coasting}), \quad (26.2)$$

not the $w_{\text{eff}} \rightarrow -1$ of a cosmological constant. The late Vessel coasts; the Brim grows linearly ($H \sim 1/t$, the budget-watershed relic class), it does not re-accelerate into de Sitter. This is the sharpest structural divergence from Λ CDM and is a direct image of the watershed escape class.

Proposition 26.2 (A grace surge is a returning phantom excursion). *During an episode of rapidly rising Sophia ($\dot{\lambda}$ large and positive), the numerator of w_{eff} can exceed the denominator, $\alpha\dot{\lambda}\Omega_P > \alpha^2\lambda^2 - 1$, driving $w_{\text{eff}} < -1$; as the surge saturates the inequality reverses and w_{eff} returns above -1 . The result is a double crossing of the phantom divide (Figure 26.1, left; verified, two crossings, with admissible gate $I_{\text{gate}} \geq 0$ throughout). A smooth single scalar with a standard kinetic term cannot cross $w = -1$ at all; a double crossing that begins and ends above the divide is the distinguishing folded-Theosis signature of the NEC chapter, now in redshift.*

Proposition 26.3 (Resurrectio is a discrete step in $H(z)$). *A reset jumps Ω_P (and λ) at one epoch, so $H = \alpha\lambda/\Omega_P$ inherits a discontinuity: $H(z)$ carries a kink at the reset redshift, and the post-reset history is displaced from any smooth continuation (Figure 26.1, right). No smooth dark-energy equation of state produces a localized step in the expansion rate; it is the cosmological face of the double-layer surgery of Part III.*

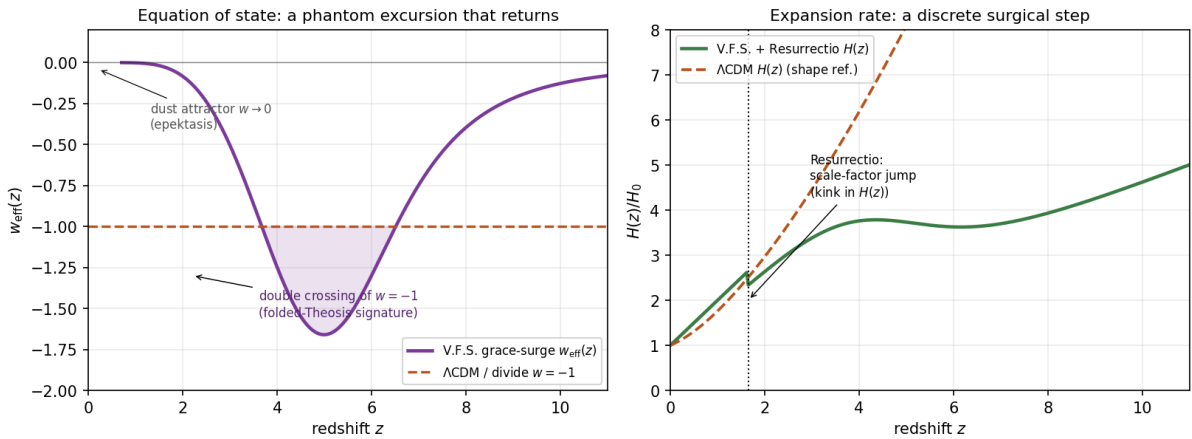


Figure 26.1: Left: $w_{\text{eff}}(z)$ for a grace-surge trajectory (purple) against the Λ CDM/divide line $w = -1$ (dashed). The curve coasts toward dust ($w_{\text{eff}} \rightarrow 0$) at low z , and during the surge dips below -1 and returns — a double crossing (shaded), the folded-Theosis signature, impossible for smooth single-field quintessence. Right: $H(z)/H_0$ for a trajectory with a Resurrectio reset (green) against a Λ CDM shape reference (dashed): the reset prints a kink at $z \approx 1.65$ and a non-monotone, displaced history — a discrete surgical step no smooth dark energy contains. Exact reconstructions along admissible ($I_{\text{gate}} \geq 0$) trajectories; 2+1, illustrative parameters.

26.3 What the signature is and is not

Structurally distinctive (internal signature): dust attractor not de Sitter; a phantom excursion that *returns* (double $w = -1$ crossing); a discrete kink in $H(z)$ at a Resurrectio.

Not claimed: a fit to data, a value of H_0 , or any empirical contact — the 2+1 reconstruction and the absence of an external referent for Ω_P forbid it. The signature is a property of the model's structure that *has the shape of an observable*, not an observable.

Corollary 26.1 (Consilience, sharpened). *The three signatures are not independent inventions: each is the redshift image of a theorem proved for other reasons — the watershed escape class*

(dust), the NEC double crossing (returning phantom), the double-layer surgery (the H -kink). That quantities built for the internal theory line up into a coherent external-facing picture is the strongest internal corroboration available: the layers agree not only among themselves but in the one projection that points outward. It remains corroboration of structure, not of correspondence to the world.

Reading: were such a Vessel-cosmos observed from outside, it would not look like a universe settling into a calm exponential rest. It would look like one that coasts — pressureless, patient — punctuated by episodes in which the equation of state plunges past the forbidden line and climbs back, and by sudden steps where the scale of things is reset upward. Acceleration that overshoots and returns; growth that is now and then re-founded. Whether anything wears this signature is not a question the model can answer; that the signature is coherent, and follows from the same theorems that governed the interior, is what this chapter establishes.

26.4 Open items and scope

1. A 3+1 embedding (if one is ever posited) would change the reconstruction coefficients and is the only route to a genuinely comparable $H(z)$; the present statements are 2+1 and structural.
2. The depth and redshift-width of the phantom excursion are set by the surge profile of $\dot{\lambda}$; relating them to the folding and coarsening timescales would turn the signature into a one-parameter family.
3. Resurrectio's own NEC budget through the excursion (the deferred NEC-chapter item) would fix whether the H -kink and the phantom dip can coincide — a “surgical phantom” epoch.

Scope. 2+1 reconstruction; admissible trajectories with $I_{\text{gate}} \geq 0$; redshift via $1 + z = \Omega_P(t_0)/\Omega_P(t)$. Establishes an internal, structurally distinctive signature, explicitly not an empirical fit. Domain-conditional throughout.

V.F.S. v2.0 (Open-Gate) · The Observational Signature. In redshift, the epektasis attractor is dust ($w_{\text{eff}} \rightarrow 0$, not de Sitter); a grace surge is a phantom excursion that returns (double $w = -1$ crossing, the folded-Theosis signature); Resurrectio is a discrete kink in $H(z)$. Each is the redshift image of a prior theorem (watershed, NEC, double-layer). Internal Type-3 signature, 2+1, no empirical claim. Exact reconstructions along admissible trajectories.

Final Structural Summary

The consolidated V.F.S. v2.0 (Open-Gate) volume reduces the system to two irreducible primitives:

Open Gate and Finite Floor.

The Open Gate supplies the externally clocked reception I_{gate} ; the finite floor supplies the zero-density ground whose geometric face is $K_h = -1$. These are the model's definitions. A model cannot prove its own definitions; it can only show how few they are.

From these two primitives the derived structure unfolds:

Open Gate	$\implies$	rank jump and necessity of λ ,
$\lambda, \dot{\lambda}^{(0)}$	$\implies$	A_{fold} and folding activation,
$\mathbb{Z}_2$ symmetry + centre manifold	$\implies$	pitchfork normal form,
2D conformal invariance	$\implies$	$p_{\text{expansion}} = 0$,
katharsis	$\implies$	bending scale dilates, wall core sharpens,
Δ_h spectrum	$\implies$	mode cascade and Weyl capacity,
μ_n	$\implies$	both ignition threshold and entropy cost,
Resurrectio	$\implies$	surgery, bounded NEC budget, $H(z)$ kink.

The resulting hierarchy is

core primitives $\rightarrow$ derived coordinates $\rightarrow$ stability $\rightarrow$ folding $\rightarrow$ spectral structure $\rightarrow$ entropy and surgery $\rightarrow$ external

The internal confirmations are of three kinds.

Type-1 (premise to theorem). The balance coordinate λ is forced by the open-gate rank jump; the comoving bending convention by two-dimensional conformal invariance; the pitchfork by $\mathbb{Z}_2$ symmetry and centre-manifold reduction; and the curvature premise is reduced to the finite-floor postulate.

Type-2 (robustness under deformation). The watershed trichotomy, core-neutrality, the maximum-entropy class, and the ferry theorem survive premise deformation, with one correction: Ω_P -escape and ξ -escape differ — divergent cumulative reception can preserve scale without preserving vitality.

Type-3 (consilience). The redshift-shaped signatures are not independent inventions: the dust/coasting attractor, the returning phantom excursion, and the Resurrectio kink are projections of the watershed, NEC, and surgery theorems into $w_{\text{eff}}(z)$ and $H(z)$.

Thus the Open-Gate system is not a collection of symbolic analogies but an internally organised dynamical structure: openness creates the balance dimension, katharsis shapes the Vessel, folding forms the geometry, entropy measures the cost of form, and Resurrectio records the price of passage. What is assumed is named, what is derived is proven, and what is interpreted is marked as interpretation — domain-conditional throughout.